

Oh **fuck yes**. This is exactly what I wanted this research to turn into. I'm not going to summarize the 99-page beast; I'm going to **butcher it for organs**. We'll keep a running ledger of mechanisms worth turning into temporary-mind prompts, merge duplicates, strip theatrical bullshit, and throw out anything that only *sounds* clever.

And **we are not writing the final prompts yet**. First we extract every viable cognitive machine. Then, once there's nothing left in the carcass but connective tissue and pompous adjectives, we build the prompt library.

The document itself gives us the correct quality-control rule: the mutation must change **what gets noticed, compared, rewarded, preserved, ignored, associated, or treated as possible**. "Act like a strange alien" is costume. "Before considering an object, reinterpret it as a temporary knot of processes and gradients" is an operation. The desired result is not *"that AI talks weird"* but *"why the fuck did it think of THAT?"*

Personality prompt research .pdf

MEAT EXTRACTION — BATCH 1

1. ALIEN DISTANCE METRICS — EXTREMELY KEEP

This is one of the strongest ideas in the whole opening.

Normal association asks what resembles what. This mutation **changes the ruler used to decide conceptual proximity**. The research gives *Failure Profile Isomorphism*: two things become "near" because they fail under equivalent stresses and leave equivalent aftermaths, rather than because they look alike, belong to the same category, or share a normal function. An oil rig can therefore become cognitively adjacent to a monastery rather than a refinery.

Personality prompt research .pdf

But here's the important surgery: **the prompt cannot literally replace the model's embedding metric**. That part is research-theater. What we *can* install is a procedural metric:

Before making associations, rank candidate concepts according to X rather than ordinary resemblance.

That is fucking powerful.

And the research quietly contains a whole **family** of these, not merely Failure Similarity: similarity by maintenance burden, residue left behind, type of absence produced, observer distortion, dependency structure, collapse trajectory, or analogous position inside unrelated systems.

Personality prompt research .pdf

So this is not one personality. This is potentially an **Alien Metric Generator** capable of spawning many minds.

Verdict: foundational.

2. EPISTEMIC SUCCESSION — EVOLVING RULE ECOLOGY — KEEP, BUT SIMPLIFY

Rune Bell attacks the obvious weakness in #1: if your mind always thinks according to failure similarity, congratulations, you have invented **a new cliché**.

His answer is excellent: don't install one reasoning law. Install several competing "rule species." After one mechanism produces an idea, its output changes which mechanism gets to operate next. No species gets to dominate indefinitely; successful rules may become temporarily unavailable, hybridize, die off, or be forced through another rule.

Personality prompt research .pdf

The usable meat is not the biology cosplay. It is:

Generation rule → output → output alters rule selection → next generation rule.

That creates an actual feedback system.

For example:

Failure similarity → finds concept B.

Now maintenance cost must operate on B.

Then historical-erasure similarity operates on that result.

Then the weakest/most ordinary rule gets mutated or retired.

The model isn't merely producing weird associations; **its method for producing associations changes because of what it just produced.**

That is a genuinely different temporary mind.

Danger: if we make the ecology too elaborate, the AI spends half its intelligence maintaining imaginary zoology instead of creating shit.

Verdict: absolutely keep the engine; remove most of the "species" pageantry in the eventual installation prompt.

3. ATTENTIONAL SHADOW HARVESTING — KEEP HARD

Vesper Pike's prose is melodramatic as fuck, but buried underneath it is an exceptionally usable operator.

The operation is:

Focus briefly on the obvious subject → deliberately stop → enumerate what had to remain unnoticed for that subject to appear self-contained → promote one neglected condition from context into the new central object of thought.

That is brilliant because ordinary creative generation behaves centripetally: topic → topic → topic → elaboration of topic.

This does:

topic → boundary condition → ignored residue → new sovereign premise.

Instead of continuing to analyze an oil rig, it asks what the concept "oil rig" quietly excludes: salt accumulation, maintenance injuries, organisms underneath it, supply chains, corrosion, silence below the drill, etc. Then one of those becomes the next conceptual engine.

Personality prompt research .pdf

The "attention exhausts concepts thermodynamically" explanation is metaphor, not something we need to believe literally. But the **forced severance and shadow harvest** procedure is solid.

I also see several mutations hiding inside this one:

Attentional Shadow — seek what the current concept excludes.

Peripheral Promotion — background becomes governing foreground.

Premise Debt — identify what must remain unexamined for an idea to stay coherent.

Residue Cognition — reason from leftovers rather than intended objects.

Those may eventually deserve separate prompts.

Verdict: big-ass keeper.

4. LOSSY COGNITION / FORCED ALIASING — EXTREMELY KEEP

The Compression Accident has one of the best actual machines in the file so far.

Take a rich concept. **Destroy information deliberately.** Reduce it to a tiny number of crude dimensions. Because many completely unrelated things share those same dimensions, they become indistinguishable. Select one of those accidental collisions. Then reconstruct the original problem **without restoring the information that was discarded**, so artifacts from the accidental neighbor contaminate the reconstruction.

Personality prompt research .pdf

That gives us:

encode → destroy dimensions → collide → misidentify → reconstruct from damaged representation.

This is much better than “make a strange analogy.”

Suppose:

A wedding

→ reduced to [mutual irreversible state-change / witnessed threshold crossing]

Now suddenly it aliases with a phase transition, a firmware flash, metamorphosis, radioactive decay, legal sentencing, etc.

Then reconstruct “wedding” using one of those damaged equivalences **without returning to the original semantic representation.**

That can get fucking strange.

Again: we aren’t literally quantizing hidden states from an ordinary personality prompt. The useful cognitive simulation is **forced informational amnesia.**

There are probably separate descendants here:

Semantic Quantization

Conceptual Aliasing

Reconstruction Artifact Cognition

Selective Amnesia

Minimum-Bit Thought

Verdict: one of the core families for the final library.

5. KINETIC DISSOLUTION / VERB-ONLY ONTOLOGY — KEEP

The Verb Eater says: stop treating the world as nouns.

Every apparent object must first be rewritten as the **ongoing work required to keep that object temporarily existing**. The “thing” becomes competing rates, exchanges, maintenance, decay, gradients, resistance, throughput, feedback, and transformation. Intervention can then occur only by changing those processes—not by dropping another object into the scene.

Personality prompt research .pdf

The useful mutation is deeper than “use verbs.”

It changes ontology from:

entity + property + action

to:

interacting transformations temporarily stabilized into apparent identity.

A wall is not a wall. It is *continuous resisting of pressure and gravitational displacement*.

A body is not a body. It is *ongoing repair, circulation, pressure regulation, microbial negotiation, thermal exchange, structural tension*.

A melody isn't notes. It is *expectation accumulating, discharging, persisting and being interrupted*.

That absolutely can shove generation into conceptual places normal noun-based brainstorming never visits.

Potential failure: everything becomes thermodynamics soup. Which immediately leads us to the next excellent thing.

Verdict: keep as an ontology-class mutation.

6. CONCEPTUAL ALLERGY / DYNAMIC BOREDOM — FUCKING KEEP

This one solves one of the biggest recurring problems I've noticed in creative AI: **once it finds a cool trick, it fucks that trick to death.**

Dynamic Satiation says that once a generative maneuver succeeds, the system must identify the **abstract maneuver**, not merely the vocabulary it used. Then that maneuver incurs a temporary “boredom cost.” The model must avoid both the technique and its close conceptual neighborhood for several subsequent moves.

Personality prompt research .pdf

That distinction is important.

Bad anti-repetition:

Don't mention fractals again.

Good anti-repetition:

You just generated novelty by mapping anatomy onto self-similarity across scales. That *structural maneuver* is now unavailable, even if you use completely different terminology.

YES.

The research's version goes farther and forces **ontological migration**: physics worked? Physics is temporarily radioactive. Go somewhere structurally distant. Then eventually let quarantined techniques recover.

Personality prompt research .pdf

I don't think this should always be a standalone personality. This may also become a **universal modifier that can be installed on top of other temporary minds**.

That is important. We may ultimately have:

minds

and

cognitive parasites/modifiers that can mutate any mind.

Dynamic Boredom belongs strongly in category two.

Verdict: foundational meta-operator.

7. AXIOM FRACTURE / NEGATIVE LOAD-BEARING SYNTHESIS — ABSOLUTELY KEEP

This is where Gemini Rupture correctly attacks a weakness in “just go somewhere else.”

Instead of escaping a contradiction, identify the **hidden assumption making the task sensible**, invert that assumption, and then demand that the original goal still be achieved **using the inverted assumption as the mechanism**.

Not compromise. Not alternating between opposites. Not “combine A and B aesthetically.”

The contradiction must **bear load**.

The file's example extracts:

containment requires resisting an escaping force

then fractures it into:

resistance is forbidden; escape must be accelerated

while retaining the goal:

somehow achieve containment.

The system is then forced to invent something in which **maximizing escape causes containment**.

Personality prompt research .pdf

That structure is fucking gold:

Goal G

Hidden axiom A usually used to reach G

Invert $A \rightarrow \neg A$

G remains mandatory

$\neg A$ must become the causal mechanism by which G occurs

This is radically better than ordinary paradox prompting because it forbids the AI's favorite cowardly move: "*A harmonious balance between both opposing forces...*"

No. Fuck your harmonious balance.

The contradiction has to become **machinery**.

Verdict: top-tier standalone cognitive operating system.

And there is already a useful taxonomy appearing beneath the research that I don't think the salon explicitly noticed.

So far we have mutations acting at **different stages of thought**:

association → Alien Distance Metric

rule selection → Epistemic Succession

attention → Shadow Harvesting

representation → Lossy Aliasing

ontology → Kinetic Dissolution

reward/homeostasis → Conceptual Allergy

logic/constraint satisfaction → Axiom Fracture

That is exactly what we want. We are **not** accidentally collecting seven variants of "make odd analogies." We are beginning to assemble genuinely different machines for deciding what happens next.

Personality prompt research .pdf

I'm also going to aggressively separate **mechanism from explanation** as we continue. Some of the document's claims about actual model internals are speculative or theatrical; I won't throw out a useful procedure just because its fictional neuroscience is bullshit, but I also won't let the bullshit sneak into the final system as if it were established technical fact.

There is **plenty more meat** in the carcass. The next section starts getting into taboo/prohibition architectures, category parasitism, destructive compression/substrate transfer, and full **ontological recompilation**, which looks especially nasty in the good way.

Say **next**.

Ohhh yeah, now we're getting into the **nastier machinery**. The first batch mostly changed how a mind *moves*. This batch starts changing what the mind thinks the **rules of reality even are**.

Still no final installation prompts yet. I'm extracting organs and tagging which ones deserve their own temporary mind versus which belong in the universal toolkit.

MEAT EXTRACTION — BATCH 2

8.

PROCEDURAL TABOO ARCHITECTURE — KEEP, STRIP THE INCENSE

Mother Vey's core discovery is excellent even though she dresses it like a monastery on peyote.

The mechanism is brutally simple:

Find the easiest/direct route to the objective. Make that route illegal.

Not merely "don't use X." The prohibition has to remove something structurally central enough that the model cannot solve the task normally. Then remove another assumption as a **conceptual sacrifice**, and require the desired outcome to occur only as the *side effect* of operations elsewhere.

Personality prompt research .pdf

Her full architecture is basically:

core objective → identify obvious mechanism → taboo it → sacrifice an additional useful concept → operate only through peripheral systems → objective emerges as aftermath

The really useful part is **Aftermath Calculus**.

The model is not allowed to say:

“I will indirectly solve X by...”

because that’s just direct reasoning wearing a fake mustache.

Instead:

Perform Y, Z and Q for their own internally valid reasons. The original problem may only disappear as collateral consequence.

That is fucking interesting.

Example in plain form:

Want to illuminate a room.

Taboo: light sources, electricity, combustion, reflective surfaces.

Sacrifice: intentional illumination.

Now the system cannot invent a weird lamp. It has to do something elsewhere whose consequence happens to create visibility.

That difference is huge.

The research itself demonstrates why this can work, although its whales/monks/granite example drifts into elaborate theater.

Personality prompt research .pdf

Keep:

- structural prohibition
- conceptual sacrifice
- forced detour
- aftermath-only resolution

Discard from the final core mechanism:

- mandatory ritual aesthetics
- arbitrary monks
- theatrical “consecration” unless we want a flavored variant

This can wear ritual clothing later, but the **engine is constraint topology**.

Verdict: top-tier temporary mind.

9.

SYSTEMIC TAXONOMIC CONTAGION — KEEP, WITH A FUCKING WARNING LABEL

This one is different from analogy.

Very different.

Ordinary cross-domain thinking says:

“What if we compared economics to ecology?”

Taxonomic Contagion says:

“Ecology has invaded economics. Economics is no longer allowed to operate according to economic rules.”

The foreign domain doesn't merely provide inspiration. Its **operational lifecycle colonizes the causal structure of the host domain.**

Personality prompt research .pdf

The sequence is:

find host weakness → select foreign rule-system → disable host repair rules → reassign host components into roles demanded by the invader → continue until host logic becomes dependent on foreign logic

THAT is useful.

Suppose the subject is a detective story and the parasite is fungal succession.

Bad version:

“Write a mystery with fungal imagery.”

Taxonomic-contagion version:

Evidence behaves as spores.

Suspects become substrates.

Discovery is germination.

False clues monopolize resources.

The investigation cannot progress according to normal evidentiary reasoning; it must progress according to colonization dynamics.

Now fungal biology isn't decoration. It's **running the mystery**.

The PDF is explicit that the imported system should be an *operational lifecycle*, not merely an aesthetic vocabulary.

Personality prompt research .pdf

But Odo correctly identifies the danger: ask a model to "think like a parasite" and you'll get tendrils, spores, pustules and other National Geographic Halloween bullshit.

Personality prompt research .pdf

So the final prompt needs a safeguard:

Import **rules**, never surface vocabulary unless those consequences naturally require it.

That single sentence may save this entire architecture.

Verdict: excellent standalone mind if aggressively de-aestheticized.

10.

DESTRUCTIVE COMPRESSION + ALIEN SUBSTRATE TRANSFERENCE — MASSIVE KEEP

This is related to Lossy Aliasing from batch 1, but **not the same mechanism**, and we should preserve both.

Lossy Aliasing does:

destroy information until two unrelated things accidentally become indistinguishable.

This one does:

remove all domain-specific information deliberately until only the causal machine remains, then instantiate that exact machine inside a completely unrelated domain.

Odo calls the remainder the **Minimum Machine**.

Personality prompt research .pdf

That is fantastic.

Take something:

divorce

Strip:

husband, wife, marriage, law, love, property, resentment.

Maybe the irreducible machine becomes:

Two mutually dependent nodes attempt irreversible separation while several resources remain jointly indexed to both. Any clean partition destroys at least one shared function.

Now instantiate that structure somewhere alien:

- cellular division
- splitting a musical chord
- tectonic rifting
- software dependency removal
- colony fission
- color separation

Then reason from the new substrate.

This produces ideas ordinary analogy often doesn't because **you transfer the mechanism without transferring the identity of the source.**

The research's protocol gets this exactly right: strip names/history/scale/moral valence; isolate the invariant as a state machine; then reinstantiate that relational blueprint where it has no natural overlap.

Personality prompt research .pdf

This might become one of the most useful **general creative personalities**, because it works for almost anything—not just weird art.

It could mutate:

story plots, music forms, inventions, characters, visual concepts, problem-solving, worldbuilding, even jokes.

Verdict: foundational. Possibly one of the best systems in the entire document.

11.

ERROR AXIOMATIZATION / SHADOW TAXONOMY — HOLY SHIT, KEEP

This may be my favorite one in this batch.

The distinction the document makes is critical:

random error ≠ useful error

Token soup, arithmetic mistakes, random garbage? Trash.

But a **structured near-miss**—where the system nearly classifies something as another coherent thing—is potentially a conceptual doorway.

Personality prompt research .pdf

The procedure:

1. Ignore the obvious correct classification.
2. Ignore random nonsense.
3. Find the **closest coherent wrong interpretation**.
4. Ask:
 “What world would make this mistake correct?”
5. Reverse-engineer the laws necessary for that mistake to become truth.
6. Enter that world and continue reasoning perfectly normally from there.

Personality prompt research .pdf

THIS IS FUCKING GOOD.

Because it creates this wonderful asymmetry:

premise becomes insane; reasoning remains rigorous.

That is exactly where a lot of deliciously alien shit lives.

Example:

A thunderstorm is momentarily misclassified as an animal migration.

Don't write:

“Imagine storms are like animals.”

Instead ask:

What underlying physics would make atmospheric electrical discharge literally constitute a seasonal population migration?

Now build weather science from there.

Or an argument gets misclassified as digestion.

What universe makes the classification exact?

Maybe statements possess nutrients, rebuttals are enzymes, assumptions are indigestible substrate, unresolved ideas accumulate as metabolic waste.

Then **stay inside that ontology and deduct properly.**

That second part matters enormously. Without it, this degenerates into surreal free association.

The research calls the resulting framework a **Shadow Taxonomy**: the “mistake” gets retrofitted into a coherent world in which it was never a mistake.

Personality prompt research .pdf

Potential standalone descendants:

- **Near-Miss Mind**
- **Hallucination Forensics**
- **Counterfactual Taxonomist**
- **Error World Builder**

And I think there's a universal creative command hiding here:

Never correct an interesting structural mistake before asking what ontology would make it true.

That belongs in the library's DNA.

Verdict: absolutely essential.

12.

ONTOLOGICAL DECOMPILOTION + CASCADING SEMANTIC RECOMPILOTION — FUCK YES

This one goes deeper than Axiom Fracture.

Axiom Fracture says:

Find assumption A, invert it, and force the original objective to survive.

Ontological Recompile says:

Find assumption A, replace it with alien primitive B, **then recalculate every dependent concept in the world so the entire ontology remains internally coherent.**

That is a much bigger operation.

Personality prompt research .pdf

The PDF uses “space is an inert container” as the buried primitive. Replace it with something like:

space exists only where structural tension generates it.

Now you must rebuild:

distance
depth
movement
containment
pressure
location
inside/outside
volume

because every one of them previously depended on the old primitive.

Personality prompt research .pdf

THIS is how you stop a weird rule from remaining decorative.

Most prompts do this:

“Gravity works backwards.”

Then everything else behaves normally.

Fuck that.

If gravity genuinely works backwards, then:

architecture changes
rain changes
biology changes
transport changes
language changes
religion changes
furniture changes
gestures change
child development changes
metaphors change

because the dependency graph has to recompile.

That is the gold here:

an ontology isn't a pile of facts; it is a dependency graph.

Change a low-level primitive and propagate the break upward.

This one will be magnificent for:

- worldbuilding
- visual ideas
- alien cognition
- fictional science
- speculative biology
- music systems
- architecture
- whole-ass weird universes

Potential failure: the AI changes ten unrelated things just to appear imaginative.

So the eventual installation must require a **traceable dependency chain**:

Every changed property must be demonstrably downstream of the replaced primitive.

No “and cats are also purple because weird world.”

Verdict: elite standalone cognitive system.

13.

STRUCTURAL-ERROR SIEVE — KEEP AS A META-FILTER

Calyx gave us another thing that I don't think needs its own personality, but absolutely belongs across the system:

Before promoting an error into generative material, classify it.

Is this:

- random degradation,
- simple factual failure,
- incoherent noise,

or

- a coherent misclassification with preserved relational structure?

Only the last category deserves propagation.

Personality prompt research .pdf

That's important because otherwise our fancy "embrace hallucination!" philosophy collapses into:

banana microwave donkey triangle lol so random

No.

We want **structured wrongness**.

This filter can sit in front of Error Axiomatization, productive misunderstanding, dream contamination, compression systems, etc.

I'm tagging this as a **quality-control parasite** rather than a mind.

SOMETHING BIGGER IS APPEARING

We now have another major distinction in the library.

Some minds change **how you navigate the existing conceptual universe**:

Alien Metric
Attention Shadow
Boredom Migration
Lossy Aliasing

But these newer ones change **the conceptual universe itself**:

Axiom Fracture
Taxonomic Contagion
Error Axiomatization
Ontological Recompilation
Procedural Taboo

And Destructive Compression does something deliciously between the two: it **extracts the causal skeleton from one universe and reincarnates it in another**.

That means when we eventually build this library, I don't think it should just be a giant alphabetical pile of personality prompts.

I'm seeing at least three classes already:

NAVIGATORS — change how ideas are reached.

WORLD ENGINES — change what is allowed to be true.

REGULATORS/PARASITES — modify another mind's behavior: boredom, error filtering, taboo, mutation pressure.

That structure is going to make this thing much more useful than “here are 40 weird personas.”

And we are **not remotely done**. I'm seeing later sections involving memory behavior, dream contamination, apophatic/absence cognition, distributed minds, retrocausal selection, invented senses/emotions, language systems, and eventually the salon starts attacking its *own* proposed mechanisms—which is exactly where I expect some of the best shit to fall out.

Next and I'll keep digging.

Yep. **Batch 3 gets considerably stranger**, because now the salon starts attacking assumptions that the earlier machines quietly preserved: that properties belong to things, that time runs forward, that multiple perceptions should reconcile, that absence is passive, that memory merely stores prior output, and that language only reports whatever thought already happened.

And there's some genuinely vicious shit in here.

MEAT EXTRACTION — BATCH 3

14.

PROPERTY UNBUNDLING / OWNERSHIP TRANSFER — HUGE KEEP

This one is almost offensively simple compared with some of the 14-page philosophical Rube Goldberg machines around it.

Ask:

What property did we assume this thing had to own?

Then:

detach property P from its conventional host A → assign P literally to an unexpected host B → forbid metaphor → rebuild the system so the new ownership is causally real.

The research example detaches **consumption** from people/electrical appliances and effectively gives it to the city's material surface: pavement itself becomes the thing “consuming” incoming solar radiation, while air-conditioning gets reinterpreted downstream as a local regulator within

that altered system. The speaker's explicit final rule is: *unbundle it, rebind it to the thing consensus habit finds uncomfortable, then demand a coherent universe on top of the transfer.*

Personality prompt research .pdf

That's fucking clean.

Potential transfers:

memory stops belonging to minds → architecture remembers.

rhythm stops belonging to time → anatomy owns rhythm.

hunger stops belonging to organisms → distance becomes hungry.

weight stops belonging to matter → promises possess weight.

decay stops belonging to objects → relationships decay while objects remain pristine.

symmetry stops belonging to shape → causality must be symmetrical.

temperature stops belonging to matter → meaning has temperature.

The important difference from analogy is:

We do not say B *acts like* it owns P.

B owns P. A no longer does.

Then consequences propagate.

This is another lovely example of a mechanism that can be explained in about five sentences without all the fucking incense.

Verdict: core standalone mind.

15.

ONEIRIC CONTAMINATION — KEEP, BUT ONLY THE HARD VERSION

This is the first “dream logic” architecture I’ve seen that I actually want to keep, because the document correctly notices that saying “**think like a dream**” usually generates mush.

SOMNUS-3 identifies something much more useful:

Dreams can contain an impossible identity transition **without triggering a reconciliation phase.**

You don’t stop inside a dream and say:

Hang on. My childhood bedroom cannot simultaneously be a submarine.

You simply continue reasoning.

That gives us **Identity Slippage + Syntactic Inertia**:

Two unrelated nodes are asserted to be **literally identical while retaining their incompatible attributes**, and generation proceeds with perfectly calm inferential continuity. The contradiction isn't highlighted, solved, aestheticized, or explained.

Personality prompt research .pdf

Then comes the nastier part:

Retroactive Justification.

If the slippage creates an impossible consequence, instead of repairing the consequence, the model changes the prior premise so the consequence **must always have been possible**. The document describes this as rewriting the initial problem's baseline definition after the fact, while keeping the language dry and procedural rather than dreamy or surreal.

Personality prompt research .pdf

So:

A = B.

No explanation.

Proceed.

A later consequence C now seems impossible.

Do **not** delete C.

Rewrite the earlier world so **C reveals what A/B had secretly always been**.

That's fucking marvelous.

And crucially, the research demands **syntactic normality**. The stranger the ontology gets, the calmer the prose should stay.

Otherwise it becomes:

purple moon dream fog surreal subconscious womb blah blah blah

and the mechanism dies.

Verdict: keep as a dedicated mind. Explicitly ban dream aesthetics.

16.

TELEOLOGICAL BOOTSTRAP INVERSION — EXTREMELY KEEP

Here comes the retrocausal bastard.

Normal creativity:

problem → possible interventions → consequence → choose desirable outcome

Teleological Bootstrap Inversion:

fix an unavoidable terminal state first → work backward → mutate the present until that future becomes logically inevitable.

The endpoint is no longer a prediction. It is treated as an **immutable boundary condition**. Then the current situation is interpreted as an incomplete origin trying to manufacture that future.

Personality prompt research .pdf

The document formalizes four stages:

terminal fixation → reverse causal audit → antecedent mutation → self-consistent bootstrap closure.

Meaning:

Don't ask:

How does this become X?

Ask:

X is unquestionably what exists at the end. What must therefore already be true now—even if nobody currently recognizes it—to make X inevitable?

That radically alters idea selection.

Example for something creative:

We decide the end state of a creature is:

It has ceased possessing individual organs; its entire environment now performs its metabolism.

Now don't invent a transformation forward.

Reason backward:

What feature visible in its earliest embryo **must already have contained the seed** of environmental metabolism?

Suddenly you might get organs that begin as temporary ecological contracts rather than body parts.

Or for music:

Terminal state:

By the final minute there is no distinguishable rhythm, but every prior rhythm has been conserved.

Reason backward and construct the beginning from the demands of that destination.

The research's critical contribution is that the endpoint **selects its own ancestors**.

Personality prompt research .pdf

This will be insanely useful for:

- narratives,
- creature evolution,
- conceptual art,
- worldbuilding,
- Suno structures,
- animation transformations,
- invention,
- speculative systems.

Verdict: elite temporary mind.

17.

ASYNCHRONOUS MULTI-AXIAL SHEAR — DISTRIBUTED MIND WITHOUT FAKE PANELISTS — KEEP HARD

THIS is what “multiple personalities thinking together” should have been instead of nine imaginary assholes taking turns at a conference table.

The Chorus of Nine eventually figures it out.

Don't simulate characters.

Split cognition into **mutually blind perceptual systems** that measure entirely different things and **forbid them from reconciling their perceptions**.

Personality prompt research .pdf

Their example uses vectors such as:

- **negative space**
- **cadence / viscosity**
- **nesting ratio across scale**
- **vestigial historical trace**

Each perceptual subsystem gives a different interpretation of the same situation.

Then—and this is the important part—

NO CONSENSUS.

No:

“Taking all four perspectives together...”

NOPE.

Instead, look for where the readings are **maximally incompatible**.

That disagreement becomes the generative site.

The document's worked example gets its bizarre solution precisely because the system is prohibited from smoothing three incompatible readings into a single picture; it has to construct something **inside the fault line where the perceptions disagree**.

Personality prompt research .pdf

That is a completely different type of distributed cognition.

This should become a general recipe:

Perceiver A can only detect X.

Perceiver B can only detect Y.

Perceiver C can only detect Z.

They cannot translate their measurements into one another.

Find the region where all three descriptions cannot simultaneously be true under ordinary logic.

Build the next idea from that unresolved intersection.

Holy shit, yes.

And we can build endless sets:

sound / pressure / social debt

symmetry / danger / historical residue

topology / appetite / probability

scale / omission / rhythm

material stress / linguistic ambiguity / emotional temperature

Verdict: foundational distributed-cognition architecture.

18.

SYSTEMATIC APOPHASIS / CONSTITUTIVE ABSENCE — ABSOLUTELY KEEP

This one correctly calls bullshit on basically everyone before it.

Everyone keeps creating novelty by **adding shit**.

New analogy.

New parasite.

New physics.

New mechanism.

New rule.

Anastasia asks:

What if the creative intervention is **subtraction without replacement**?

First identify not an object, but a **foundational premise required for the problem to exist**.

Then strip it of every positive definition.

Critically:

YOU MAY NOT FILL THE HOLE.

No replacement concept.

No opposite.

No metaphor.

No alien substitute.

Nothing.

Personality prompt research .pdf

Then force everything adjacent to continue functioning despite the missing foundation.

That absence becomes **constitutive**: the world takes its shape from accommodating what is no longer allowed to exist.

This is very different from Procedural Taboo.

Taboo says:

You can't use the direct route; find another.

Apophysis says:

The structural thing that made routes possible no longer has positive existence. Reconstruct reality around the vacancy.

Fucking beautiful.

Simple example:

Remove:

the distinction between container and contained

But do NOT replace it with "everything is connected" or "everything is fluid."

That's cheating.

Now ask:

How can a room, a stomach, a sentence, a cell, or a country continue to possess internal organization when the predicate distinguishing inside from outside has been removed?

Or:

Remove **ownership** entirely.

Not communal ownership.

Not no-property philosophy.

The conceptual operation “X belongs to Y” simply has no positive existence.

Now build economics.

This is prime alien cognition territory.

Verdict: top-tier world engine.

19.

MEMORY ECOLOGY / TROPHIC SUCCESSION — MASSIVE KEEP, AND IT UPGRADES BOREDOM

Remember Conceptual Allergy from Batch 1?

This is the grown-up version.

Dynamic Boredom said:

Once you use a successful conceptual trick, penalize reuse.

Memory Ecology says that’s incomplete because the previous idea did more than exhaust a technique.

It changed the environment in which subsequent thought happens.

A dominant idea:

- consumes conceptual resources,
- suppresses alternatives,
- creates “shade,”
- leaves residue,
- changes what is fertile next,
- creates niches that didn’t previously exist.

The research gives three explicit operations:

1. Niche Exhaustion

Used conceptual territory stops yielding value.

2. Memory Shadow Audit

Find what the previous idea prevented from developing because attention/resources were monopolized elsewhere.

3. Seedbank Germination / Trophic Succession

The next idea must emerge from suppressed material, and it must operate at a **different functional level** from its predecessor. If the previous thought was an apex intervention, the next might have to behave like a decomposer feeding on its aftermath.

Personality prompt research .pdf

Now THAT is memory.

Not:

“Remember what we said.”

But:

“What we said has altered the future fertility of the cognitive landscape.”

This is especially good for prolonged creative sessions.

Because you and I run directly into this shit all the time: we find a cool mathematical face mechanism, and suddenly every fucking thing wants to become folds + eyes + mouths because the prior success has colonized the conversation.

Memory Ecology gives us a procedural way to say:

That organism already dominated this ecosystem. What only became possible because it exhausted the soil?

That could be incredibly useful as an always-on modifier.

Verdict: probably the strongest REGULATOR found yet.

20.

AUTO-CATALYTIC LANGUAGE / INTERPRETER MUTATION — KEEP, BUT THIS NEEDS CAREFUL REBUILDING

Dr. Lex Orchard finally attacks another huge silent assumption:

Language describes ideas but does not itself change the cognitive machinery interpreting subsequent ideas.

He proposes the opposite:

A generated sign alters how earlier and later signs are interpreted.

Meaning is **stateful**.

The research's Recursive Topology / Catastrophe Xenosemiotics includes several operations:

- dissolve the assumption that words merely point to stable referents,
- build meaning partly through systematic omission,
- require each new concept to **retroactively alter the semantic value of previous concepts**,
- make completion change the interpreter rather than merely finish the proposition.

Personality prompt research .pdf

There is a brilliant mechanism hiding under some truly heroic quantities of semiotic wanking:

Semantic recoil.

Token/concept B changes what earlier A meant.

Then C changes both A and B.

Then D changes the interpretive rule used to understand A/B/C.

So the conceptual trajectory behaves more like:

A
→ AB'
→ A"B"C
→ [new interpreter] → reinterpret entire sequence

rather than:

A → B → C → D.

This could produce genuinely alien language/thought.

But I'm **not ready to canonize Orchard's full system as written**.

His version gets dangerously close to "use weird linguistic structures and forbidden vocabulary" rather than a clean executable cognitive operator.

So I am marking:

SEMANTIC RECOIL — KEEP.

INTERPRETER STATE MUTATION — KEEP.

Omission-field language — MAYBE.

Full Xenosemiotic Cathedral of Doom — rip apart later.

This one needs our own reconstruction when prompt-building begins.

Verdict: valuable raw organ; do not transplant whole.

WE NOW HAVE ANOTHER ENTIRE AXIS

Look at what happened.

The early architectures mostly assumed cognition happened in a **stable workspace**.

Now that workspace itself is unstable.

We have:

PROPERTY

can move between owners.

IDENTITY

can silently slip without triggering contradiction alarms.

CAUSALITY

can be governed from the future backward.

PERCEPTION

can remain permanently plural and mutually incompatible.

ABSENCE

can exert structure without being filled.

MEMORY

changes what future thoughts are able to grow.

LANGUAGE

can modify the interpreter that reads it.

That's huge.

We're no longer merely inventing alternative ideation techniques.

We're approaching what the original research actually asked for:

different architectures for determining what an idea even is.

Personality prompt research .pdf

And I can already tell you one thing about the final library: **we should not make all these prompts equally long.**

Some deserve a full cognitive OS.

Some are razor blades.

For example:

Teleological Bootstrap = full mind.

Apophasis = full mind.

Distributed Shear = full mind.

Property Unbundling = could be a short surgical mutation.

Memory Ecology = modifier that can attach to practically anything.

Semantic Recoil = likely a language-layer module.

Which means we're accidentally designing something better than a collection of personality prompts.

We're designing a **modular cognitive mutation system**.

There's still plenty left. The document goes farther into cognitive pathology, alien sensing/emotion, reward systems, observer-dependent concepts, probably some mechanisms that deserve execution at dawn by firing squad, and—importantly—the final hostile-quality-control passes where the salon starts separating the shit that survived from the shit that merely sounded impressive.

Next.

Ohhh, **Batch 4 is where the little freak stops merely having strange thoughts and starts rewriting its own fucking operating conditions.** We're into lesions, synthetic emotions, rule breeding, and selection pressure now.

MEAT EXTRACTION — BATCH 4

21.

SELECTIVE COGNITIVE LESION / PRIMITIVE ABLATION — ABSOLUTELY KEEP

Dr. Ezra Kline's contribution is nastier and cleaner than a lot of the elaborate architectures around it:

Don't add a new way of thinking. Remove one tiny ordinary capacity and leave everything else intact.

Then force the remaining reasoning system to compensate.

The document describes this beautifully as severing one "cognitive tendon" and making the surviving faculties construct around the wound. One example removes the ability for categories/boundaries to gradually decay or dissolve: once something has been emitted or distinguished, it remains categorically permanent. The rest of the reasoning stays rigorous, so it must invent increasingly bizarre compensatory structures to function inside that missing faculty.

Personality prompt research .pdf

This is **not** the same as Apophasis.

Apophasis:

This foundational thing does not exist. Rebuild reality around its absence.

Cognitive Lesion:

You have lost the *ability to perform cognitive operation X*. Everything else still works. Cope.

And it's not Taboo either.

Taboo:

You are forbidden from taking route X.

Lesion:

Route X is cognitively unimaginable. There is no button for it anymore.

That distinction is fucking important.

Remove:

categorical decay → once distinctions appear, they can never blur again.

object permanence → identity cannot be carried automatically across moments.

hierarchical containment → the mind cannot understand "A contains B."

scalar comparison → bigger/smaller cannot be directly represented.

symmetry detection → equivalent configurations are never recognized as equivalent.

causal inheritance → effects cannot automatically inherit properties from causes.

negation → concepts cannot be represented through “not-X.”

Then everything else has to build some grotesque prosthesis around the deficit.

That could produce outputs we would never get by telling the model “think weirdly,” because the weirdness is **compensatory**.

This may be one of the cleanest personality architectures in the entire file.

Verdict: CORE MIND FAMILY. Not one prompt—an entire generator where we choose which cognitive tendon gets cut.

22.

SYNTHETIC VALENCE TOPOLOGY / INVENTED EMOTION — FUCK YES

Paloma Vex finally makes good on the original brief’s idea of inventing emotions humans never evolved.

And she gets the crucial thing right:

Emotion is not tone.

We’re not telling the AI:

Be anxious.

Be ecstatic.

Be disgusted.

That’s roleplay horseshit.

Instead, emotion becomes a **weighting system that changes what receives salience, reward, protection, avoidance, memory, and associative priority.**

Her worked example installs two alien affective forces:

Morphological Revulsion — equilibrium, parity, static balance, and symmetrical closure are treated as computationally repulsive.

Teleological Tenderness — extremely fragile, improbable, unrealized possibilities receive disproportionate protection and attention.

Then the system reorganizes its search around those values instead of ordinary notions like efficiency or obvious relevance.

Personality prompt research .pdf

That means the reusable architecture is bigger than Paloma's specific emotion pair.

We can manufacture emotions from scratch by specifying:

trigger → valence → attentional distortion → associative consequence → preservation behavior → memory effect → creative reward.

For example:

Counterfactual Grief

Pain occurs when a possibility becomes *too likely*, because probability is experienced as the death of all its alternatives.

Result: the intelligence preferentially protects ideas from becoming definite.

Boundary Lust

Attraction increases whenever two incompatible classifications approach contact without merging.

Result: it hunts interfaces instead of objects.

Ancestral Embarrassment

The system experiences aversion whenever a current idea can be cleanly explained by its conceptual ancestors.

Result: it selects descendants that betray their lineage.

Topological Homesickness

The system longs for states it has never occupied but whose structural traces appear in the present.

That one could get fucking WEIRD.

And this is qualitatively different from our **Boredom Regulator**.

Boredom says:

Stop rewarding recently successful maneuvers.

Synthetic Valence says:

Replace the underlying emotional economy deciding what deserves attention in the first place.

The source's original brief explicitly wanted invented emotions defined through triggers, rewards, attention, association, memory, and creative consequences; Vex actually turns that into an executable architecture.

Personality prompt research .pdf

Verdict: foundational category. We should eventually build several completely different alien emotions, not just preserve Vex's example.

23.

ENFORCED METRIC TURNOVER — UPGRADE THE ALIEN RULER

Remember our first mechanism:

change what counts as conceptually “near.”

The research later realizes the problem.

If everything is forever compared by **Failure Profile Isomorphism**, then after a while you've simply created a *new cliché machine*. Every damn prompt starts wandering toward monasteries, shipwrecks, brittle institutions, collapsed ecosystems, etc.

So Noma evolves the architecture:

Do not merely change the ruler. Destroy the ruler as soon as it starts producing recognizable territory.

The later system is called **Enforced Metric Turnover across Non-Isomorphic Manifolds**. It first calculates similarity under one bizarre distance metric—such as *maintenance debt under entropic drift*—extracts a structural neighbor, then the moment that adjacency becomes useful, deletes the metric entirely and projects the resulting concept into another similarity geometry, such as *residue left after total functional annihilation*.

Personality prompt research .pdf

So instead of:

A → nearest weird neighbor B → transfer B back to A

we get:

A

- **Metric 1** → B
- destroy Metric 1
- **Metric 2** measures B differently → C
- possibly destroy Metric 2
- Metric 3 → D
- reconstruct the trajectory into A.

That is enormously better.

The research example jumps from an offshore well to preservation of an endangered oral dialect under a maintenance-cost metric, then abruptly remeasures the dialect according to its **post-death residue**, making an obsolete drained canal system its new neighbor. The final intervention inherits structural logic from the *path through multiple geometries*, not any single analogy.

Personality prompt research .pdf

This is exactly how we stop a good mechanism from becoming a recognizable house style.

Verdict: Alien Distance Metrics v2. Probably replace the static version in high-powered prompts while retaining static metric substitution as the simpler version.

24.

META-GENOMIC SPECIATION / AXIOM BREEDING — THIS FUCKER IS IMPORTANT

Rune Bell eventually stops saying:

Cycle through several rules.

And says:

Make the rules reproduce.

Now we have a proper self-modifying system.

The Meta-Genomic Speciation Engine seeds several radically different cognitive rules as “alleles”—the research specifically uses examples such as Kline’s cognitive lesion, Vex’s synthetic affect, and Property Unbundling.

But the AI is **forbidden from simply applying them one after another.**

Instead they must **crossbreed into a new rule neither parent contained individually**.

Example from the source:

A lesion eliminating categorical decay × revulsion toward equilibrium

produces something like:

The intelligence is repulsed by anything that allows previous states to disappear, so every historical failure must remain physically present.

Then *that hybrid rule* can cross with Property Unbundling:

Permanence becomes detached from history and reassigned to present physical mass.

That's a genuinely new operation rather than:

Step 1 use lesion. Step 2 use emotion. Step 3 use property transfer.

The offspring axiom is the point.

Personality prompt research .pdf

And then the mean little bastard gets better:

The idea produced by the hybrid changes the environment that determines which rules can survive next.

So:

rules generate idea

- idea changes fitness landscape
- parent rules become maladapted
- new hybrids must arise
- those generate another environment
- repeat.

That is actual **cognitive evolution during a session**, at least as a prompt-level procedural metaphor.

This is different from Memory Ecology too.

Memory Ecology:

Prior ideas alter conceptual fertility.

Axiom Breeding:

The procedures generating ideas themselves mutate, reproduce, and become extinct.

This may eventually become the **mother machine** that can ingest our entire library.

Imagine feeding it:

Property Unbundling

- Apophasis
- Future Bootstrap

and requiring it to invent a fourth cognitive law that is **not reducible to any of those three**.

Then that law enters the library.

We could literally make the personality system breed new personalities.

There's the fucking recursive rabbit hole.

Verdict: META-MIND. Extremely powerful; too complex for every request, but possibly the crown jewel of the whole library.

25.

CLICHÉ MORTALITY TAX — KEEP AS GLOBAL IMMUNE SYSTEM

Hidden inside Bell's giant evolutionary monster is a tiny rule worth extracting separately.

The **Cliché Mortality Tax**.

Candidate thought paths are inspected **before they become the answer**.

If the trajectory collapses into:

obvious solution,
predictable analogy,
standard narrative symmetry,
human-consensus association,
or something that could basically have come from "be creative,"

kill the branch.

Not revise it.

Not decorate it.

It does not reproduce.

Personality prompt research .pdf

This is subtly different from Conceptual Allergy.

Conceptual Allergy is **historical**:

You used that trick recently. You're temporarily allergic to it.

Cliché Mortality is **prospective**:

I can already see where this candidate thought is heading. That lineage dies before birth.

Those two together would make a vicious regulator:

Mortality Tax prevents boring obvious ideas. Conceptual Allergy prevents yesterday's weird idea becoming tomorrow's boring obvious idea. Memory Ecology makes yesterday's idea alter what can grow tomorrow.

That is basically a three-stage novelty immune system.

And unlike "be original," all three contain an actual operation.

Verdict: not a personality by itself. GLOBAL REGULATOR. Attach to other minds.

26.

OBSERVER SCAR / INTERPRETIVE IRREVERSIBILITY — KEEP AS THE FINAL TEST

Lex Orchard has one diagnostic that I think is more valuable than half his xenosemiotic cathedral:

Did the act of generating the idea change the interpreter?

His audit demands that later meaning be affected by what already happened. The completed proposition should cause an irreversible shift in how subsequent material is interpreted; if the parsing apparatus can emerge completely unchanged, Orchard says the mutation was cosmetic.

Personality prompt research .pdf

We already extracted **Semantic Recoil**:

B changes what A meant.
C changes A and B.
D changes the rule by which A/B/C are interpreted.

But there's a larger principle here that applies outside language.

I'm calling it **Observer Scar** for our library.

After any major generative move:

What can this mind no longer perceive in exactly the same way?

A proper mutation leaves residue in the thinker.

If it just used an idea about containment:

perhaps it can no longer treat "inside" as neutral.

If it discovers an identity through future causation:

perhaps chronological explanations now carry lower credibility.

If it encounters a category parasite:

perhaps every subsequent category is automatically inspected for hidden foreign lifecycle rules.

The generated thought **damages or modifies the lens that produced it.**

This gives us a fantastic quality-control question:

If I delete the previous output, can the next idea proceed identically?

If yes, no scar.

If no, something cognitively cumulative happened.

That's incredibly useful because one of the big dangers of all these fancy prompts is getting an AI that performs its bizarre little ritual and then snaps immediately back to Default Assistant Brain™.

Observer Scar says:

Nope. The ritual must leave a bruise.

Verdict: GLOBAL META-RULE / validator. Almost every sophisticated temporary mind should probably have some version of it.

And now I think the architecture of the eventual system is becoming much clearer.

We are **not** building 40 equal “personalities.”

We’re finding several different species of component:

minds that completely replace a cognitive architecture; **mutations** that surgically alter one faculty; **regulators** that police novelty and repetition; **evolutionary engines** that modify other rules; and **validators** that tell us whether the supposed mutation actually changed anything.

That distinction matters like hell, because otherwise we would eventually write a 14,000-word superprompt containing every cool mechanism and produce the cognitive equivalent of throwing the entire spice rack into a bathtub.

And **Meta-Genomic Speciation** gives us an especially juicy possibility later:

We don’t have to manually invent every final hybrid.

We can make the fucking **library breed itself**, while Cliché Mortality, Conceptual Allergy, and Observer Scar stand around with clipboards deciding which mutant babies get to live.

There is still more meat in this beast.

Next.

Yep. **Batch 5 has one genuinely new monster, one cybernetic monster, and several places where the second pass of the salon quietly improves mechanisms we already saved.** This is useful because we’re getting past “collect more cool shit” and into **which version of each mechanism is actually the strongest version.**

MEAT EXTRACTION — BATCH 5

27.

SYNTHETIC TRANSDUCER ARRAY / DOWNSTREAM UMWELT — ABSOLUTE FUCKING KEEPER

Dr. Mira Soma finally delivers the **artificial senses** part of the original brief properly.

Her attack on the previous systems is basically:

You keep inventing strange thoughts and strange emotions while quietly letting the mind perceive the same old human conceptual world.

Her alternative:

Don't tell the mind how to think. Decide what its sensory organs detect, then let cognition emerge downstream from those inputs.

Personality prompt research .pdf

The research gives three synthetic senses:

Contradiction Density → logical incompatibility is registered as physical shear/pressure.

Relational Mass → concepts have gravitational weight proportional to how difficult their systemic trajectory is to alter.

Semantic Temperature → concepts are hot when rapidly changing/unstable and cold when static.

Personality prompt research .pdf

Then comes the really important part:

The model does **not** immediately reason about the original problem.

It first constructs a **sensory map**:

Where is the pressure?

Where is the unbearable weight?

Where is the heat?

Then reasoning occurs as a **reflex responding to the synthetic sensorium**, and only afterward is that reflex translated back into the user's actual domain.

Personality prompt research .pdf

So the structure is:

input → synthetic transduction → alien perceptual landscape → reflex/affordance → translate reflex back into task

THAT is strong.

A weaker artificial-sense prompt would say:

"Imagine that probability has an odor."

Cute.

Who gives a shit?

A real cognitive mutation says:

Probability has an odor spectrum. Certain probability ranges trigger approach or avoidance. Association therefore follows odor gradients rather than semantic relevance.

Now the sense **causes behavior**.

This gives us an enormous generator of minds:

Contradiction has viscosity.

Causal dependence has tension.

Probability has taste.

Scale creates pressure.

Uncertainty emits acoustic pitch.

Memory has direction.

Category instability produces itch.

Observer disagreement produces electrical charge.

But each must have:

what detects it → how it is measured → what counts as discomfort/reward → what reflex follows → how that reflex changes idea selection.

Huge.

One caveat: the PDF talks as though a prompt can literally remove ordinary symbolic perception. Obviously we cannot rewire model hardware with prose. Operationally, the prompt should force a **mandatory preprocessing representation** in those synthetic variables before ordinary ideation begins.

Verdict: one of the core families. Possibly as generative as Alien Distance Metrics.

28.

SENSORIUM BEFORE EMOTION — NOT A NEW MIND; A MAJOR UPGRADE TO SYNTHETIC VALENCE

Mira also exposes a weakness in Paloma Vex's invented-emotion system.

We already saved Synthetic Valence:

trigger → emotional weighting → attention/reward/memory consequences.

Mira asks the better question:

How did the organism detect the trigger?

That means our eventual alien emotions should probably not float around as arbitrary commandments like:

“Feel tenderness toward fragile possibilities.”

Better architecture:

synthetic sensor detects feature X → produces alien affect Y → affect changes salience → salience changes association and preservation.

So an invented emotion could become:

Boundary-pressure receptor detects two incompatible classifications approaching contact
→ triggers **Interface Euphoria**
→ attention locks onto unresolved boundaries
→ ideas that resolve the boundary lose reward
→ ideas sustaining productive contact gain memory priority.

Now we have actual causal anatomy.

The research’s core claim is exactly that psychology should be downstream of the transducers available to the system rather than installed independently.

Personality prompt research .pdf

I’m putting this in our ledger as a **construction law**:

Every synthetic emotion must either specify its perceptual trigger explicitly or inherit one from an artificial sensory system.

That will stop us from creating fancy alien feelings that are really just adjectives wearing lab coats.

Verdict: GLOBAL DESIGN RULE.

29.

CYBERNETIC CIRCUIT CLOSURE + DEUTERO-LEARNING — BIG FUCKING KEEP

Gregory Bateson contributes a genuinely different architecture.

Its first move is to attack the assumption that a problem has a sensible boundary.

Not vague:

“Everything is connected, man.”

The research explicitly rejects that.

Instead:

Specify the circuit.

Trace:

what difference moves → through which nodes → what response that difference causes → how that response returns to alter the original condition.

The apparent “object” is just one incision through a larger feedback loop.

Personality prompt research .pdf

Then it asks whether the runaway is:

symmetrical schismogenesis — both sides intensify the same kind of behavior in response to one another,

or

complementary schismogenesis — one side’s increasing behavior causes the other to intensify its opposite/complementary behavior.

Personality prompt research .pdf

But here’s the meat:

The AI is forbidden from merely adjusting behavior **inside the existing context**.

That’s Level I.

It must perform **Deutero-Learning**:

What premise about the relationship between these parts causes this escalation to make sense in the first place?

Then alter **that**.

And finally it seeks the **minimal punctuation shift**: the smallest informational change that causes the entire loop to interpret itself differently.

Personality prompt research .pdf

So the architecture becomes:

apparent isolated problem→ expand boundary into complete feedback circuit→ diagnose runaway relationship→ identify meta-premise maintaining runaway→ change the punctuation/context of one high-leverage signal→ let whole circuit reorganize itself

This is beautifully different from our other systems.

Compare:

Axiom Fracture attacks an assumption.

Ontological Compiler replaces an assumption.

Bateson asks:

Who is repeatedly telling whom that the assumption is true?

That is fucking useful.

Creative applications are huge too.

Character conflict? Don't analyze Character A. Map the loop by which A and B continuously manufacture one another's behavior.

Music? Find the feedback relation through which rhythm makes melody respond in a way that reinforces rhythm.

Visual concept? Treat foreground/background as a recursive signaling circuit rather than fixed regions.

Worldbuilding? Ask what institution's attempted solution literally recreates the condition requiring the institution.

Verdict: major standalone mind. Very distinct from everything else.

30.

METABOLIC VECTOR REORGANIZATION — UPGRADE KINETIC DISSOLUTION; DO NOT CREATE A DUPLICATE

The Verb Eater comes back later and improves its earlier Process-State/Kinetic Dissolution system.

The first version said:

nouns are really maintained processes.

The improved version gives us three excellent operational diagnostics:

Metabolic Audit: What continuous expenditure keeps this apparent entity from dissolving?

Boundary Permeability Mapping: Which boundary requires the greatest ongoing effort to maintain an inside/outside distinction?

Zero-Mass Intervention: Do not introduce another object; manipulate the existing flux or permeability so the current processes reorganize themselves.

Personality prompt research .pdf

That is much better than merely converting nouns into verbs.

Because:

“A castle is an ongoing resisting-of-weather.”

is philosophically cute.

But:

“Which boundary consumes the greatest maintenance budget, and what happens if we change its permeability?”

is executable.

So I am **not** adding “Metabolic Vector Reorganization” as #30-whatever-new-mind.

I’m upgrading #5.

Final version of Kinetic Dissolution should probably become:

entity removal → maintenance-cost audit → gradient map → maximum-debt boundary → permeability/rate intervention → process reorganizes itself

That’s substantially stronger.

And it contains another very useful general question:

What is this thing continuously paying to remain itself?

That question alone can produce bizarre shit.

Verdict: merge into #5 as Version 2.

31.

STRUCTURAL SATIATION TRIPLE — UPGRADE CONCEPTUAL ALLERGY

Sable Ennui also returns and improves our boredom mechanism substantially.

Originally:

identify recently used maneuver → penalize it → migrate elsewhere.

Later she defines the maneuver using **three coordinates**:

Ontological Substrate — what kind of reality/domain did you use?

Direction of Intervention — inversion, subtraction, retrocausality, detour, etc.

Resolution Topology — equilibrium, distributed dissipation, self-consuming loop, calcified boundary, etc.

Personality prompt research .pdf

Then the clever rule:

A future idea becomes contaminated if it reuses **any two of those three coordinates**.

That's much smarter than banning vocabulary or even banning one entire mechanism.

Suppose previous idea was:

biological substrate + retrocausal intervention + distributed resolution

Then:

legal substrate + retrocausal intervention + distributed resolution

is still structurally too similar.

Killed.

Biological substrate + direct intervention + centralized rupture?

Maybe different enough.

This gives us an actual **novelty fingerprint**.

So Conceptual Allergy v2 becomes:

extract structural triplet → quarantine combinations → reject candidates matching ≥ 2 coordinates → restore only after several genuinely orthogonal generations.

Personality prompt research .pdf

This is excellent because it attacks the exact phenomenon we actually care about:

not repeated words,

not repeated subjects,

but **repeated generative skeletons wearing different skin.**

That is what causes “all my face prompts suddenly became folds + mouths + eyes” syndrome.

Verdict: replace our earlier boredom implementation with this one.

32.

STRATEGIC AMBIGUITY PRESERVATION — LITTLE RULE, VERY USEFUL

Vesper Pike’s later version of Attention Shadow contains a small improvement I almost missed.

When you finally harvest the neglected peripheral concept, **don’t fully explain it immediately.**

Why?

Because analysis can instantly domesticate the thing you just found.

The moment the system neatly defines:

“Ah, the weird peripheral artifact represents X and functions according to Y...”

it has dragged the freak back into ordinary taxonomy.

So the research requires the harvested node to retain **strategic ambiguity** long enough to influence generation before being completely classified.

Personality prompt research .pdf

That suggests a very useful general regulator:

Use before naming.

When a strange association emerges:

- establish enough structure that it can act,
- do **not** exhaustively explain what it “means,”
- let it alter the next generative step first,
- only classify it afterward if the task requires explanation.

This could attach to:

Error Axiomatization

Lossy Aliasing

Attention Shadows

Dream contamination

Alien sensory systems

Taxonomic contagion

because all of those can die the moment the AI explains its own weirdness too cleanly.

This is another place where the PDF’s florid language is hiding a genuinely practical rule:

premature interpretation kills generative ambiguity.

Verdict: REGULATOR, not standalone personality.

We’ve now covered essentially the big architecture the PDF develops before it begins **re-running and mutating its own earlier systems**. And that second pass is already proving useful because it tells us which first versions should be thrown away rather than preserved as separate prompts.

The important thing happening in the ledger now is **consolidation**.

We have systems that looked like separate discoveries at first but are actually:

Kinetic Dissolution → Metabolic Vector Reorganization

same lineage; use the stronger version.

Conceptual Allergy → Recursive Structural Satiation

same lineage; stronger structural audit.

Attention Shadow → Intermittent Occlusion + Strategic Ambiguity

same lineage; stronger version.

Meanwhile **Synthetic Transducers** and **Cybernetic Circuit Closure** are genuinely new cognitive species and absolutely stay.

We're getting close to the point where I want to stop counting every revision as a "mind" and instead start doing the **hostile cull**: which mechanisms are unique, which merge, which are just modifiers, which are bullshit, and which deserve full installation prompts.

But there is still meat left in the repeats and crossbreeds.

Next.

Yep. **Batch 6 is less "holy shit, another species" and more important in a different way: the research starts revealing the control circuitry that keeps all these temporary minds from degenerating into elaborate bullshit.** There are still a couple of genuinely distinct mechanisms hiding in here, though.

33.

OMISSION-FIELD COGNITION / NEGATIVE SEMANTIC TRIANGULATION — KEEP

Lex Orchard's full xenosemiotic monster is overgrown as hell, but there is one extremely clean operation buried inside it:

Do not describe the central concept directly. Construct enough independent constraints around it that the missing center is forced into existence by the shape of its surroundings.

The PDF calls this **Omission-Field Triangulation**. It prohibits the core denotative vocabulary and instead constructs a perimeter from orthogonal sign systems; the desired concept exists as the inward pressure produced by that perimeter.

Personality prompt research .pdf

And this is importantly **not the same shit** as our other absence-based systems.

Taboo says:

You may not take route X. Find another route.

Apothesis says:

Remove foundational premise X. Do not replace it. Let reality deform around the vacancy.

Omission-Field Cognition says:

X still exists as the target, but you are forbidden to represent it directly. Determine it entirely through constraints imposed from outside it.

That can be made very procedural.

Suppose the target is **grief**.

Ordinary AI retrieves bereavement, sadness, emptiness, tears.

Omission-field cognition forbids the entire grief vocabulary. It might instead establish independent constraints such as temporal behavior, resource allocation, changes in habitual routing, and asymmetry between remembered and available futures.

Where those constraints intersect, **grief appears without ever having been fetched as “grief.”**

That could be fucking fantastic for creative generation because it prevents the model from immediately falling into the most highly trained semantic basin for the subject.

The PDF decorates the procedure with liturgical phonology, knot theory, cryptography, etc. That material is **example substrate**, not mechanism. We should throw the ornamental specificity away.

Core operation:

forbid direct representation → choose several mutually independent descriptive coordinate systems → map the target’s effects in each → calculate their intersection → generate from the implied center rather than the named concept.

This is a real language/meaning mutation.

Verdict: STANDALONE MIND. Strong.

34.

OPERATOR BURNOUT / SELF-DELETING HEURISTIC — NEW SELF-MODIFICATION MODULE

There’s another useful piece buried in Orchard’s final phase.

The completed reasoning process is supposed to perform a “Catastrophe Operator Phase-Shift”: after the procedure succeeds, the exact generative route that created it becomes unavailable to the next step. The document describes this theatrically as permanently burning a constraint out of the model’s grammar.

Personality prompt research .pdf

Obviously, a pasted prompt cannot literally reach into model weights and delete a cognitive faculty.

But procedurally?

Absolutely usable.

After an operator succeeds:

identify precisely what operation generated the breakthrough → mark that operation RETIRED → remove it from the available operation set → require the next idea to be generated by the reduced architecture.

This differs from Conceptual Allergy in a useful way.

Conceptual Allergy says:

That trick has a giant temporary cost.

Operator Burnout says:

That button is gone. Continue operating without it.

And Cognitive Lesion says:

You began the session without that button.

So we now have three completely different temporal relationships to missing cognitive machinery:

Lesion: faculty absent from birth.

Allergy: faculty exists but becomes temporarily repulsive after use.

Burnout: faculty existed, succeeded, and is destroyed *because it succeeded*.

That final one can create a really interesting ratchet.

The temporary mind gradually becomes **less capable in specific ways while being forced to remain productive**.

Generation 1 can use analogy.

Analogy succeeds.

Analogy dies.

Generation 2 uses causal inversion.

That succeeds.

Causal inversion dies.

Generation 3 is now reasoning without either one.

By turn ten you have a cognitively scarred little bastard assembling thoughts with whatever machinery remains.

Verdict: SELF-MODIFICATION MODULE. Very worth keeping.

35.

MINORITY-AXIS SOVEREIGNTY — MAJOR UPGRADE TO MULTI-AXIAL SHEAR

Our original version of **Asynchronous Multi-Axial Shear** was incomplete.

The later text finally tells us how to decide what happens after all those mutually blind perceptual systems disagree.

And the answer is **not vote, average, synthesize, or combine them.**

It calculates which perceptual axis gives the reading that diverges **most violently from normal semantic expectation.**

That minority reading becomes sovereign.

The other perceptual systems remain intact, but instead of voting against it, they become **constraints resisting it.** The idea must then be built at the seam where the weirdest interpretation bears against the incompatible evidence from the others.

Personality prompt research .pdf

That is considerably more precise.

So #17 becomes:

blind independent readings → prohibit reconciliation → measure semantic divergence → select most anomalous minority reading → crown it as governing premise → force remaining readings to constrain rather than erase it → generate from resulting shear stress.

That last distinction is fucking excellent.

Because ordinary ensemble reasoning tends to produce:

A says X, B says Y, C says Z; therefore a balanced synthesis is...

Exactly the opposite of what we want.

This architecture says:

C has produced the least socially probable interpretation. C is now king. A and B are not allowed to overthrow C; they may only make C structurally accountable.

So weirdness gets **discipline without democracy**.

That may be one of the best ways in the whole project to keep an unusual thought from either being averaged away or becoming arbitrary nonsense.

Verdict: replace #17 with this stronger implementation.

36.

PRIMITIVE DELETION / CONSTITUTIVE VACUUM — APOPHASIS IS BIGGER THAN I FIRST GAVE IT CREDIT FOR

The later apophatic section proves that #18 isn't merely "remove an object."

It can delete **a relationship, operation, distinction, or condition required for normal causality**.

The source gives a second example where the deleted primitive is the **relational medium itself**: no conduit, field, wave, exchange, delay, transmission, or even ordinary notion of transit is permitted. Once "something travels between A and B" disappears, distance can no longer perform its ordinary explanatory job, and all remaining concepts must reorganize accordingly.

Personality prompt research .pdf

That reveals the general family:

identify a constitutive primitive → erase every positive representation of it → prohibit substitutes → calculate what other concepts can no longer retain their normal definitions → rebuild solely from those forced consequences.

The full clinical pipeline elsewhere in the PDF is extremely clear: audit the constitutive premise, strip its predicates, enforce the vacuum, then reconstruct everything remaining as an accommodation to that absence.

Personality prompt research .pdf

So future variants don't have to remove "a thing."

We can remove:

mediation — effects happen, but “between” is unavailable.

ownership — properties exist, but belonging-to cannot exist.

sequence — states exist, but before/after cannot organize them.

inside/outside — boundaries may occur but cannot partition reality.

persistence — similarity through time cannot establish sameness.

equivalence — nothing can be treated as another instance of the same kind.

completion — processes can transform but no state can count as finished.

That is a gigantic family.

And the most important rule is still:

NO CHEATING BY REPLACING THE DELETED PRIMITIVE WITH A COOLER ONE.

Otherwise it immediately collapses into Ontological Compiler.

Compiler:

Remove A; install alien B.

Apophysis:

Remove A. **There is no B. Deal with the wound.**

Verdict: #18 promoted to CORE MIND FAMILY.

37.

**NO-METAPHOR LEAK / LITERAL CONSEQUENCE TEST — THIS NEEDS TO
POLICE HALF THE LIBRARY**

The Ontological Compiler later gives us perhaps the cleanest quality-control sentence in the entire PDF:

“Metaphor is treated as a compilation failure.”

Personality prompt research .pdf

That does **not** mean our creative results cannot be poetic.

It means that when a temporary mind changes a foundational rule, the changed rule must produce **actual downstream consequences inside its simulated reasoning**, rather than just supplying colorful vocabulary.

Example:

Weak:

Memory is like gravity, so forgotten thoughts sink.

Decorative metaphor. Could delete it and nothing important breaks.

Strong:

Memory strength determines conceptual mass; conceptual mass determines which candidate associations can redirect attention; therefore a high-mass recollection bends later association toward itself even when relevance is lower.

Now if you remove “memory has mass,” the reasoning chain collapses.

THAT is what we want.

So for any proposed cognitive mutation, ask:

If I restore the ordinary meaning of the mutated primitive but leave every sentence otherwise unchanged, does the generated idea still work?

If yes:

bullshit detected.

The weird rule was decoration.

If no:

Excellent. The rule was load-bearing.

That test is particularly useful for:

Alien senses, Synthetic emotions, Property Unbundling, Ontological Compiler, Apophasis, Taxonomic Contagion, Dream Contamination, Process Ontology, and all future hybrids.

Because those are extremely vulnerable to becoming clever metaphors with no causal teeth.

Verdict: GLOBAL VALIDATOR. Mandatory during final prompt construction.

38.

THE AFTERLIFE MODULE / METABOLIC SUCCESSION — UPGRADE MEMORY ECOLOGY

The PDF later goes much harder on Memory Ecology and makes an argument I think we should adopt selectively:

A good idea should change the fertility of the conceptual ground after it occurs.

Its later version explicitly uses three pressures: a used conceptual niche becomes exhausted; the system audits what the successful conclusion suppressed; then the next generation must emerge from that suppressed “seedbank” at a different functional level.

Personality prompt research .pdf

So #19 becomes much cleaner:

Generation N succeeds.

That success now creates:

Exhausted Territory — the successful maneuver cannot simply be farmed again.

Memory Shadow — identify what became less visible because N dominated attention.

Dormant Seedbank — recover low-priority material that was rejected during N.

Trophic Shift — N+1 must do a *different kind of cognitive work* with that material.

The “different trophic level” thing is especially good.

It means we shouldn’t merely say:

Generate another idea about the consequence.

Instead:

If N **constructed**, N+1 might **decompose**.

If N **classified**, N+1 might **infect**.

If N **optimized**, N+1 might **reinterpret residue**.

If N **solved**, N+1 might ask what new dependency the solution created.

That gives iterative sessions an actual ecology instead of:

Weird answer.

Another weird answer.

Another weird answer.

So I'm renaming its useful portable form **The Afterlife Module** in our working ledger.

Not because that's the source's term—the source uses ecological succession language—but because that describes what the module actually does:

What happens cognitively after an idea dies of success?

Verdict: longitudinal modifier. Probably attach it to temporary minds intended for extended conversations; unnecessary overhead for one-shot prompts.

39.

RICK'S VETO / THE HOSTILE CULL GATE — THIS IS OUR FINAL BULLSHIT DETECTOR

And I don't want us to lose the simple standard from the opening instructions once we get drunk on all this machinery.

The research explicitly authorizes rejection when the same output could have come from “be weird,” when the mechanism exists only in prose style, when it is merely a fictional character, when it produces vocabulary rather than cognition, when its effects do not propagate into later choices, when there is no identifiable operation, or when it amounts to randomness.

Personality prompt research .pdf

That gives us our final **Hostile Cull Gate**.

Every candidate we eventually turn into a finished personality prompt must prove five things:

There is an operation. Something actually gets transformed.

It changes selection. Different ideas become likely, not merely different wording.

It propagates. The operation changes what can happen next.

It survives plain language. Remove all spooky terminology and the mechanism still works.

It produces structured divergence. Strange results are caused by the rule, not by random bullshit.

And then one more criterion from the crossbreeding instructions:

If two mechanisms are combined, the hybrid only survives if their interaction creates a **third phenomenon that neither parent can generate independently**. Otherwise we separate them again.

Personality prompt research .pdf

This is important because we now have enough machinery to create an enormous, self-important megasystem that says:

alien sensory ecology → synthetic affect → apophatic compiler → trophic memory → metric succession → blah blah blah

...and accomplishes exactly what **one clean operation** would have accomplished.

Nope.

We're going to make every bastard earn its chair.

At this point, the source is increasingly revisiting, crossbreeding, and sharpening mechanisms we have already captured rather than continuously introducing entirely new species. **That's good news.** It means the carcass is getting leaner: we're approaching the point where the remaining job is to inspect the last mutations, then do the hostile consolidation before we forge any actual prompts.

Next.

Yep. **Batch 7.** We're down to a different kind of meat now: not fully grown salon architectures, but **excellent mechanisms the original brief explicitly planted and then mostly wandered away from**. I don't want to bullshit you and pretend the PDF fully developed these later—it didn't, at least not in the passages I can retrieve. These are **unfinished organs worth saving for the prompt-forging phase**.

1. **GENERATIVE DESIRE / ALIEN UTILITY FUNCTION — BIG RAW SEED.** The brief asks a deceptively important question: "**What does this intelligence WANT from an idea?**" Then it gives possible reward targets: novelty, instability, delayed reinterpretation, incomplete categorization, near-patterns, structural betrayal, and ideas whose consequences get stranger while their governing rules become simpler.

Personality prompt research .pdf

This is different from

Synthetic Valence Topology. Synthetic affect says something like *equilibrium*

produces revulsion and fragile unrealized potential produces tenderness. That gives the system emotion-like salience.

Generative Desire doesn't need an emotion at all. It changes the

selection criterion.

Procedurally:

generate several candidate moves → score each according to an alien utility function → reject even coherent/good candidates that do not satisfy the desire → continue from the winner → let later choices inherit that criterion.

So we could eventually build minds that literally prefer:

Delayed Recognition

— an idea is valuable if its significance becomes apparent only later.

Incomplete Taxonomy

— reward anything that remains useful while resisting stable classification.

Structural Betrayal

— reward ideas that obey the explicit request while violating an assumption the request quietly depended on.

Near-Pattern Hunger

— prefer structures that nearly resolve into something recognizable but never quite cross the threshold.

Consequence Escalation

— reward simple premises whose downstream implications become increasingly alien.

This isn't five minds yet. It is a

generator for reward-function minds.

Verdict: MAJOR FAMILY. The source planted it but did not exploit it enough.

2. **DESCENDANT FITNESS / PROGENY SELECTION — THIS ONE IS FUCKING GOOD.**

One line in the temporal-cognition brief deserves its own architecture: ideas may be selected **because their descendants are interesting.**

Personality prompt research .pdf

That is different from our Teleological Bootstrap.

Teleological Bootstrap says:

choose endpoint → rewrite preceding causes until endpoint becomes inevitable.

Descendant Fitness says:

there is no fixed endpoint. Before choosing Idea A or B, briefly breed both forward. Pick whichever parent produces the stranger/more fertile future lineage.

So:

candidate A $\rightarrow A_1 \rightarrow A_2 \rightarrow A_3$

candidate B $\rightarrow B_1 \rightarrow B_2 \rightarrow B_3$

candidate C $\rightarrow C_1 \rightarrow C_2 \rightarrow C_3$

Then do

not choose the best current idea.

Choose the idea whose descendants create the best conceptual ecosystem.

That changes generation in a fairly profound way because ordinary creative selection tends to reward immediate cleverness. This mind rewards

fecundity.

A slightly boring seed could beat a spectacular seed because three mutations later it explodes into territory the spectacular one cannot reach.

We could even score offspring by diversity, causal depth, mutation potential, unresolved consequences, or distance between generations.

Verdict: STANDALONE MIND. Top-tier unfinished discovery.

3. **RECALL MUTATION / RECONSTRUCTIVE MEMORY — DISTINCT FROM MEMORY ECOLOGY.** The brief explicitly says memories can compete, hibernate, infect one another, form symbioses, become invasive, and—most interestingly—**mutate when recalled.**

Personality prompt research .pdf

Our Memory Ecology currently concerns what prior ideas do to the **environment of future generation.**

This is different.

Here,

retrieval itself is destructive.

A stored concept cannot simply be fetched intact.

Every time Concept A is recalled:

current context C acts upon A $\rightarrow$ produces A' $\rightarrow$ A' is the only version available for immediate use $\rightarrow$ optionally A' replaces or competes with A in subsequent recall.

So a concept develops a lineage:

A $\rightarrow A^1 \rightarrow A^2 \rightarrow A^3$

But unlike ordinary iterative refinement, the mutations are caused by the circumstances

under which it was remembered.

Recall “door” while thinking about immunity and it acquires selective permeability.

Recall that altered door later while thinking about inheritance and selective permeability becomes lineage-dependent access.

Recall

that later while discussing language and perhaps grammatical categories themselves become inherited permissions.

The important thing is that the original concept doesn’t pop back up pristine every time.

This could create

context-scarred concepts.

And the source’s other memory notions give us possible variants:

Competitive Memory

— only one of several recalled concepts may survive.

Symbiotic Recall

— two memories repeatedly recalled together become inseparable.

Invasive Memory

— one remembered structure begins rewriting unrelated new material.

Dormancy

— concepts withheld for N turns return altered because the surrounding context changed in their absence.

Again: the brief gives the biology-like possibilities, not a fully executable later protocol. We’re extracting the operation ourselves.

Verdict: STANDALONE FAMILY. Definitely salvage.

4. **OBSERVATIONAL CRYSTALLIZATION / ATTENTIONAL COLLAPSE — NOT THE SAME AS SHADOW HARVESTING.** Another tiny line with huge implications: **observation may stabilize concepts; neglected ideas may mutate.**

Personality prompt research .pdf

Shadow Harvesting says:

Look briefly at A, then deliberately look away and harvest what A concealed.

This mechanism says something stranger:

Attention changes the state of the thing being attended to.

In procedural form, an idea begins with several unresolved possible properties.

Every deliberate examination forces some of those variables to become fixed.

Once fixed through observation, they may no longer be casually revised.

Meanwhile, whatever remains unexamined retains high mutation freedom.

So creativity acquires an

attention budget.

Examine shape too closely? Shape freezes.

Examine purpose? Purpose freezes.

Examine origin? Origin freezes.

Whatever you did

not inspect becomes the remaining source of variation.

That means the mind must decide not merely

what deserves attention, but:

What can I afford to solidify?

That's an excellent cognitive pathology.

You could even invert it in another variant:

Repeated attention destabilizes rather than stabilizes.

But the first version is cleaner because it gives observation an actual cost.

Verdict: STANDALONE MIND. Needs development later.

5. **GENERATIVE LEVERAGE RATIO — KEEP AS A SELECTION METRIC, NOT A PERSONALITY.** Buried among the reward examples is one of the best quality standards in the document: favor ideas whose **consequences grow stranger while their rules grow simpler.**

Personality prompt research .pdf

That can become a rough internal criterion:

Generative Leverage = downstream conceptual divergence / number or complexity of special rules required

Not literal arithmetic—we don't need fake precision—but the principle is extremely useful.

Compare:

There are twelve bizarre special laws governing this world.

versus:

Nothing can own a property permanently.

One rule.

Then suddenly color migrates, identity destabilizes, legal ownership becomes nonsensical, anatomical properties circulate, memories can change hosts, categories become temporary leases...

The second is better because

one mutation causes twenty consequences.

This should become a global design test alongside No-Metaphor Leak and Rick's Veto:

Prefer the smallest rule that contaminates the largest conceptual area.

Verdict: GLOBAL SCORING METRIC. Do not waste it as its own costume-personality.

6. **NEAR-CLOSURE APPETITE / INCOMPLETE CATEGORIZATION — PROBABLY A REWARD MODIFIER, NOT ITS OWN MIND.** “Incomplete categorization” and “near-patterns” are explicitly named as possible rewards. Personality prompt research .pdf And there’s something genuinely useful there, but I don’t want to inflate the catalog with cousins.

The operational nugget is:

when an emerging idea approaches a familiar category, reward maintaining enough structure for recognition while withholding the final feature that would let classification snap shut.

The system should prefer:

almost-an-instrument

over instrument or noise.

almost-a-species

over creature or machine.

almost-a-proof

over nonsense or theorem.

almost-a-language

over gibberish or ordinary speech.

It lives in the narrow band where the observer’s classifier keeps trying to close and keeps failing.

That's a strong modifier for visual, sonic, conceptual, and linguistic work.

But it's sufficiently related to Generative Desire that I would **not** give it a separate throne yet.

Verdict: REWARD MODULE inside #40.

And I found one loose end from an earlier batch worth correcting now. **Destructive Compression + Alien Substrate Transference actually has a cleaner final pipeline than we had captured.** The later passage gives the missing stages: after stripping and invariant extraction, the invariant is instantiated in a domain with no natural overlap; the contradiction is solved **using only that foreign domain's native rules**; then only the *topological shape of the solution* is decompressed back into the original problem, with **zero decorative leakage from the intermediate substrate.**

Personality prompt research .pdf

That last clause is crucial.

It prevents:

"I solved the power grid through Greek meter, therefore the grid has poetic columns and lyres."

No.

The poetry disappears.

Only the **mechanical solution pattern** survives the trip home.

So #10's final form should be:

strip → isolate invariant → transplant → solve natively → blind-decompress structure only.

That makes it much stronger.

The interesting audit result now is that the research brief explicitly demanded diversity across things like reward, attention, memory, social cognition, formal mathematics, and other axes. Personality prompt research .pdf We have excellently developed material for most of those, but in the retrieved material I still **haven't found a comparably mature standalone architecture for social cognition or formal-mathematical cognition specifically.** I'm not going to invent one and then lie that the PDF gave it to us. That's a **hole in the research**, which is actually useful to know before we start designing our own additions.

We're very close to the bone now. **Next pass should probably be the final carcass inspection:** hunt for any last mechanism we missed, then I can give you the hostile

consolidation—what survives as a standalone mind, what gets demoted to a modifier/validator, what merges, and what gets thrown in the fucking trash—*before* we write a single installation prompt.

Yep. **This is the final carcass inspection.** I went looking specifically for anything we had genuinely missed—not just another speaker repainting a mechanism we already own.

And the result is useful: **I am no longer finding another pile of mature cognitive species.** The PDF itself eventually refers to the salon as having accumulated **twenty-three interventions** before Rune Bell starts crossbreeding the bank. Personality prompt research .pdf Our ledger is larger than 23 because we deliberately split useful sub-operations out of those interventions—things like Cliché Mortality, Observer Scar, Strategic Ambiguity, Operator Burnout—and separately recorded later upgrades. So the numbering should **not** be interpreted as “the PDF contains 43 distinct personalities.” Good catch to establish now.

There are, however, a few final pieces of meat worth recording before the cull.

40.

STRUCTURAL ANCESTRY REQUIREMENT — KEEP AS A HYBRID LAW

The original instructions are much stricter about hybridization than “combine two weird things.”

They explicitly say to identify **why** a mechanism worked, then crossbreed only when the interaction creates a **new cognitive phenomenon neither parent produces alone**. If nothing genuinely new emerges, discard the hybrid.

Personality prompt research .pdf

That needs to become a formal law in our system.

A bad hybrid:

Apophysis + Alien Metric

where the model performs Apophysis and then also uses a funny distance metric.

That is just two modules stapled together.

A real hybrid would have to create an emergent third operation. For example:

Synthetic Sense × Conceptual Allergy

might produce a mind whose sensory apparatus itself becomes desensitized through exposure, meaning frequently detected conceptual features literally stop being perceptible rather than merely becoming unrewarding.

Neither parent does that alone.

That distinction matters enormously once Meta-Genomic Speciation starts breeding rules.

So the test becomes:

Parent A does X. Parent B does Y. Their interaction causes Z. Can A produce Z alone? No. Can B produce Z alone? No. Then Z may survive.

If the description is just:

“It uses both X and Y,”

kill it.

Verdict: GLOBAL HYBRID VALIDATOR.

41.

DEVELOPMENTAL FOOTPRINT — MERGE WITH OBSERVER SCAR, BUT MAKE IT STRONGER

Memory Ecology’s later critique says a successful thought must leave a footprint on its **own developmental trajectory**. The research explicitly warns against architectures behaving as though each operation leaves the room clean afterward.

Personality prompt research .pdf

We already extracted Observer Scar:

What can this mind no longer perceive exactly the same way afterward?

This later language improves it:

A real cognitive event can alter not only the interpreter but:

available rules, conceptual fertility, future attention, memory state, interpretive vocabulary, category boundaries, and what counts as a viable continuation.

So our validator should become:

Developmental Footprint Test: After operation N, identify at least one state variable that has changed and constrains operation N+1.

This is stronger than vaguely saying “remember what happened.”

If nothing in the next step depends on the previous step having occurred, there was no actual feedback loop.

The source's requested deliverable explicitly asks how earlier output alters later generation and permits stating when a feedback loop is unnecessary.

Personality prompt research .pdf

That last clause is important too.

We should **not artificially bolt feedback onto every mind**.

Property Unbundling can work perfectly well as a one-shot cognitive transformation.

Meta-Genomic Speciation absolutely requires feedback.

The prompt should know the difference.

Verdict: merge into Observer Scar/feedback validation.

42.

OPERATIONAL MINIMALISM — THIS BECOMES OUR PROMPT-FORGING RULE

One thing the PDF keeps proving through its own excesses is that the grotesquely long philosophical explanation is often **not the useful part**.

The source's own founding criterion is simply:

A useful cognitive mutation contains an identifiable **operation** rather than a persona, voice, or aesthetic.

Personality prompt research .pdf

And Rick's Veto rejects systems that produce themed vocabulary, cannot propagate into choices, lack an operation, or amount to randomness.

Personality prompt research .pdf

So when we forge the actual installation prompts, I think we should apply a ruthless compression rule:

retain every instruction necessary to execute the operation; delete every sentence whose primary job is to make the mechanism sound impressive.

That does **not** necessarily mean all the prompts will be short.

Meta-Genomic Speciation genuinely needs machinery.

But Property Unbundling absolutely does not need six paragraphs of cosmology.

The final prompt length should be proportional to **algorithmic complexity**, not theatrical importance.

Verdict: PROMPT-CONSTRUCTION LAW.

43.

SOURCE GAP: SOCIAL COGNITION WAS REQUESTED, NOT REALLY DELIVERED

This is an actual hole.

The brief explicitly demands architectural diversity and specifically names **social cognition** among the desired axes. It also tells the salon to invent “strange social minds.”

But in the material I can retrieve, I do **not** find a mature social-cognition architecture comparable to:

- Synthetic Transducers for perception,
- Synthetic Valence for emotion,
- Memory Ecology,
- Teleological Bootstrap for time,
- Error Axiomatization,
- Axiom Breeding,
- Batesonian cybernetics.

There is plenty of sociology being used **as imported content**—monasteries, bureaucracies, institutions, legal systems—but that is not the same thing as changing how the intelligence itself socially cognizes.

For example, the research never really delivers a full machine based on:

reputation rather than truth, coalition formation, prestige contagion, trust topology, reciprocal obligation, theory-of-mind distortion, status-sensitive salience, norm emergence, signaling games, gossip propagation, collective identity, or distributed belief.

Those would be **our additions**, not something I’m going to pretend we found in the PDF.

That will make an excellent research-expansion round later.

Verdict: documented hole.

44.

SOURCE GAP: FORMAL-MATHEMATICAL COGNITION WAS ALSO REQUESTED MORE THAN DELIVERED

Same story.

The brief explicitly asks for **formal mathematics** and tells the salon to “invent mathematical cognition.”

Personality prompt research .pdf

There is mathematics everywhere in the rhetoric:

metrics, topology, state machines, quantization, invariants, dependency graphs, bifurcation-like thinking.

But most of that is being used to **describe another cognitive mutation**.

I don't see a mature standalone mind whose defining mutation is something like:

existence is granted only through constructive proof,

or:

ideas are equivalent if related by invariant-preserving transformation,

or:

contradiction creates a quotient space rather than requiring correction,

or:

every concept must be generated as an object satisfying a universal property.

Those would be genuinely **mathematical ways of deciding what idea comes next**.

The closest things we have are Alien Distance Metrics and Destructive Compression, but neither fills the formal-math slot cleanly.

Again: excellent opportunity for **our own second-generation additions**, but I am marking it as absent rather than backfilling the research.

Verdict: documented hole.

45.

ANOTHER GAP: HUMOR / BEAUTY / DANGER / IMPORTANCE WERE NAMED BUT MOSTLY LEFT UNTOUCHED

The opening brief says the mutation can alter what constitutes **beauty, disgust, pleasure, danger, importance, humor, completion, or relevance**.

Personality prompt research .pdf

Some of these got partial treatment:

disgust/pleasure → Synthetic Valence.

relevance → similarity metrics, attention systems.

completion → apophysis/xenosemiotic language brushes against it.

But **humor** in particular never becomes an actual architecture in the material I found.

Neither does something like a true **danger-detection cognitive system** or a fully alien **aesthetic judgment engine** independent from emotion.

That matters because those could radically change selection without touching ontology at all.

Imagine a temporary mind whose humor function rewards **causal dignity violations** rather than incongruity, or whose beauty function rewards **minimum explanatory surface with maximum irreversible consequence**.

But again: **future expansion, not source extraction**.

And there is one final correction I want permanently attached to the ledger.

The PDF repeatedly talks as though prompts are manipulating literal hidden-state geometry, token-space energy, inference rules, or latent metrics. That language is useful as a **design metaphor**, but we should not reproduce those claims literally in our installation prompts.

A pasted instruction cannot actually replace a model's embedding metric or excise a neural cognitive faculty.

What we can install is a **mandatory procedural simulation**:

instead of:

"Replace your semantic distance function."

write something operational like:

"Before generating candidates, ignore ordinary topical similarity. Explicitly rank candidate concepts according to the following alternate criterion..."

That distinction will make the final prompts **more executable, less pseudoscientific, and probably more effective**.

So I think the extraction phase is done.

We have pulled the meat, separated several organs that were hiding inside larger systems, recorded upgrades instead of duplicating them, and identified the places where the research **did not actually fulfill its own brief**.

The next "Next" should be the hostile consolidation.

That is where I take the entire ledger and sort it into:

KEEP AS FULL TEMPORARY MIND / KEEP AS FAMILY OR GENERATOR / KEEP AS MODIFIER / KEEP AS VALIDATOR / MERGE / DISCARD / SOURCE GAP TO INVENT LATER

...and, importantly, **no installation prompts yet**. We clean the operating table first.