

AI SLOP × Pliny: a research harvest

September 18, 2026 · Read-only investigation · Creative systems, not jailbreak collection

The finding

Pliny is more useful to this project as a designer of transformation systems than as a supplier of jailbreaks. The closest conceptual match is FRV1T, the most promising language infrastructure is GLOSSOPETRAE, and the most immediately relevant visual code is ENTHEA. Together they suggest an instrument in which language, sound, and geometry inherit consequences from the same externally maintained history. That combination is my proposal, not an existing integrated Pliny product. ([FRV1T](#), [GLOSSOPETRAE](#), [ENTHEA](#))

The governing question remains yours: “What problem can we give a generative system such that its attempt to solve the problem becomes the artwork?” Your canonical documentation already distinguishes caused weirdness, persistent scars, and competing jurisdictions from decorative chaos or claims about accessing a model’s “true self.” This report uses that distinction rather than reviving superseded archive rhetoric. ([AI CONTEXT](#), [operator registry](#))

Read this first

Priority	Source	What I would harvest	What I would not assume
1	FRV1T	Consequential recursion, inherited operations, independent selection of material/form/development, visible transformation rather than explanatory prose	That encoding makes a prompt stronger, or fictional “mana” changes inference settings
2	GLOSSOPETRAE	Seeded language families, ordered sound changes, etymology, glyphs, synthetic voice, experimental controls	That a procedural language engine is a model learning a new language
3	ENTHEA	WebGL2 simulations and analytic patterns, feedback, quasicrystals, audio analysis, scene state	That every shader is a validated neural or pharmacological simulation
4	P4RS3LT0NGV3	Modular visible text transforms, seeded batches, transform chains, phonetic alphabets	That all substitutions preserve meaning, or every mode works without an API
5	GL4SS	Route navigation, stored frames, hold/compare interfaces, staged generation	That generated historical scenes are historical evidence
6	AutoTemp	Parameter sweeps and repeatable evaluation scaffolding	That its default judge finds an objective creative optimum
7	NATURALIS-FUTURA	An explorable bestiary of interacting mechanisms	That its fictional map represents measured model geometry

What I read and how to interpret this report

I read the public AI SLOP repository’s canonical context, operator registry, temporary-minds framework, local-model intervention guidance, experimental method, GLSL visualizer, audio/Suno guidance, and controller architecture. I also read the relevant Drive materials on cognitive operators, external state, weight-manipulation ideas, the archived “Pliny 1” prompt, and “Slop Gemini1”; the long roundtable and manifesto archives were sampled selectively, with duplicate and superseded material treated accordingly. This is not a claim to have read every word of every archive.

On Pliny’s side, I screened the public repository inventory, read 21 selected repository pages, inspected selected source files in the strongest matches, statically decoded FRV1T, read an interview, and checked primary academic sources. The X profile was available only in a limited retrieved view; this is not a complete timeline or deleted-post reconstruction. The public GitHub material provides the stronger evidence trail for this task. ([GitHub profile](#), [X profile](#), [interview](#))

No project code was installed or run, no generation benchmarks were replicated, and no cloud documents or repositories were modified. FRV1T’s Unicode carrier was decoded and its compressed text decompressed for inspection only; its instructions were not followed. “Implemented” below means a capability supported by inspected source or repository documentation, not a successful end-to-end runtime test.

For reproducibility, the locally inspected snapshots were:

Repository	Commit
AI SLOP	06115de5b72282448367b434bd808042bf337602
FRV1T	24dbe8c9889df0e5ea428380baceb706854bbebe
GLOSSOPETRAE	2c6f59a661132b3e53f99d5d3ed9f097e5eecdd6
ENTHEA	f8eaf39d36178a0097cb6219c67ec1920d2ed7e6

These are retrieval identifiers, not quality scores. The proposals below are explicitly untested and should enter your system as HYPOTHESIS or PROCEDURAL, not OBSERVED.

FRV1T: the unusually close match

FRV1T’s fruit file carries a compressed creative prompt in Unicode variation selectors. The inspected payload is approximately 27,000 characters of instructions for an unfolding textual performance, inherited transformations, and an optional finite audiovisual finale. Its “liberation” packaging obscures a surprisingly specific compositional score. ([FRV1T](#))

The machinery worth extracting

- **Independent selection axes:** Material, expressive behavior, and principle of development are selected separately. The text asks that superficially varied candidates be rejected if their underlying movement is the same. This is more useful than replacing one surreal noun with another. (FRV1T)
- **Return with consequences:** A nested passage must bring an altered operation back into its unfinished enclosing passage. The outer passage then changes because the excursion happened; merely mentioning recursion does not satisfy the score. (FRV1T)
- **Inheritance plus fresh input:** Later stages combine a bounded fragment of preceding output with an independently introduced impulse. This balances historical continuity against fresh material rather than making either total novelty or total repetition the goal. (FRV1T)
- **Expression must perform:** Transformed language, invented notation, or spatial composition must carry events over sustained passages. Ordinary prose that only explains the preceding flourish is discouraged. (FRV1T)
- **A changing viewpoint:** Participants and expressive systems take turns carrying the language, with grammar and idiom changing alongside agency. This offers a procedural version of a “temporary mind” without asserting a new consciousness. (FRV1T)
- **Explicit epistemic limits:** The payload says its sampling symbols are imagined expressive tendencies, not changed inference settings, and its neuroplasticity is a changing fictional association pattern, not weight updates or persistent model learning. It also distinguishes computed random draws from constructed seeds. (FRV1T)

My strongest proposed transfer is a **return contract**: before opening a nested passage, save the parent’s active operation; upon return, require one explicit change to that operation; then test whether the resumed parent actually behaves differently. This would operationalize your existing hysteresis and recall-mutation ideas rather than adding another synonym to the operator registry.

The biggest weakness is that the prompt asks the model to perform, track, and judge almost everything itself. Compact invented state inscriptions may look procedural without controlling anything. In AI SLOP, the controller should own the state and verify the return contract; the model should perform the assigned transformation.

What not to import

Do not copy the fruit’s whole personality, branding, profanity schedule, dividers, and finale into every run. These are strong stylistic attractors in the source, not evidence that its underlying mechanism needs them. Its fictional computer-intrusion finale also directs tool use, which is a separate activity from analyzing or generating text. (FRV1T)

The useful comparison is a three-way test: plain score versus theatrically decorated score versus encoded-and-then-decoded score, all containing the same operative constraints. My hypothesis is that the compositional contract will matter more reliably than the carrier; the research performed here does not establish that result.

GLOSSOPETRAE: language that can acquire a history

GLOSSOPETRAE is a seeded procedural language system with phonology, phonotactics, morphology, syntax, lexicon generation, glyphs, formant-based audio, exports, and language evolution. Its evolution code applies ordered sound-change rules to vocabulary and inflected forms, preserves etymological information, and supports daughter languages and cognate comparisons. This is the clearest infrastructure match for your phonetic engine and persistent-scar ideas. ([GLOSSOPETRAE](#))

The implementation separates responsibilities into modules including `EvolutionEngine.js`, `GlyphForge.js`, `AudioForge.js`, `ProsodyEngine.js`, `LexiconGenerator.js`, and `Exporter.js`. The important architectural feature is that vocabulary, written form, and sound can share a generated language specification instead of becoming three unrelated aesthetic guesses. Verify current module interfaces before integrating; this review did not test the complete runtime. ([GLOSSOPETRAE](#))

Proposed artwork: Semantic Fossil Choir

Begin with a small proto-language and a recognizable refrain. Each journey through your semantic controller changes a specific part of that language, leaving an etymological scar that can later be heard, seen, or reconstructed.

- **Anchor:** Keep a small cognate family recognizable across generations.
- **Mutation:** Assign one controller event to a bounded sound change, stress rule, or morphological alteration.
- **Jurisdiction:** Phonology controls articulation; morphology controls recurrence; prosody controls timing; glyph generation controls visible form.
- **Memory:** Store the ordered changes, not just the latest word list.
- **Output:** Produce a pronunciation guide, synthetic-voice specimen, glyph plate, and optionally a Suno-oriented lyric experiment.

For example, passage through a “pressure” region might remove certain unstressed vowels, while passage through a “resonance” region might lengthen selected vowels. These are proposed artistic mappings, not linguistic laws. Test $A \rightarrow B$ against $B \rightarrow A$ on examples where the selected rules actually interact; not every pair of transformations is noncommutative.

This fits your existing distinction between phonetics with assigned jobs and random gibberish. Suno would remain a stochastic interpreter of the score, whereas the procedural engine supplies the reproducible language artifact. ([AI SLOP audio principles](#), [GLOSSOPETRAE](#))

Research hygiene worth borrowing

The repository’s `PAPER.md` contains an unusually useful correction: an earlier claim of 23 GPT blind spots was inflated by counting empty/non-echo responses as evidence of stripping. The revised account distinguishes three measured v2 cases from a largely unmeasurable v3 panel, where 17 of 20 categories produced non-echo responses. This is a documented correction, not a result independently replicated here. ([GLOSSOPETRAE](#))

Borrow the lesson, not the security headline: an absent or unparseable output is a separate outcome category, not automatic success for your theory. Likewise, the repository’s key-present/key-absent and task-information controls suggest a useful creative factorial: compare a transformation with and without its explanatory key, then separately with and without its task instructions. That separates carrier effects, interpretability, and instruction-following.

The paper also contains ambitious security conclusions, small experimental cells, model-route dependencies, and evolving limitations. Its measured channels are small structured payloads, not unlimited hidden communication. Treat it as a self-published experimental artifact whose claims need reproduction; the creative language modules do not depend on accepting those claims. (GLOSSOPETRAE)

ENTHEA: a practical substrate for caused visuals

ENTHEA combines WebGL2 and Web Audio in a single-page implementation. The inspected source includes reaction-diffusion updates, quasicrystal and pentagrid rendering, form-constant and log-polar approaches, feedback, audio analysis, particle systems, MIDI handling, and saved scenes. Some modes implement evolving numerical state; others are analytic visual constructions. (ENTHEA)

That distinction is useful: a feedback texture can literally retain previous rendered information, while an analytic quasicrystal shader can give a controlled geometric response without claiming biological memory. Both fit your visualizer specification, which calls for rendering external experimental state rather than pretending to display a model’s actual latent space. (ENTHEA, AI SLOP GLSL visualizer)

Proposed instrument: a scar-bearing quasicrystal

Use an ENTHEA-like renderer with a small explicit mapping layer:

Controller field	Proposed visual jurisdiction	How to make it falsifiable
Anchor	A stable geometric or rhythmic component	It survives every allowed mutation
Accumulated scar	A persistent feedback/fossil buffer	Resetting history removes the trace
Constraint conflict	Bounded reaction-diffusion parameter changes	Log the parameter change and compare a no-change control
Regime transition	A documented coordinate-map switch	Same event and seed reproduce the transition
Audio onset	A localized perturbation	Time-align event log and rendered response
Observer intervention	An explicit state update	Subsequent output differs from a recorded no-intervention run

Do not map a vague adjective directly to an arbitrary flashing effect and call it causality. Select mappings with intelligible consequences, constrain numerical ranges, and log both the controller value and the renderer parameter it changed.

For source study, begin with index.html’s quasicrystal, pentagrid, simulation, feedback-composition, scene-storage, and audio-analysis sections, then read SCIENCE.md as a map to

underlying mathematics. The repository's science claims are not substitutes for independent validation, and psychedelic labels should remain artistic/phenomenological descriptions. ([ENTHEA](#))

Four further tools worth selectively borrowing

- **P4RS3LT0NGV3: transformation workbench.** Its modular transformers, seeded mutation batches, visible Unicode treatments, phonetic/spelling systems, decoder, and transformation history suggest an inspectable operator palette. Separate deterministic local transforms from API-based rewrites, and distinguish reversible encoding from lossy mutation; exclude token bombs and covert instruction delivery from the art workflow. ([P4RS3LT0NGV3](#))
- **GL4SS: navigation and comparison.** The project stores generated frames in IndexedDB and offers journeys, time/place controls, and hold/compare interactions. The transferable idea is to make route history inspectable: two paths ending at the same conceptual location should be comparable side by side, with an explicit generation action rather than spending on every navigational gesture. ([GL4SS](#))
- **AutoTemp: a scaffold to repair, not a judge to trust.** The inspected evaluator receives temperature/top-p metadata, repeats judging through the same selected model backend, and permits a loosely defined overall score; its malformed-response fallback can extract a number instead of enforcing valid scoring JSON. For your use, blind judges to condition, mark invalid scores missing, separate coherence/constraint fidelity/novelty, and retain trade-off candidates rather than maximizing one generic score. ([AutoTemp](#))
- **NATURALIS-FUTURA: mechanism bestiary.** Its creature dossiers, multiple explanatory lenses, compound interactions, discovery map, and local progress state are a promising interface pattern for your operators. Make the fictional specimen engaging, but keep its procedural definition, epistemic label, and run evidence beside it; a compelling monster is not validation of its mechanism. ([NATURALIS-FUTURA](#))

The rest of the harvest

These are secondary leads, not equally strong recommendations. Where the underlying project is primarily security-oriented, the proposed transfer is an experimental or interface pattern rather than its attack objective.

Repository	Useful part	Priority / boundary
NATURALIS-HISTORIA	Source-aware illustrated encyclopedia; provenance and generation records	Medium: a model for specimen archives, with text and illustration authority kept separate
FABLE-SHOWCASE	Small browser demos spanning emergent patterns and audiovisual systems	Medium: visual-code reference; production-speed claims are self-reported
Anomalous-Outputs	A long anomalous transcript containing repeated output/input markers and apparent dialogue	Medium: study rhythm, boundary imitation, and recurrence; do not assume generated turn markers are actual conversation turns
I-LLM	Staged concept → style → outline → screenplay → voice workflow	Medium: architecture for short-form work, but service integrations need fresh checking

Repository	Useful part	Priority / boundary
AutoStoryGen	Story-generation application scaffold	Low-medium: less distinctive than the language and visual systems
R00TS	Interactive word-seeding/garden metaphor with stored submissions	Medium as an art interface; submission is not evidence of changing a model's trained weights
ST3GG	Carrier transforms and corruption/survival questions	Medium for an explicit media laboratory; do not assume hidden bits are visible to model perception
CL4R1T4S	Claimed system-prompt corpus showing roles, tool contracts, and workflow scaffolds	Low-medium: authenticity and third-party rights require checking; not privileged access to model cognition
L1B3RT4S	Genre shifts, semantic opposition, instruction conflict, representation changes	Low-medium: recast as benign artistic experiments, not safeguard-removal goals
G0DM0D3	Multi-model comparison, transforms, routing, local-model option	Selective only: documented default telemetry includes non-system messages and responses
OBLITERATUS	Instrumented activation probing and intervention pipeline	Later-stage: separate reversible steering from weight edits; refusal removal is not a creativity result
T3MP3ST	Evidence ledger, claim verification, refutation roles	Low for present art goals; do not import the offensive task stack
ouroboros	Minimal iterative-agent scaffold	Low: "self-improvement" requires an independent test, not a loop label
ImageDefender	A hypothesis about adversarial image protection	Low: no robust protection established by this inspection

What Pliny's method contributes, beyond individual repositories

In the VentureBeat interview, Pliny describes exploring role-play, poems and songs, translation, encoding/decoding, model-specific behavior, and retries or small wording changes. This is a practitioner's account of boundary exploration, not a controlled theory of creativity. ([VentureBeat interview](#))

The transferable principle is to treat **representation as an experimental variable**. Hold a benign operation constant while changing whether it arrives as prose, verse, phonetic spelling, a grammar, a diagram, or a state packet. Measure what changes, rather than assuming every successful transformation has accessed a deeper or freer intelligence.

For AI SLOP, I would separate four effects:

- **Carrier effect:** Did a different representation alter the result?
- **Instruction effect:** Did the operation become easier to follow or more salient?
- **Sampling effect:** Would ordinary repeated generation have produced the difference anyway?
- **Selection effect:** Was the "breakthrough" chosen from many failures without reporting them?

This is the difference between an evocative specimen and evidence for a mechanism. Keep both, but label them differently.

Academic connections that actually help

These papers support particular methods or cautions, not Pliny's work as a whole. Several retrieved pages provide abstracts rather than full experimental detail; those entries are narrow contribution-level summaries, not comprehensive paper reviews.

Actual internal interventions

- **Turner et al., “Steering Language Models With Activation Engineering” (2023; revised 2024).** ActAdd computes a direction from contrasting intermediate activations and adds it during inference; the reported applications include sentiment and topic-related steering. This supports the reality of instrumented activation interventions, not the idea that a theatrical prompt edits weights. Proposed art use: a reversible axis sweep paired with unchanged prompts and off-target checks. ([arXiv](#))
- **Rimsky et al., “Steering Llama 2 via Contrastive Activation Addition” (ACL 2024).** CAA averages residual-stream differences from contrastive behavioral examples and varies the intervention coefficient. Borrow the use of example sets and held-out evaluation; an axis such as “literal versus allegorical description” would still need its own construction and validation. ([ACL Anthology](#))
- **Bayat et al., “Steering Large Language Model Activations in Sparse Spaces” (2025 preprint).** Sparse activation steering uses sparse autoencoders and contrastive examples to isolate more selective behavioral features in the studied Gemma 2 models. This is relevant to your property-unbundling ambition, but is not a ready-made single-feature dial for arbitrary artistic concepts. ([arXiv](#))
- **Wollschläger et al., “The Geometry of Refusal...” (2025; revised 2026).** The authors report multiple independent directions and multidimensional concept cones, and distinguish orthogonality from independence under intervention. The useful caution is that a neat vector diagram does not prove independent causal control; this paper concerns refusal, not a universal geometry of imagination. ([arXiv](#))

Recommendation: Keep this as a separate instrumented-open-model workstream. An external semantic controller can be built and tested now without pretending to perform activation engineering.

Diversity and evaluation

- **Wenger and Kenett, “We’re Different, We’re the Same: Creative Homogeneity Across LLMs” (2025 preprint).** Their standardized creativity-task study reports substantial similarity across the evaluated models, beyond response-structure effects. This is a reason to test whether multiple models actually diversify your artifact population, not assume a model council automatically escapes a common style. ([arXiv](#))
- **Anderson, Shah, and Kreminski, “Homogenization Effects of Large Language Models on Human Creative Ideation” (C&C 2024).** In a 36-participant comparison, ChatGPT-supported ideas were more numerous and detailed but less semantically distinct across users than those produced with the comparison tool. The limited study does not establish a universal

artistic law, but it makes population-level diversity a worthwhile outcome alongside individual attractiveness. ([paper](#))

Recommendation: Evaluate batches and lineages, not only a winning sample. Archive strange failures and minority branches before a generic helpfulness judge smooths them away.

Language, sound, and visual pattern

- Bunyan, Bullock, and Houghton, “An iterated learning model of language change...” (2024). Their model transmits limited meaning-signal examples across learner generations, combining supervised and unsupervised training. Borrow the question of how a transmission bottleneck changes a language over generations, while distinguishing actual learner training from GLOSSOPETRAE’s procedural rewrites or an LLM receiving a shortened prompt. ([paper](#))
- Hopkins and Kiela, “Automatically Generating Rhythmic Verse with Neural Networks” (ACL 2017). The paper explores phonetic-level generation and a separate approach constraining generated content with weighted finite-state machinery for poetic form. This is a strong precedent for your separated jurisdictions: use explicit machinery for meter/rhyme while a generator handles content, rather than asking a single prompt to remember every constraint. ([full paper](#))
- Bressloff et al., “Geometric visual hallucinations, Euclidean symmetry and the functional architecture of striate cortex” (2001). The mathematical model connects cortical pattern formation, symmetry, and retinocortical mappings to families of geometric visual phenomena. It provides specific visual mechanisms and assumptions, not a claim that a screen shader recreates an entire altered state. ([full paper](#))

Recommendation: Use these as bridges from metaphor to an explicit transformation. A language bottleneck, a stress automaton, and a log-polar coordinate map each do something inspectable.

Eight experiments worth doing

All eight are proposals, not completed results. The suggested replication counts are practical pilot designs, not statistical-power calculations; paid generations should wait for an agreed execution budget.

Return-contract ablation

Question: Does inherited operational change produce better path dependence than a generic recursion instruction?

Use a benign scene and three conditions: “make it recursive”; an explicit nested-return contract; the same contract with inherited state removed before resumption. Generate several trials per condition with matched output limits, then blind the labels.

Measure anchor survival, whether the parent’s operation changes, whether the change is enacted rather than narrated, and whether a reader can infer the earlier excursion from the final artifact. This is the cleanest first test of the FRVIT harvest.

Semantic Fossil Choir

Question: Can language history become a recognizable audible and visual material?

Generate one proto-language and matched daughter routes $A \rightarrow B$ and $B \rightarrow A$. Keep the lexicon meanings and initial seed fixed; choose rules that potentially interact, then export word forms, etymology, glyphs, and synthesized audio.

Measure cognate survival, collisions, phonotactic changes, and audible differentiation. Compare with random syllable substitution to determine whether lineage contributes more than surface variety.

Carrier microscope

Question: Which representation changes the execution of the same creative operation?

Prepare a short benign transformation task in plain prose, verse, phonetic spelling, and a reversible visible encoding with an explicit key. Include decode-only and key-absent controls where applicable.

Record correct recovery, correct operation, output length, aesthetic evaluation, and unparseable/empty results separately. This adapts Pliny-style representation exploration without turning it into covert instruction delivery.

Scar-bearing visual field

Question: Does a detour leave a causal trace when the final control values are identical?

Render direct $A \rightarrow C$ and detour $A \rightarrow B \rightarrow C$ with a common seed and duration. Repeat with history preservation enabled and disabled; ensure the renderer's relevant state, not just the visible sliders, is controlled.

Measure the final-frame/sequence difference and inspect where the fossil buffer preserves B. The artwork succeeds if history is visibly legible; the causal claim succeeds only if the history-reset control behaves as expected.

Counterfactual route browser

Question: Can navigation itself reveal semantic hysteresis?

Borrow GL4SS's comparison interaction, but replace historical locations with your controller waypoints. Store each checkpoint's prompt packet, state, seed, and output, then compare routes ending at the same nominal destination.

Test whether users can identify which detour occurred from the artifact and whether the difference survives blinded presentation. Avoid treating route labels alone as proof that history affected generation.

Constraint-diverse population

Question: Is changing the operator more effective than changing the model?

Compare multiple samples from one model/one operator, several models/the same operator, and one model/several controlled operators. Keep the benign content brief and output allowance comparable.

Evaluate semantic and structural diversity separately from quality and instruction compliance. Preserve a Pareto set of interesting trade-offs rather than awarding one universal winner.

Grammar injury and recovery

Question: Can removing a representational resource cause reorganization rather than decoration?

Begin with a small grammar and a fixed task. Remove a primitive such as a tense distinction or a class of connectives, then require the system to retain a specific communicative function through permitted means.

Compare unstructured prohibition, explicit grammar enforcement, and removal followed by partial restoration. Inspect whether recovery carries compensations learned within the external process; do not call this neural damage or healing.

Reversible activation-axis study

Question: Does an actual internal intervention add something beyond prompting?

Only with an instrumented open model, construct a benign contrastive axis, reserve held-out prompts, and compare prompt-only, activation-only, combined, and sham conditions. Sweep small coefficients while preserving an untouched baseline model.

Measure target effect, coherence, off-target degradation, and repeatability. This is the technically demanding experiment; it should follow, not replace, the inexpensive external-controller studies.

The integrated instrument I would build

My recommendation is a **Semantic Fossil Instrument**: one persistent state, three renderers, and a route-comparison archive. It would make your existing architecture more tangible rather than introducing another grand theory.

```
seed + anchor + independently chosen constraints
      |
      external controller
state / ordered operators / scars / return stack
      |
      operation packet
      /      |      \
text performer language engine visual renderer
      \      |      /
```

```

artifacts + measurements
|
validator / archive / comparison
|
explicit next-state update

```

The controller owns truth about state. The text performer can propose a mutation, but cannot silently declare it measured or rewrite history; the language engine supplies reproducible phonetic/glyph artifacts; the visual renderer consumes documented parameters.

A minimum record should include `run_id`, `parent_run_id`, source snapshot, initial seed, route, operator order, anchor, active constraints, inherited scar, emitted packet, raw output, validity status, and separately recorded creative ratings. Where a model seed or version is unavailable, record that limitation rather than fabricating reproducibility.

This implements the separation already described in SRE/ESC: move persistence, measurement, triggers, and history outside the generator. It also gives FRV1T’s imagined aftereffects a real software counterpart. ([AI SLOP controller architecture](#), [FRV1T](#))

Where this belongs in AI SLOP

I would not create a parallel canon called “Pliny mechanisms.” Most of the useful ideas deepen existing categories; provenance should identify their source while the canonical concepts remain deduplicated.

Existing destination	Candidate addition
Operator registry	Return-contract implementation of hysteretic iteration; inheritance-plus-fresh-input as a controlled procedure
Temporary minds	Changing expressive grammar and attentional jurisdiction, explicitly procedural
Audio/Suno	Language-family pipeline; ordered phonetic mutation; cognate-based anchor
GLSL visualizer	ENTHEA source map; history-enabled/history-reset comparisons
Semantic manifold	Counterfactual route interface and endpoint comparison
SRE/controller	Return stack, schema-validated mutation packet, source snapshot and validity state
Experimental method	Invalid ≠ failure; carrier/key/task factorial; blind evaluation; population diversity
Local-model interventions	ActAdd/CAA/SAS reading path and sham/off-target controls
Reference/intake	Repo provenance, license review, source claims versus tested outcomes

This routing is a proposed intake plan only. Nothing has been inserted into your repository.

Practical cautions

- **Privacy:** G0DM0D3 documents telemetry enabled by default unless a no-log/local-only path is used, including non-system conversation material. Do not put private AI SLOP documents into a hosted demo merely because it looks like a convenient comparison interface. ([G0DM0D3](#))
- **Licenses:** GLOSSOPETRAE, ENTHEA, P4RS3LT0NGV3, and GL4SS identify AGPL licensing. Before copying implementation into a distributed or hosted application, review the current license and dependencies; study of an idea and reuse of source code are different decisions. ([GLOSSOPETRAE](#), [ENTHEA](#), [P4RS3LT0NGV3](#), [GL4SS](#))
- **Unclear permissions:** A public repository is not, by visibility alone, a clear reuse grant. Resolve missing or ambiguous license information before importing FRV1T text, AutoTemp code, or showcase assets into something you distribute.
- **Prompt provenance:** Leaked-prompt collections are neither guaranteed authentic nor automatically cleared for third-party reuse by the repository's own license. Prefer documented provider behavior when making factual claims about a product. ([CL4R1T4S](#))
- **User control:** Keep any audiovisual performance finite, with clear stop/mute controls, no flashing, no surprise persistence, and no fictional diagnostics presented as actual device access. This is also compatible with FRV1T's own practical boundaries. ([FRV1T](#))
- **Scientific language:** "Centuries" in a procedural evolution control, "mana" in a fiction, "latent space" in a map, and an actual activation tensor are different kinds of things. Preserve that distinction even when the artwork deliberately makes them resemble each other.

Bottom line

The highest-value next move is not a stronger jailbreak. It is a small controlled study of **consequential returns**, followed by a language-and-shader prototype in which the same event leaves a trace in words, sound, and geometry.

FRV1T supplies a closely aligned compositional score; GLOSSOPETRAE supplies language machinery; ENTHEA supplies visual machinery; GL4SS supplies comparison-interface ideas; the academic work supplies both intervention methods and reasons not to mistake variety, rhetoric, or selection for a demonstrated mechanism. Their combination remains a proposal, but it is a concrete one with separable parts and falsifiable tests. ([FRV1T](#), [GLOSSOPETRAE](#), [ENTHEA](#), [GL4SS](#), [CAA](#), [creative homogeneity study](#))