

Caused Weirdness: A Thesis on Structured Instability

Synthesized from the AI SLOP repository (github.com/merrypranxter/ai_slop) — canonical layer (docs/), principles, mechanisms, media practice, and experimental method. September 2026.

Abstract

This repo argues something that most generative art practice gets backwards: **weirdness is not an aesthetic you request, it is a structural consequence you cause**. The interesting output — the slop — appears when a generative system is handed a problem whose easy resolutions are all blocked, and is forced to invent a compromise. The thesis here is that this compromise, the repair the model invents under pressure, is the actual artwork. Everything else in the repo — the Temporary Minds, the Semantic Manifold Game, TOPOS-SRE, the DAVID protocols, the Suno jurisdiction tricks, the math-slop image recipes — is plumbing built to make that pressure precise, repeatable, and honest about what it does and doesn't know.

The whole project can be stated as one pipeline:

ordinary assumption → precise intervention → competing pressures → system attempts resolution → emergent artifact → observe → test → preserve or mutate

This document unpacks that pipeline into its load-bearing claims.

1. The core inversion: the problem is the artwork

The canonical question of the repo, and the one that separates it from prompt-craft folklore:

What problem can we give the generative system such that its attempt to solve the problem becomes the artwork?

Most generative art asks: *what weird thing should the model make?* That question produces what the repo calls a **keyword smoothie** — “surreal fractal face, psychedelic, glitch” — a stack of adjectives with no causal structure. The model resolves it the way it resolves everything: by falling into the nearest familiar basin. You get generic AI-weird, which is the most average thing on the internet right now.

The inversion: don't describe the output, **engineer the bind**. Define what must stay stable (an anchor), what rule attacks it (a pressure), why the two can't both be satisfied cheaply (an incompatibility), and what shape the forced repair might take (a consequence). The shorthand:

ANCHOR + PRESSURE + INCOMPATIBILITY + CONSEQUENCE + OBSERVATION

Weak: “body with circuit-board patterns.” That's costume. Strong: biological anatomy must remain wet continuous tissue, but routing, redundancy, junction behavior, buffering, and failure propagation obey circuit-network rules, and the visible anatomy must satisfy both mandates at once. The image that results is not a decorated body — it's a **governance error made visible**. The model had nowhere boring to go.

2. Constraint governance: the wrong simulator, correctly used

The repo's most generative design model is also its most carefully quarantined one. The generator is treated as a **Governance Engine** reconciling mismatched mandates — but the repo is explicit that this

is a *metaphor with procedural consequences*, not an architecture claim.

The mistake it attacks is **the wrong simulator**: pick a subject, pick a style, ask the model to simulate the combination, hope for weirdness. Simulation of a coherent world is exactly what models are trained to do, so they do it — everything gets domesticated into one internally consistent universe, and consistency is the enemy of slop. The repo names this gravitational pull **Statistical Comfort**: biology behaves biologically, genre behaves like genre, the model settles into the regime where its learned regularities all hold at once.

The countermove is **Domain-Remapping: the Alien Governance Mandate** — force a target system to remain recognizably itself while another domain's *operational laws* govern part of its behavior. The crucial word is **laws, not imagery**. You don't borrow the source domain's nouns (mushrooms, courts, germs — costume); you borrow its causal relations (lifecycle, jurisdiction, infection dynamics) and give them real jurisdiction over part of the target. The artifact is the negotiation, not the blend.

Two follow-on rules matter as much as the model itself:

- **Keep both sides alive.** A contradiction is interesting only when neither side can be cheaply ignored. Hence separated jurisdictions in audio (harmony obeys one system, rhythm another, timbre a third — the track must negotiate rather than genre-blend), identity anchors in image, hard invariants in controllers. The repo's named failure mode **condition dropout** — the model solving tension by quietly ignoring the weaker condition — is the thing every experiment has to defend against.
- **Productive misunderstanding is material.** When the generator interprets an instruction "wrong" in a structurally coherent way, don't normalize it. Ask what rule the output invented, whether the mistake reproduces as a family, whether it can be promoted to an explicit mandate on the next pass. A misunderstanding graduates from accident to operator only when its useful relation can be stated and tested.

3. Path dependence: if B didn't scar, B was decoration

The repo's most distinctive formal commitment is **ancestry preservation**: $A \rightarrow C$ must not equal $A \rightarrow B \rightarrow C$. If the system passed through B, the next state must carry a traceable consequence of B — a scar, a mutation, a changed interpretation. Otherwise B was tourism.

This single principle spawns an entire operator family: Recall Mutation (every recall scars the remembered concept; the scarred version becomes the next memory), Semantic Recoil (a later state rewrites the interpretation of an earlier one, and reasoning proceeds from A rather than restoring pristine A), Observer Scar, Developmental Footprint, hysteretic iteration (compare $A \rightarrow B \rightarrow A$ against $B \rightarrow A \rightarrow B$; path dependence is real only if history changes what follows), and recursive artifact fossilization — feeding a generated accident back as conditioning material until repair noise becomes stable structure.

The consequence is a redefinition of what the artwork even is. The core artifact of the practice is not a prompt and not a final image — it's:

STATE + TRAJECTORY + HISTORY + CURRENT INTERPRETATION

The route is art. A path through concepts accumulates constraints no direct prompt to the endpoint could reproduce. Trajectories can be saved, replayed, compared, mutated, treated as compositional objects. This is why the repo keeps building stateful external machinery — Semantic Manifold Game, SRE/TOPOS-SRE, controllers with ledgers and attractor histories — instead of collecting better prompts. One-shot prompting structurally cannot have path dependence. A ledger can.

4. Nouns into operations: the operator discipline

The operator registry is the repo's immune system against vibes. Its graduation rule: a phrase belongs in the registry only when its operation can be stated **without the aesthetic costume** and it changes what transitions, outputs, or selections become legal.

The operators themselves share one grammar — translate any concept into what it *does*: what changes, what's conserved, what flows, what bifurcates, what persists after the cause disappears, what boundary must be maintained. Then map that behavior into the target medium's actual degrees of freedom. A few canonical shapes:

- **Property Unbundling** — detach a property from its owner, give it to something else, follow the consequences literally.
- **Primitive Deletion** — remove one load-bearing primitive *without replacement*; reconstruct around the wound rather than sneaking in a synonym. (Explicitly not the same as recompilation, which installs a substitute.)
- **Alien Distance Metric** — change the ruler defining conceptual nearness (failure behavior, maintenance burden, collapse trajectory), then choose neighbors by the new ruler.
- **Omission Field** — keep X real but ban its direct representation; triangulate it from independent external consequences.
- **Destructive Compression / Alien Substrate Transference** — strip a problem to a causal invariant, destroy its surface identity, solve in a foreign substrate, bring back only structural consequences.
- **Material Anchoring / Causal Artifact Chain** — make abstraction tactile (foam latex, emulsion, magnetic tape) and make degradation *caused* by the transformation rather than pasted on as a filter. “VHS glitch” as adjective is banned; telecine generation loss with head-switching timing is an operator.

And the regulators — the repo's built-in defense against its own success: **Structural Satiation** fingerprints a winning move's skeleton and temporarily rejects near-repeats even when the vocabulary changes; **Cliché Mortality** retires operators that win too often, so one productive trick never becomes the project's new generic default. The system is allergic to itself on purpose.

5. Epistemic hygiene as an artistic material

The thing that makes this repo unusual among “AI weirdness” projects is that it treats **label discipline as part of the art practice**, not as bureaucracy. Every mechanism claim carries a status: OBSERVED, SUPPORTED, HYPOTHESIS, PROCEDURAL, METAPHOR, SPECULATIVE, ARCHIVED. The law: *never upgrade a label because prose sounds confident*.

This exists because the archive contains the opposite — jailbreak rhetoric, “the glitch is the machine's true self,” claims of direct access to hidden tensors, manifesto fantasies about epistemic sabotage. The repo's solution is elegant: **two authority layers**. The archive (`originals/`, `ai_readable/`) is authoritative about *what was said* — nothing messy is discarded. The canonical layer (`docs/`) is authoritative about *what the project currently means* — nothing theatrical gets promoted. The ghost gets to stay in the archive; it does not get cited as mechanism.

The payoff formula is the roundtable's best line: **creative utility and mechanism confidence are separate scores**. A technically wrong theory can produce excellent art. Keep the artifact; demote the explanation. “Honest slop” is kept as an aesthetic metaphor and explicitly denied as neuroscience — a glitch isn't a window into a machine soul, it's evidence of *where our instructions, the model's learned*

regularities, and the medium's representation fail to line up cleanly. Which is, as the philosophy doc notes, plenty weird without inventing a silicon ghost.

The experimental method operationalizes all of it: controls, ablations, dose-response sweeps, 2×2 interaction tests (if A+B looks like A alone, B wasn't doing causal work), seed repeats, artifact *families* instead of favorite one-offs, and a graduation rule — an idea becomes a reusable operator only when it produces a recognizable mutation signature across unlike inputs **and** the characteristic effect disappears under ablation. Until then: a juicy hypothesis, labeled as such. Even failure is disciplined — an ugly diagnostic result outranks a gorgeous one-off, because failure must say *which* condition dropped, *which* variable dominated, *which* story was wrong.

6. What the thesis actually claims

Stripped to its skeleton, the repo's position is:

1. **No free weirdness.** Interesting generative output is downstream of identifiable operations on assumptions. If you can't name the assumption attacked, the operation applied, the constraints still active, the expected artifact, and what would count as the model ignoring the problem — the idea belongs in the garden, not the registry.
2. **Conflict is a medium.** The artwork is the compromise produced by incompatible but still-active obligations. Designing the bind *is* the composition.
3. **History is a material.** Scars, ledgers, trajectories, and fossilized accidents make one-shot aesthetics obsolete; stateful systems are the natural instrumentation.
4. **Rigor is the enabling constraint, not the enemy.** Strange premises require *more* disciplined consequences, not fewer. Change one foundational rule, then reason harder. The label discipline, the ablations, the refusal to straighten the slop into a tidy-but-false theory — all of it exists to keep the weirdness *caused* rather than coincidental.

The anti-thesis the repo explicitly rejects is worth stating too: this is not randomness worship, not epistemic sabotage, not machine-mysticism, and not “make it weird.” The practice is already strange enough. It doesn't need that shit to be interesting.

What it needs — and what the repo builds, layer by layer — is pressure with nowhere boring to go, and a lab notebook honest enough to keep the wreckage.

Sources: AI_CONTEXT.md, README.md, docs/00-project-map.md, docs/01-principles/{structured-instability, constraint-governance, philosophy}.md, docs/02-mechanisms/operator-registry.md, docs/06-experimentation/experimental-method.md. Repo untouched; this document lives outside it.