

Diffusion Inversion, Attention & Latent Manipulation

Trajectory surgery, attention binding conflicts, and noised-latent interference as mechanisms of structured instability

Compiled: 2026-09-19 | For: merrypranxter/ai_slop (research-only; repository untouched)
Status labels follow the project's canonical epistemic scheme.

Contents

- 1 Executive summary
- 2 Technical background
- 3 Mechanism catalog (M1–M9)
- 4 Failure surfaces as artistic material
- 5 Promising experiments (E1–E4)
- 6 Cross-modal applications
- 7 What appears most useful for AI SLOP
- 8 Weak / dubious ideas worth rejecting
- 9 Open questions
- 10 Sources / bibliography
- 11 Licensing / provenance appendix

1 Executive summary

Diffusion models offer three distinct intervention substrates, and this topic covers all of them: the **trajectory** (the sequence of noised latents an inversion constructs back to a starting noise), the **attention maps** (spatial binding tables between text tokens and image regions), and the **latent field itself** (which can be spliced, blended, and arithmetically combined between denoising runs). Each substrate has a documented failure surface, and each failure surface is aimable: inversion error accumulation, attention leakage, the reconstruction–editability trade-off, and non-Gaussian intermediate latents are all characterized in peer-reviewed or benchmarked work — not folklore.

The strongest verified findings: (a) cross-attention control (Prompt-to-Prompt^[1]) established that cross-attention maps bind prompt tokens to spatial layout, and that injecting, replacing, or re-weighting those maps edits images while preserving structure — with three named operations that are, in effect, a tiny formal grammar for constraint conflict; (b) inversion is a leaky state channel: naive DDIM inversion accumulates error (concentrated in early steps, amplified by classifier-free guidance^{[7][12]}), and every modern method (null-text^[2], PnP inversion^[3], ReNoise^[4], eta-inversion^[5]) is a different trade between reconstruction fidelity and editability — meaning 'how wrong is the return path' is a tunable artistic parameter, not just a bug; (c) the modern rectified-flow stack (RF-Inversion, RF-Solver, FireFlow^{[9][10][11]}) brings these interventions to FLUX-class models in 8–15 steps, with diffusers and ComfyUI support, so the mechanisms are practically reachable in 2026; (d) attention leakage — edits bleeding into untargeted regions through inaccurate cross-attention maps — is documented as the root cause of editing failure^[8], which inverts into a mechanism: leakage is controllable semantic contamination.

Highest-value project connection: inversion gives AI SLOP a literal **state channel with a return path** — image → noise → edited image — where path dependence (a canonical project principle) is physically realized: $A \rightarrow C$ is provably not $A \rightarrow B \rightarrow C$ unless B leaves a scar in the trajectory. Attention injection is a binding-conflict instrument (two prompts fighting over the same pixels). Latent splicing is the lowest-cost, highest-chaos member of the family, already available as experimental ComfyUI nodes. Together these are DAVID-operator material for the image medium with genuine dose-response structure.

Honesty constraints honored: no claim that latent space is a literal navigable world (the project already treats it as a 'useful lie'); interpolation artifacts are explained mechanistically (off-manifold, non-Gaussian intermediates) rather than mystically; every starter-lead repo was individually verified for activity and license, and one provenance trap was found — there are **two unrelated papers named StyleDiffusion** (§3, M8).

2 Technical background

A latent diffusion model denoises a latent $z_T \sim N(0, I)$ into an image over T steps, conditioned on text embeddings via cross-attention. Three facts make it steerable. First, the sampling trajectory is a deterministic ODE path (DDIM, or the rectified-flow ODE in FLUX-class models) — it can be run backwards (*inversion*) to recover a noise latent that regenerates a given image, approximately. The approximation is the whole story: each inversion step locally linearizes a nonlinear update, errors accumulate, and classifier-free guidance amplifies them because the guided score was never meant to be inverted^{[7][12]}. Second, cross-attention maps M_t (pixels × tokens, per step) encode which pixels

listen to which words; self-attention features encode structure (Q, K) and appearance (V) separately^[13]. Third, the latent is just a tensor: between steps it can be blended, masked, spliced, or offset, and the VAE decoder will attempt to repair whatever it receives toward a plausible image — the decoder is a repair organ, and its repair failures are artifacts.

Verified starter leads (checked via GitHub API / HF on 2026-09-19):

Lead	Verified status	Assessment
bloc97/CrossAttentionControl	Real; 1,336 stars; MIT; frozen 2022-10	Historically important unofficial P2P implementation; superseded by google/prompt-to-prompt and diffusers attention processors. Reference only
cure-lab/PnPInversion	Real; 414 stars; frozen 2024-03; official ICLR 2024 code	Includes PIE-Bench (700 images) and runnable scripts for 9 editing methods — a ready-made comparative harness
wangkai930418/DPL	Real; 111 stars; NeurIPS 2023	The leakage-repair paper; most valuable to this project as <i>evidence about the failure surface</i> , not as a tool
sen-mao/StyleDiffusion	Real; 84 stars; active 2026-08	Prompt-embedding inversion for text-based editing (CVMJ 2024) — NOT the ICCV 2023 style-transfer paper of the same name (see M8)
furiosa-ai/eta-inversion	Real; 34 stars; ECCV 2024	Unified harness: 8 inversion methods × 4 editing methods via CLI flags — ideal for systematic artifact-family studies
AustinMroz/ComfyUI-Splice Tools	Real; 6 stars; GPL-3.0; frozen 2024-06	Experimental noised-latent manipulation nodes; tiny adoption, no validation — treat as SPECULATIVE tooling, interesting because it is raw
lajjadred/comfyui-lrw-nodes	Real; 12 stars; created 2026-05	Riemannian/Bayesian latent manipulation nodes; new, unvalidated — SPECULATIVE
Angusliuuu/Awesome-Controllable-Generative-Models-Papers	Real; 42 stars; active 2026-07	Decent current map of controllable generation incl. attention interpretation and spectral manipulation
HF: anshuln/visualizing_diffusion_attention	Author profile lists 'Visualizing Diffusion Attention' space; API slug unresolvable in this session	Exists per web index; treat link as needs-manual-confirm
HF: AswinMathew/latent-space-explorer	UNVERIFIED — not resolvable via HF API	Do not cite until confirmed

Table 1 — Starter-lead verification. Star counts and dates from the GitHub API, 2026-09-19.

Modern tooling reality: Hugging Face **diffusers** exposes pluggable attention processors (`set_attn_processor`) as a first-class API — custom cross-attention control is a supported extension point, not a monkey-patch^[14]. RF-Inversion ships in diffusers and was re-implemented by the ComfyUI community^[9]; FireFlow does FLUX inversion+editing in 8 steps^[11]. The old SD 1.5 stack (P2P, null-text) still works but is legacy infrastructure; the 2026 working stack is FLUX/SD3-class

rectified-flow models plus ComfyUI for node-level latent access.

3 Mechanism catalog

Same template as Topic 01. Tiers: ESTABLISHED / DOCUMENTED, PLAUSIBLE INTERPRETATION, PROJECT HYPOTHESIS, SPECULATION. Creative utility and mechanism confidence scored separately.

M1 — Cross-attention injection (Prompt-to-Prompt)

SOURCE: Hertz et al., ICLR 2023^[1]; code: google/prompt-to-prompt; bloc97/CrossAttentionControl (unofficial, frozen)

SOURCE TYPE: Peer-reviewed paper + official and community implementations

LICENSE / USAGE NOTES: MIT (bloc97 repo); google repo Apache-2.0 (check); diffusers pipeline support varies by model

WHAT IT ACTUALLY DOES: Runs diffusion twice on prompt pairs sharing a seed; injects/replaces/reweights cross-attention maps between runs. Three operators: word swap (replace token, keep layout), prompt refinement (add tokens, freeze old attention), attention re-weighting (scale one token's map)

CONTROLLABLE VARIABLES: cross_replace_steps and self_replace_steps (fraction of steps to control, per word if desired), local_blend masks, equalizer coefficients per token

WHAT STATE OR STRUCTURE IT PRESERVES: Spatial layout and geometry — the cross-attention maps ARE the layout; freezing them freezes composition

WHAT IT CHANGES: Semantic content bound to the replaced/added/reweighted tokens

FAILURE SURFACE: Attention leakage: maps are inaccurate, edits bleed into background and semantically related distractors^[8]; over-injection freezes too much and the edit cannot materialize

ORIGINAL PURPOSE: Text-only image editing without masks

POSSIBLE AI-SLOP MUTATION: Binding conflict: swap a token to a word whose semantics contradict the frozen spatial structure ('castle' → 'flame' with castle layout frozen) — the model must render a flame shaped like a castle. The frozen map is one jurisdiction, the new token another

EXPECTED OBSERVABLE ARTIFACT: Impossible-object renderings; category collisions with coherent geometry; the exact replace-step fraction where the edit half-exists

AMPLIFICATION / ITERATION METHOD: Sweep cross_replace_steps from 0→1 for a fixed contradictory swap; the transition curve maps where binding breaks

CONTROL OR ABLATION TEST: Same swap without injection (expect full layout change); injection with identical prompts (expect bit-identical reproduction)

LIMITATION: Built for U-Net cross-attention; MMDiT models (SD3/FLUX) fuse text and image tokens, so the mechanism needs re-derivation per architecture^[15]

MECHANISM CONFIDENCE: High

APPLICABLE MEDIA: IMAGE; video derivatives exist (Video-P2P); audio diffusion uses the same attention pattern (Topic 06)

AI_SLOP CONNECTION: DAVID operator (binding conflict); mutation operator; experiment

CONFIDENCE TIER: ESTABLISHED / DOCUMENTED

CREATIVE UTILITY: Very high

M2 — Null-text inversion (pivotal null embedding optimization)

SOURCE: Mokady et al., CVPR 2023^[2]

SOURCE TYPE: Peer-reviewed paper + code

LICENSE / USAGE NOTES: Open code; per-image optimization (~3M extra parameters stored per image at 50 steps); minutes per image

WHAT IT ACTUALLY DOES: Runs DDIM inversion, then optimizes the unconditional ('null text') embedding at each step (pivoting on the source-prompt trajectory) so the guided reverse path reproduces the real image

CONTROLLABLE VARIABLES: Optimization steps, pivot strength, number of pivotal embeddings stored, guidance scale during reconstruction

WHAT STATE OR STRUCTURE IT PRESERVES: The real image itself — this is the highest-fidelity classical reconstruction route for SD 1.5-class models

WHAT IT CHANGES: Makes real photographs editable by P2P-class methods

FAILURE SURFACE: The optimization IS a pressure valve: the null embeddings absorb all accumulated inversion error. Editing after null-text inversion works partly because the error was paid off in advance; reconstruction of unedited regions can still destabilize under complex prompts^[7]

ORIGINAL PURPOSE: Faithful real-image editing

POSSIBLE AI-SLOP MUTATION: Read the optimized null embeddings as a per-image 'error autobiography': 50 vectors recording exactly what the model had to believe, step by step, to make this image possible. Interpolate two images' null-text trajectories to hybridize their debts rather than their content

EXPECTED OBSERVABLE ARTIFACT: Images that render the compromise between two unrelated scenes' correction histories — debt hybridization instead of image blending

AMPLIFICATION / ITERATION METHOD: Trajectory interpolation coefficient sweep; compare against direct latent interpolation as ablation

CONTROL OR ABLATION TEST: Same interpolation on raw DDIM-inverted latents (no null-text); identity reconstruction ($\alpha=0,1$ endpoints)

LIMITATION: Slow per image; storing per-step embeddings is unwieldy; null-text vectors are not designed to be semantically meaningful — their 'autobiography' reading is an interpretation, not a guarantee

MECHANISM CONFIDENCE: High for the mechanism; the debt-hybridization variant is PROJECT HYPOTHESIS

APPLICABLE MEDIA: IMAGE

AI SLOP CONNECTION: State preservation study; experiment; validator (error as measurable quantity)

CONFIDENCE TIER: ESTABLISHED (mechanism) / PROJECT HYPOTHESIS (variant)

CREATIVE UTILITY: High

M3 — PnP Inversion / Direct Inversion (branch disentangling)

SOURCE: Ju et al., ICLR 2024^[3]; code: cure-lab/PnPInversion + PIE-Bench

SOURCE TYPE: Peer-reviewed paper + benchmark + code

LICENSE / USAGE NOTES: Open; repo frozen 2024-03 but complete; benchmark downloadable

WHAT IT ACTUALLY DOES: Three lines: during denoising of the inverted latent, adds the per-step deviation between the guided reverse path and the recorded inversion path to the SOURCE branch only, leaving the target/edit branch untouched

CONTROLLABLE VARIABLES: Which branch receives corrections, per-step correction on/off, choice of base inversion and editing method (9 editors wired in)

WHAT STATE OR STRUCTURE IT PRESERVES: Separates responsibilities: source branch preserves essential content, target branch preserves edit fidelity

WHAT IT CHANGES: Editing quality at zero optimization cost; order-of-magnitude faster than null-text

FAILURE SURFACE: The deviation term is a memory of how wrong the inversion was. Deliberately under-correcting (partial correction factor) reintroduces controlled inversion debt into the edit

ORIGINAL PURPOSE: Cheap high-fidelity inversion-based editing

POSSIBLE AI-SLOP MUTATION: Partial-correction dosing: apply fraction λ of the PnP deviation term and sweep λ . At $\lambda=1$ faithful editing; at $\lambda=0$ raw inversion error; between them, the ghost of the original image fights the edit

EXPECTED OBSERVABLE ARTIFACT: Spectral persistence of the original subject through radical edits; double-exposure-like interference that is mechanically caused, not prompted

AMPLIFICATION / ITERATION METHOD: λ sweep $\times$ edit strength (target prompt guidance) as a 2D phase map; PIE-Bench categories give standardized territory

CONTROL OR ABLATION TEST: $\lambda=0$ and $\lambda=1$ endpoints; same λ with null-text inversion as base (debt pre-paid $\rightarrow$ ghost should vanish)

LIMITATION: SD 1.5-era benchmark stack; λ -scheduling is not part of the published method (it is our intervention)

MECHANISM CONFIDENCE: High that λ controls debt; Medium for specific ghost aesthetics

APPLICABLE MEDIA: IMAGE

AI_SLOP CONNECTION: DAVID operator (debt dosing); state-controller rule (λ schedule); experiment

CONFIDENCE TIER: ESTABLISHED / DOCUMENTED (method) / PROJECT HYPOTHESIS (λ dosing)

CREATIVE UTILITY: Very high

M4 — Iterative / fixed-point inversion (ReNoise, AIDI, Eta)

SOURCE: Garibi et al., ReNoise (ECCV 2024)^[4]; Pan et al., AIDI (ICCV 2023); furiosa-ai/eta-inversion (ECCV 2024)^[5]

SOURCE TYPE: Peer-reviewed papers + code; eta-inversion repo unifies 8 inversion methods $\times$ 4 editors

LICENSE / USAGE NOTES: Open; eta-inversion CLI exposes `inv_method` and `edit_method` flags directly

WHAT IT ACTUALLY DOES: Treats each inversion step as an implicit equation and solves it by fixed-point iteration / iterative renoising, averaging predictions to stabilize the forward trajectory — improves reconstruction at fixed compute, works on few-step distilled models

CONTROLLABLE VARIABLES: Renoising iterations per step, averaging window, edit-enhancement and noise-correction scales (model-specific but stable)

WHAT STATE OR STRUCTURE IT PRESERVES: The return path to the image's own noise seed, convergent and numerically stable

WHAT IT CHANGES: Raises the fidelity floor; also demonstrates that inversion quality is a dial with measurable reconstruction-error readouts (PSNR/LPIPS/structure distance)

FAILURE SURFACE: Few-step distilled models invert worst (linearization error explodes with step size) — the regime where artifacts are guaranteed unless corrected^[4]

ORIGINAL PURPOSE: Accurate inversion incl. accelerated models

POSSIBLE AI-SLOP MUTATION: Deliberate under-convergence: stop the fixed-point iteration after 1 pass on a model that needs 8 — a principled 'half-solved equation' generator whose error structure is

documented rather than random

EXPECTED OBSERVABLE ARTIFACT: Reconstructions with systematic directional smearing (the residual of the unsolved implicit step), distinct from generic blur or noise

AMPLIFICATION / ITERATION METHOD: Iteration-count sweep per step; eta-inversion harness gives 8-method comparison under identical conditions — artifact families can be cross-method fingerprinted

CONTROL OR ABLATION TEST: Full convergence baseline; naive DDIM at same compute

LIMITATION: Hyperparameters are model-specific; under-convergence artifacts will vary per model and need re-mapping

MECHANISM CONFIDENCE: High

APPLICABLE MEDIA: IMAGE; video (same inversion problem in T2V models)

AI_SLOP CONNECTION: Mutation operator (controlled under-solving); experiment; validator (reconstruction metrics)

CONFIDENCE TIER: ESTABLISHED / DOCUMENTED

CREATIVE UTILITY: High

M5 — Rectified-flow inversion stack (RF-Inversion / RF-Solver / FireFlow)

SOURCE: Rout et al., ICLR 2025^[9]; Wang et al., RF-Solver (ICML 2025)^[10]; He et al., FireFlow (2024)^[11]

SOURCE TYPE: Peer-reviewed papers + code + diffusers/ComfyUI integrations

LICENSE / USAGE NOTES: RF-Inversion in diffusers (pipe.invert API); ComfyUI community port exists; FireFlow code public

WHAT IT ACTUALLY DOES: Inverts FLUX-class rectified-flow models: RF-Inv via dynamic optimal control, RF-Solver via high-order Taylor expansion of the ODE, FireFlow via a midpoint solver with first-order cost and second-order precision. Editing: replace self-attention V-features at early denoising steps with features stored during inversion

CONTROLLABLE VARIABLES: Solver order/steps, gamma (RF-Inv), V-injection steps and layers, edit prompt

WHAT STATE OR STRUCTURE IT PRESERVES: V-feature injection anchors layout and appearance while the new prompt steers semantics — structure/content jurisdiction split at the feature level

WHAT IT CHANGES: Brings inversion editing to 2026-frontier open models in seconds (8–18 NFEs)

FAILURE SURFACE: RF-Solver's error can INCREASE after ~25 steps post-convergence^[11] — an over-solving artifact regime; V-injection timing controls the structure/edit tension

ORIGINAL PURPOSE: Fast, accurate inversion+editing for rectified-flow models

POSSIBLE AI-SLOP MUTATION: Cross-image V-transfusion: inject the V-features of image A's inversion into the denoising of unrelated image B's latent under prompt C — three jurisdictions, one stream

EXPECTED OBSERVABLE ARTIFACT: A's appearance grammar haunting B's structure under C's semantics; FireFlow's own finding that first-step-only injection works suggests a sharp temporal sensitivity to exploit

AMPLIFICATION / ITERATION METHOD: Sweep which denoising steps receive foreign V-features; chain transfusions (output becomes next inversion source)

CONTROL OR ABLATION TEST: Self-V injection (must reproduce standard editing); V from a blank/noise inversion

LIMITATION: Findings are young (2024–25); per-model behavior varies; MMDiT joint attention differs from U-Net assumptions

MECHANISM CONFIDENCE: High for the stack; Medium for the three-way transfusion
APPLICABLE MEDIA: IMAGE; VIDEO (RF-Edit demonstrated on video^[10] — Topic 10 bridge)
AI_SLOP CONNECTION: DAVID operator (transfusion); experiment; connection to Topic 10
CONFIDENCE TIER: ESTABLISHED (stack) / PROJECT HYPOTHESIS (transfusion)
CREATIVE UTILITY: Very high

M6 — Self-attention feature injection (PnP features, MasaCtrl)

SOURCE: Tumanyan et al., CVPR 2023 (Plug-and-Play); Cao et al., ICCV 2023 (MasaCtrl); systematized in the 2025 attention survey^[13]
SOURCE TYPE: Peer-reviewed papers + code + survey
LICENSE / USAGE NOTES: Open; both runnable from the PnPInversion harness
WHAT IT ACTUALLY DOES: During the edit run, replaces self-attention components with those recorded from the source run: PnP injects Q/K (structure), MasaCtrl swaps K/V (mutual self-attention for non-rigid pose/action changes). Survey-confirmed division: Q,K encode structure/spatial arrangement; V encodes appearance (color, texture, shape)
CONTROLLABLE VARIABLES: Which of Q/K/V is replaced, which layers, which denoising steps (injection schedule)
WHAT STATE OR STRUCTURE IT PRESERVES: Source image structure (Q/K injection) or appearance (V injection) — selective jurisdiction
WHAT IT CHANGES: Pose/action/identity under the edit prompt while the injected channel resists
FAILURE SURFACE: Wrong channel for the task = the failure IS informative: injecting V when structure should move yields 'rubber-stamped' content; over-injection collapses the edit entirely
ORIGINAL PURPOSE: Non-rigid and structure-preserving editing
POSSIBLE AI-SLOP MUTATION: QKV transplant matrix: all 8 cross-substitution combinations between two source images (A's Q into B, A's K into B, ...) — a systematic taxonomy of hybrid images with known anatomies
EXPECTED OBSERVABLE ARTIFACT: A grid of 'impossible kinships': same-structure-different-physics, same-appearance-different-geometry, and the genuinely alien combinations (Q-only swaps)
AMPLIFICATION / ITERATION METHOD: Injection-schedule sweep per combination; recursive: a hybrid's features become next-round donor
CONTROL OR ABLATION TEST: Identity transplant ($A \rightarrow A$) per channel; no-injection edit baseline
LIMITATION: U-Net-centric evidence; DiT self-attention roles less cleanly separated (FireFlow's V-only success hints the division survives but shifted)
MECHANISM CONFIDENCE: High
APPLICABLE MEDIA: IMAGE; VIDEO (spatio-temporal variants documented^[13])
AI_SLOP CONNECTION: DAVID operator (transplant matrix); experiment; new mechanism class (channel anatomy)
CONFIDENCE TIER: ESTABLISHED / DOCUMENTED
CREATIVE UTILITY: Very high

M7 — Cross-attention leakage and its repair (DPL)

SOURCE: Wang et al., NeurIPS 2023 (DPL)^[8]; code: wangkai930418/DPL
SOURCE TYPE: Peer-reviewed paper + code

LICENSE / USAGE NOTES: Open; SD-based

WHAT IT ACTUALLY DOES: Diagnoses that editing failures trace to inaccurate cross-attention maps: the edited token's map covers distractor regions, so edits bleed. Repairs by optimizing per-noun dynamic tokens with leakage-repairment losses that force maps onto the correct nouns

CONTROLLABLE VARIABLES: Leakage loss weights, targeted nouns, repair strength

WHAT STATE OR STRUCTURE IT PRESERVES: Untargeted regions — repair = jurisdictional enforcement

WHAT IT CHANGES: Localizes edits that vanilla P2P smears

FAILURE SURFACE: Leakage itself is the project's interest: the semantic-gravity field that makes 'edit the cat' also repaint the sofa. DPL quantifies it; we can decline to repair it — or tune the repair halfway

ORIGINAL PURPOSE: Fixing bleed-through in text-based editing

POSSIBLE AI-SLOP MUTATION: Leakage dial: interpolate between DPL-repaired and unrepaired attention (repair strength λ); at intermediate λ the edit leaks exactly as far as semantic similarity carries it — a measured contamination radius

EXPECTED OBSERVABLE ARTIFACT: Contagion aesthetics: edits spreading along semantic relatedness gradients (editing 'king' also crowns the queen and the throne)

AMPLIFICATION / ITERATION METHOD: λ sweep mapped against semantic-distance probes of affected regions

CONTROL OR ABLATION TEST: Full repair; zero repair; repair applied to a distractor noun instead

LIMITATION: Repair is per-prompt optimization; the λ -dial is our construction, not the paper's

MECHANISM CONFIDENCE: High for the diagnosis; Medium for the dial

APPLICABLE MEDIA: IMAGE

AI_SLOP CONNECTION: Mutation operator; state-controller rule (λ); validator (semantic contamination radius as metric)

CONFIDENCE TIER: ESTABLISHED (diagnosis) / PROJECT HYPOTHESIS (dial)

CREATIVE UTILITY: High

M8 — PROVENANCE TRAP: two StyleDiffusions

SOURCE: (a) Wang, Zhao, Xing — ICCV 2023: CLIP-space content/style disentanglement + diffusion style-removal^[6]. (b) sen-mao/StyleDiffusion — CVMJ 2024: prompt-embedding inversion for text-based editing (the starter lead; active 2026-08)

SOURCE TYPE: Two unrelated peer-reviewed papers sharing a name

LICENSE / USAGE NOTES: Both open; unrelated codebases

WHAT IT ACTUALLY DOES: (a) Removes style via partial diffusion on luma-grayscale images, then re-injects learned style (style removal module with controllable T_{remove}). (b) Inverts the prompt embedding space (P2Plus modifies both conditional and unconditional self-attention)

CONTROLLABLE VARIABLES: (a) T_{remove} (how much style is dispelled), CLIP disentanglement weight; (b) embedding inversion steps, P2Plus injection schedules

WHAT STATE OR STRUCTURE IT PRESERVES: (a) Content structure through style removal; (b) layout through dual-branch attention control

WHAT IT CHANGES: (a) Style channel only; (b) semantic content via embedding surgery

FAILURE SURFACE: The citation ambiguity itself is the failure surface for research hygiene; technically, (a)'s partial removal has a removal-depth artifact zone and (b)'s embedding inversion can land off-manifold

ORIGINAL PURPOSE: Style transfer / editing, respectively

POSSIBLE AI-SLOP MUTATION: (a)'s style-removal dial run in reverse: stop style removal at intermediate depth so an image carries a half-removed style — stylistic phantom limb

EXPECTED OBSERVABLE ARTIFACT: Images whose brushwork is present-but-decoupled from their geometry; removal-depth sweep gives the family

AMPLIFICATION / ITERATION METHOD: T_remove sweep; iterate removal on already-restyled outputs (compounding removal debt)

CONTROL OR ABLATION TEST: Full removal baseline; no-removal baseline

LIMITATION: (a) requires its specific training setup; the phantom-limb reading is PROJECT HYPOTHESIS

MECHANISM CONFIDENCE: High (both papers documented); Medium for the mutation

APPLICABLE MEDIA: IMAGE

AI_SLOP CONNECTION: Mutation operator; provenance lesson for source-audit.md

CONFIDENCE TIER: ESTABLISHED / DOCUMENTED + PROVENANCE WARNING

CREATIVE UTILITY: Medium-High

M9 — Noised-latent splicing / blending (ComfyUI practice)

SOURCE: Blended Latent Diffusion (SIGGRAPH 2023)^[16]; AustinMroz/ComfyUI-SpliceTools (experimental nodes); lajjadred/comfyui-lrw-nodes (2026, Riemannian/Bayesian latent ops)

SOURCE TYPE: Peer-reviewed paper + unvalidated community nodes

LICENSE / USAGE NOTES: SpliceTools GPL-3.0 (copyleft — note for any distributed derivative); lrw-nodes license check needed; Blended Latent Diffusion code open

WHAT IT ACTUALLY DOES: Operates on the latent mid-trajectory: spatially blend the noised original with the generated latent each step (masked local editing), or splice/perturb noised latents directly (SpliceTools), or treat latent interpolation with geometric corrections (lrw)

CONTROLLABLE VARIABLES: Splice mask, noise level at splice time, blend schedule, interpolation path (linear vs spherical vs Riemannian)

WHAT STATE OR STRUCTURE IT PRESERVES: Unmasked regions' trajectory; the VAE decoder attempts global repair of the seam

WHAT IT CHANGES: The composite image's coherence boundary — the seam is where two incompatible trajectories meet

FAILURE SURFACE: Seam artifacts, pixel-level noise residue (documented in the paper^[16]), and non-Gaussian intermediates: linear interpolation between latents leaves the decoder's comfort zone, producing washed-out or structurally alien images

ORIGINAL PURPOSE: Masked local editing; latent-space tooling

POSSIBLE AI-SLOP MUTATION: Trajectory splicing instead of image splicing: take the first half of prompt A's denoising trajectory and the second half of prompt B's (temporal splice in noise space). The decoder must resolve a latent whose early structure decisions and late detail decisions belong to different worlds

EXPECTED OBSERVABLE ARTIFACT: Chimera images where composition and texture belong to different prompts; splice-timestep sweep maps the dominance frontier (early splice = B-flavored, late = A-flavored)

AMPLIFICATION / ITERATION METHOD: Splice-timestep sweep × mask density; recursive splicing of outputs (artifact inheritance)

CONTROL OR ABLATION TEST: Same-timestep splice with matched prompt (A/A — should reproduce); pixel-space compositing comparison

LIMITATION: Community nodes are unvalidated (6–12 stars); every splice claim must be re-derived empirically per model

MECHANISM CONFIDENCE: High for blending; Medium for temporal trajectory splicing

APPLICABLE MEDIA: IMAGE; VIDEO (frame-latent splicing); GLSL-ADJACENT (noise-field thinking)

AI_SLOP CONNECTION: DAVID operator (temporal splice); mutation operator; state-controller rule (mask schedules); experiment

CONFIDENCE TIER: ESTABLISHED (blending) / PLAUSIBLE INTERPRETATION (traversal splicing) / SPECULATION (lrw geometry)

CREATIVE UTILITY: Very high — cheapest mechanism in the catalog

4 Failure surfaces as artistic material

Five failure surfaces in this topic are characterized well enough to be aimed at. Each row is an instrument.

Failure surface	Documented behavior	Aiming parameters	Evidence tier
Inversion debt (error accumulation)	DDIM inversion errors concentrate in early steps and compound; CFG amplifies them; few-step distilled models invert worst. Reconstruction $\neq$ editability: paying debt (null-text) and tolerating debt (naive DDIM) edit differently ^{[7][12]}	Inversion steps, guidance scale, iteration count, correction fraction λ (PnP)	ESTABLISHED
Attention leakage	Edited token's cross-attention map covers semantically related distractors; edits bleed into untargeted regions. Root cause documented, repair quantified (DPL) ^[8]	Repair strength, target noun choice, re-weighting coefficient	ESTABLISHED
Over-injection freeze	Injecting attention/features for too many steps locks the image to the source so hard the edit cannot materialize — constraint suffocation	cross_replace_steps / injection step fraction	ESTABLISHED ^[1]
Non-Gaussian intermediate latents	Naive interpolation/averaging of latents leaves the decoder's trained manifold: washed-out, low-contrast, structurally alien outputs; inversion latents carry spatial correlations that diverge from Gaussian seed statistics ^[12]	Interpolation path, blend schedule, splice timestep	ESTABLISHED
CFG mode-shifting / over-saturation	High guidance scales shift modes and concentrate samples (over-saturation); CFG's composite score differs from any true score of the model ^[17]	Guidance scale, guidance rescale	ESTABLISHED
Solver over-correction	RF-Solver's reconstruction error can rise again past ~ 25 steps post-convergence; more precision is not monotonically better ^[11]	NFE count, solver order	ESTABLISHED (single paper — replicate before relying)

Table 2 — Failure surfaces with aiming parameters.

5 Promising experiments

E1 — The Debt Dial (controlled inversion error as material)

HYPOTHESIS: Inversion debt — the gap between the true return path and the computed one — behaves as a graded artistic material: a λ -dose of PnP correction produces an ordered artifact family from ghost-persistence to clean edit, not random degradation

REPRESENTATIONAL TENSION: The correction term is the memory of the original image's true path; withholding it forces the edit to negotiate with a wrong history

COMPETING CONSTRAINTS: Edit prompt fidelity vs. trajectory truthfulness to the source image

WHY THE SYSTEM MAY STRUGGLE: The denoiser was trained on legal Gaussian trajectories; an under-corrected trajectory is legal-shaped but semantically false

EXPECTED RESULT: λ sweep yields: $\lambda \approx 1$ clean edit; mid- λ spectral double-exposure artifacts (original subject bleeding through radical semantic change); $\lambda \approx 0$ near-failed edit dominated by source structure

NEGATIVE CONTROL: Same λ schedule applied to a fully-converged ReNoise inversion (debt pre-paid → mid- λ artifacts should vanish)

ABLATION: $\lambda=0$, $\lambda=1$, λ randomized per step (schedule vs constant)

DOSE RESPONSE: $\lambda \in \{0, 0.25, 0.5, 0.75, 1\} \times$ guidance scale $\in \{3, 7.5, 12\}$ on PIE-Bench categories; metrics: structure distance to source, CLIP-to-target, artifact rating

ITERATIVE VERSION: Edit the output again with the same λ — debt compounds across generations (artifact inheritance; links Topic 09)

FALSIFICATION CONDITION: Mid- λ outputs are indistinguishable from λ -randomized controls on artifact-structure probes — debt is just noise, not structure

E2 — Binding War (contradictory word-swap under frozen attention)

HYPOTHESIS: A P2P word swap whose replacement semantically contradicts the frozen layout produces a stable artifact family — 'impossible objects with coherent geometry' — whose onset is a measurable function of injection depth

REPRESENTATIONAL TENSION: Frozen cross-attention map = spatial jurisdiction of the old word; new token embedding = semantic jurisdiction of the new word; both claim the same pixels

COMPETING CONSTRAINTS: Layout preservation pressure vs. semantic fidelity pressure

WHY THE SYSTEM MAY STRUGGLE: The model has no training distribution of flames shaped like castles; every denoising step must re-arbitrate

EXPECTED RESULT: Injection-fraction sweep reveals three regimes: swap fails (layout wins), hybrid (contradiction rendered — the artifact zone), layout breaks (semantics wins). Hybrid zone boundaries are prompt-pair-dependent but reproducible

NEGATIVE CONTROL: Compatible swap at matched injection depth ('cat'→'lion'); no-injection swap

ABLATION: Freeze cross-attention only vs cross+self; single-layer injection vs all layers

DOSE RESPONSE: cross_replace_steps $\in [0,1]$ in 0.1 steps $\times$ contradiction strength (semantically near/far pairs)

ITERATIVE VERSION: Re-invert the hybrid and swap again toward a third concept — contradiction layering

FALSIFICATION CONDITION: No stable hybrid regime: every contradiction pair collapses directly from swap-fails to layout-breaks

E3 — The QKV Transplant Matrix

HYPOTHESIS: Cross-image self-attention transplants (Q, K, V independently from donor A into host B) produce distinguishable, anatomically predictable hybrid classes, confirming and extending the documented structure/appearance channel division into compositional territory

REPRESENTATIONAL TENSION: Host's generative trajectory vs. donor's injected channel — one channel carries foreign structure or appearance while the rest of the network computes as if nothing happened

COMPETING CONSTRAINTS: Q/K (structure) vs V (appearance) jurisdictions per the attention survey^[13]

WHY THE SYSTEM MAY STRUGGLE: Donor features were computed for donor latents; in host context they are legal-shaped but contextually false

EXPECTED RESULT: Q/K transplants deform B's geometry toward A's layout; V transplants repaint B in A's appearance grammar; single-channel transplants (Q only) yield the least coherent, most novel artifacts

NEGATIVE CONTROL: Self-transplant A→A per channel (must reproduce A); noise-donor transplant

ABLATION: Injection layer subsets; injection step windows (FireFlow's first-step finding as anchor^[11])

DOSE RESPONSE: Injection step fraction $\times$ channel, per donor-host pair class (similar vs dissimilar content)

ITERATIVE VERSION: Hybrids become donors in round 2 — channel lineage tracking across generations

FALSIFICATION CONDITION: Transplant class has no consistent effect beyond generic corruption (compare against shuffled-feature donors)

E4 — Temporal Trajectory Splice

HYPOTHESIS: Splicing two prompts' denoising trajectories at timestep τ (first half A, second half B) yields chimeras whose dominance frontier is a smooth, reproducible function of τ — early structure decisions and late detail decisions can be assigned to different worlds deliberately

REPRESENTATIONAL TENSION: A's low-frequency structural commitments vs. B's high-frequency detail commitments, fused in one latent

COMPETING CONSTRAINTS: Trajectory continuity (the ODE wants a legal path) vs. the splice discontinuity

WHY THE SYSTEM MAY STRUGGLE: At τ , the latent is legal-shaped noise encoding A's intentions; the B-conditioned denoiser reads its own intentions into it

EXPECTED RESULT: τ sweep maps composition-vs-texture assignment: early splice = B's structure with A's textures; mid = true chimeras; late = A with B's surface. VAE repair marks the splice as specific artifact textures

NEGATIVE CONTROL: A/A self-splice (must reproduce A exactly — the identity test); pixel-space composite of finished A and B (shows what is NOT trajectory-level)

ABLATION: Splice with matched seed vs unmatched seeds; splice on latents vs splice on CFG branches

DOSE RESPONSE: $\tau \in \{10\%, \dots, 90\%\}$ of steps $\times$ prompt-pair semantic distance

ITERATIVE VERSION: Chain: output of splice k re-noised to τ and spliced with prompt C — iterated contamination with documented lineage

FALSIFICATION CONDITION: Splice artifacts match the pixel-composite control — nothing trajectory-specific exists

6 Cross-modal applications

Target medium	Transfer path	Status	Notes
Video	Spatio-temporal attention control documented (Video-P2P, FateZero, VideoGrain); RF-Edit demonstrated on video; temporal consistency = attention across frames	DOCUMENTED — Topic 10 will own this	Temporal failure surfaces are the richest variant; deferred
Audio	Audio diffusion models (Stable Audio, teticio-style latent audio diffusion) use the same U-Net/DiT attention + latent recipe; diffusers ships StableAudioAttnProcessor	PLAUSIBLE — Topic 06 will own this	Attention binding of text to spectrogram regions is terra incognita for art
LLM	Structural analogy to Topic 01: trajectory splice $\approx$ activation graft; attention injection $\approx$ steering; inversion debt $\approx$ steering residual	ANALOGY (METAPHOR label)	Use as design language only — mechanisms differ
GLSL / shaders	Noise-field composition and repair-by-iteration are native shader idioms; splice-timestep thinking transfers as time-stepped procedural composition	METAPHOR	Method transfer, not mechanism transfer
ComfyUI / controllers	Every mechanism here exposes node-level state (latents, maps, schedules) — direct fit for TOPOS-SRE-style external controllers reading artifact metrics and re-dosing	READY	E1 and E4 are designed for this integration

Table 3 — Cross-modal transfer assessment.

7 What appears most useful for AI SLOP

1. Inversion is the project's missing return path. AI SLOP's path-dependence principle ($A \rightarrow B \rightarrow C$ must differ from $A \rightarrow C$) is usually enforced procedurally, by controllers and archives. Inversion makes it physical: the trajectory through noise space IS a state with history, and its error (debt) is measurable with standard metrics. E1's Debt Dial is the cheapest strong experiment in this document — the eta-inversion harness already wires 8 inversion methods $\times$ 4 editors behind CLI flags^[5].

2. Attention injection is constraint governance made literal. The project's 'jurisdictions' metaphor becomes concrete: cross-attention maps are spatial jurisdictions over pixels, Q/K/V channels are functional jurisdictions over structure and appearance, and injection schedules are temporal jurisdictions over the denoising process. Binding War (E2) and the Transplant Matrix (E3) are jurisdictional conflicts with documented mechanics.

3. The 2026 working stack is settled enough to build on. diffusers attention processors (official API), RF-Inversion in diffusers, ComfyUI community ports, FireFlow-speed editing on FLUX — the mechanisms in this PDF run on current open infrastructure. The SD 1.5-era papers remain the clearest mechanistic documentation (U-Net attention is cleanly separated), so a dual-stack practice is recommended: SD 1.5 for mechanism studies, FLUX for production artifacts.

4. Failure-surface discipline transfers directly. The project's experimental method (controls, ablations, dose-response, artifact families) matches this literature's own norms — PIE-Bench categories and structure/background metrics can be reused as the validator layer for image experiments without

inventing new measurement machinery.

8 Weak / dubious ideas worth rejecting

Claim / idea	Verdict	Why
'Latent space is a smooth semantic space; linear interpolation gives meaning'	REJECT as stated	Intermediate points leave the trained manifold (non-Gaussian intermediates); interpolation artifacts are decoder failure, not latent semantics. Useful failure — wrong explanation
'Naive DDIM inversion under high CFG reconstructs real images'	REJECT	Error accumulation under CFG is documented across the whole inversion literature; it is the reason null-text, PnP, ReNoise, and RF methods exist
'Attention maps show where objects are'	OVER-SIMPLIFIED	Maps are inaccurate enough to cause leakage; treating them as ground truth masks reproduces the failure DPL was built to repair
'StyleDiffusion' as a single method	REJECT (provenance trap)	Two unrelated papers share the name (ICCV 2023 style transfer vs CVMJ 2024 embedding inversion). Cite with author+year or lose the thread
bloc97/CrossAttentionControl as living infrastructure	DEPRECATE	Frozen 2022-10; use google/prompt-to-prompt or diffusers processors instead
ComfyUI-SpliceTools / lrw-nodes as validated mechanisms	CAUTION	6–12 stars, no validation, one is GPL-3.0 (copyleft). Fine as raw material for self-run experiments; never cite as evidence
'More inversion steps always help'	REJECT	RF-Solver error can rise again past convergence (~25 steps); monotonicity is false ^[11]
HF: AswinMathew/latent-space-explorer	UNVERIFIED	Not resolvable via HF API in this session; excluded from all claims

Table 4 — Rejection register.

9 Open questions

- Does inversion debt have a consistent visual signature distinguishable from generic noise corruption? (E1 is designed to answer this; no published artifact taxonomy for under-corrected inversion was found.)
- How do P2P-class mechanisms re-derive on MMDiT/joint-attention architectures (SD3, FLUX)? Text and image tokens share attention — the 'cross-attention map' is no longer a separable object^[15].
- Is the Q/K/V structure–appearance division stable across architectures, or U-Net-specific? FireFlow's V-only success on FLUX hints it survives in shifted form — the Transplant Matrix would map this.
- What is the right metric for 'structured vs mushy' hybrid artifacts? PIE-Bench metrics measure editing success, not artifact interest — the project may need its own artifact-family probes.
- Can attention leakage radius be predicted from embedding-space semantic distance? If yes, contamination becomes composable; if no, it stays an empirical dial.

- Do trajectory splices produce anything that prompt blending cannot? (E4's falsification condition — if splice $\approx$ prompt mix, the mechanism is decorative, not structural.)

10 Sources / bibliography

- Hertz, A., Mokady, R., Tenenbaum, J., Aberman, K., Pritch, Y., & Cohen-Or, D. (2023). *Prompt-to-Prompt Image Editing with Cross Attention Control*. ICLR 2023. arXiv:2208.01626. Code: github.com/google/prompt-to-prompt. <https://arxiv.org/abs/2208.01626>
- Mokady, R., Hertz, A., Aberman, K., Pritch, Y., & Cohen-Or, D. (2023). *Null-text Inversion for Editing Real Images using Guided Diffusion Models*. CVPR 2023. arXiv:2211.09794. <https://arxiv.org/abs/2211.09794>
- Ju, X., Zeng, A., Bian, Y., Liu, S., & Xu, Q. (2024). *PnP Inversion: Boosting Diffusion-based Editing with 3 Lines of Code*. ICLR 2024. arXiv:2310.01506. Code+benchmark: github.com/cure-lab/PnPInversion. <https://arxiv.org/abs/2310.01506>
- Garibi, D., et al. (2024). *ReNoise: Real Image Inversion Through Iterative Noising*. ECCV 2024. arXiv:2403.14602. <https://arxiv.org/abs/2403.14602>
- Furiosa AI (2024). *Eta Inversion: Designing an Optimal Eta Function for Diffusion-based Real Image Editing*. ECCV 2024. Code: github.com/furiosa-ai/eta-inversion (8 inversion methods × 4 editors). <https://github.com/furiosa-ai/eta-inversion>
- Wang, Z., Zhao, L., & Xing, W. (2023). *StyleDiffusion: Controllable Disentangled Style Transfer via Diffusion Models*. ICCV 2023. arXiv:2308.07863. <https://arxiv.org/abs/2308.07863>
- InverseMeetInsert / GEO (2024). *Robust Real Image Editing via Geometric Accumulation Inversion* (documents DDIM cumulative error and CFG null-text instability). arXiv:2409.11734. <https://arxiv.org/abs/2409.11734>
- Wang, K., et al. (2023). *Dynamic Prompt Learning: Addressing Cross-Attention Leakage for Text-Based Image Editing*. NeurIPS 2023. Code: github.com/wangkai930418/DPL. <https://github.com/wangkai930418/DPL>
- Rout, L., Chen, Y., Ruiz, N., Caramanis, C., Shakkottai, S., & Chu, W.-S. (2025). *Semantic Image Inversion and Editing using Rectified Stochastic Differential Equations (RF-Inversion)*. ICLR 2025. arXiv:2410.10792. Code: github.com/LituRout/RF-Inversion (diffusers + ComfyUI ports). <https://arxiv.org/abs/2410.10792>
- Wang, J., et al. (2025). *Taming Rectified Flow for Inversion and Editing (RF-Solver / RF-Edit)*. ICML 2025. arXiv:2411.04746. <https://arxiv.org/abs/2411.04746>
- He, X., Mei, C., Wang, P., & Tang, F. (2024). *FireFlow: Fast Inversion of Rectified Flow for Image Semantic Editing*. arXiv:2412.07517. <https://arxiv.org/abs/2412.07517>
- Emergent Mind (2025). *DDIM Inversion: Methods & Applications* — synthesis on error accumulation, early-step concentration, non-Gaussian structure of recovered latents (Staniszewski et al. 2024; Wallace et al. EDICT 2023; BDIA 2023). <https://www.emergentmind.com/topics/denoising-diffusion-implicit-models-ddim-inversion>
- Sun, K., et al. (2025). *Attention in Diffusion Model: A Survey (Q/K/V role division; injection taxonomies across image and video)*. arXiv:2504.03738. <https://arxiv.org/abs/2504.03738>
- Hugging Face (2023–). *Diffusers: Attention Processor API* (pluggable set_attn_processor extension point). huggingface.co/docs/diffusers/en/api/attnprocessor. <https://huggingface.co/docs/diffusers/en/api/attnprocessor>
- Paul, S. (2025). *Flavors of attention in modern diffusion models* (cross-attention vs joint-attention/MMDiT structural differences). sayak.dev. <https://sayak.dev/posts/attn-diffusion.html>
- Avrahami, O., Fried, O., & Lischinski, D. (2023). *Blended Latent Diffusion*. SIGGRAPH 2023. arXiv:2206.02779. Code: github.com/omriav/blended-latent-diffusion. <https://arxiv.org/abs/2206.02779>
- Kiwhan, S. (2024). *Correcting Classifier-Free Guidance* (mode-shifting and variance analysis of CFG). kiwhan.dev. <https://kiwhan.dev/blog/2024/classifier-free-guidance/>
- Tumanyan, N., Geyer, M., Bagon, S., & Dekel, T. (2023). *Plug-and-Play Diffusion Features for Text-Driven Image-to-Image Translation*. CVPR 2023. arXiv:2211.12572. <https://arxiv.org/abs/2211.12572>
- Cao, M., et al. (2023). *MasaCtrl: Tuning-free Mutual Self-Attention Control for Consistent Image Synthesis and Editing*. ICCV 2023. arXiv:2304.08465. <https://arxiv.org/abs/2304.08465>

Meng, C., et al. (2022). *SDEdit: Guided Image Synthesis and Editing with Stochastic Differential Equations*. ICLR 2022. arXiv:2108.01073. <https://arxiv.org/abs/2108.01073>

All GitHub metadata verified via the GitHub API on 2026-09-19; arXiv IDs cross-checked against search indexes. Items marked UNVERIFIED in Table 4 are excluded from mechanism claims.

11 Licensing / provenance appendix

Asset	License / terms	Verified	Provenance note
google/prompt-to-prompt	Apache-2.0 (see repo)	2026-09-19	Official P2P code; Imagen + SD notebooks
bloc97/CrossAttentionControl	MIT	2026-09-19 (GitHub API)	Unofficial P2P; frozen 2022-10 — reference only
cure-lab/PnPInversion + PIE-Bench	See repo (research code)	2026-09-19 (GitHub API)	Official ICLR 2024 code; frozen 2024-03 but complete; benchmark included
furiosa-ai/eta-inversion	See repo	2026-09-19 (GitHub API)	Unified inversion/editor harness — recommended experiment bench
wangkai930418/DPL	See repo	2026-09-19 (GitHub API)	Leakage diagnosis + repair; small but complete
sen-mao/StyleDiffusion	See repo	2026-09-19 (GitHub API)	CVMJ 2024 embedding inversion — NOT the ICCV 2023 paper; active 2026-08
LituRout/RF-Inversion	See repo; diffusers integration Apache-2.0	2026-09-19	ICLR 2025; official + community ComfyUI port
AustinMroz/ComfyUI-SpliceTools	GPL-3.0 (copyleft)	2026-09-19 (GitHub API)	6 stars, experimental; GPL affects any distributed derivative workflow
lajjadred/comfyui-lrw-nodes	Check repo (not listed)	2026-06 (GitHub API)	2026-05 creation; unvalidated Riemannian/Bayesian latent ops
huggingface/diffusers	Apache-2.0	2026-09-19	Attention processor API is the supported intervention surface
HF spaces (anshuln, AswinMathew leads)	—	PARTIAL / UNVERIFIED	anshuln attention-viz space exists per web index (slug unconfirmed); latent-space-explorer unverified

Table 5 — Licensing and provenance.

Research conducted 2026-09-19 via web sources, arXiv, and the GitHub API. The ai_slop repository was read for connection mapping only — nothing was modified, added, or committed. ESTABLISHED claims trace to peer-reviewed venues or benchmarked code; PLAUSIBLE INTERPRETATION, PROJECT HYPOTHESIS, and SPECULATION are marked inline. The GPL-3.0 license on SpliceTools is flagged because node-pack provenance matters if derivative workflows are ever distributed.