



Inspect the whole GitHub repository (or the ai specific part for u) and expand upon it - search everywhere for more ideas like that - https://github.com/merrypranxter/ai_slop.git - don't change anything in the repo it is just for you to see. You may search GitHub, academic stuff, the web anything

I inspected the repository tree and concentrated on the canonical AI material, machine-readable registries, semantic controllers, Temporary Minds, experimental method, media transduction documents, backlog, and representative source-lineage files. I made **no changes** to the repository; the strongest expansion is not “more prompts,” but turning this unusually coherent conceptual library into an **executable ecology of operators, experiments, archives, and generative artworks.**

AI SLOP Expansion Map

What it already is

AI SLOP is better understood as a laboratory for **structured cognitive mutation** than as a prompt collection. Its central cycle is:

ordinary assumption → precise intervention → competing pressures → attempted resolution → artifact → observation → test → preservation or mutation

The canonical system currently contains 26 Temporary Minds, a broader cross-media operator registry, several external-state semantic controllers, media-specific translation methods, and an unusually good distinction between **creative usefulness** and **confidence in the explanation of why something worked.**

Four layers recur throughout the repository:

- **Conceptual operators:** Mutations of identity, distance, attention, memory, value, causality, categories, completion, and selection.
- **Media transduction:** Converting those mutations into actual image, video, audio, material, typography, topology, rhythm, or shader behavior.
- **Stateful systems:** Maintaining trajectories, scars, symbol tables, attractors, retired operators, and mutation histories outside the model.

- **Mechanism research:** Controlled investigation of binding failures, conditioning conflicts, compression, positional effects, reference pressure, and actual local-model interventions.

The machine-readable layer already makes this more than literary theory.

[machine/operators.json](#) can become the seed of a runtime registry rather than merely an index.

Strongest qualities

Area	What works
Operator design	Most entries specify an identifiable transformation rather than “be surreal” or “think differently.”
Epistemic hygiene	Prompt procedures are explicitly separated from claims about changing weights, embeddings, or hidden states.
Negative-space methods	Deletion, omission, taboo, lesion, and aftermath are treated as productive operations rather than decorative constraints.
Path dependence	Recall scars, semantic recoil, fossils, trajectories, and observer effects provide a real temporal dimension.
Cross-media thinking	The repository generally translates causal behavior rather than importing scientific vocabulary as visual costume.
Validation	Rick’s Veto, no-metaphor leakage, structural ancestry, observer scars, ablations, and dose-response testing are excellent foundations.
Machine readability	Operators and protocols already have identifiers, classes, summaries, and canonical sources.

The repository’s experimental method is particularly valuable: baseline, isolated variable, repeat runs, strength sweeps, ablation, failure recording, and separate mechanism-confidence and creative-utility scores. That discipline could distinguish this project from the enormous quantity of untestable “LLM consciousness hacking” material online.

Main missing layer

The largest gap is **connective tissue**. The repository describes operators, validators, controllers, ledgers, compilers, and state transitions, but they do not yet form one runnable environment in which a user can:

1. Install an operator.
2. Generate several candidates.
3. Observe what structurally changed.
4. Compare against a baseline.
5. Place the result in a behavioral archive.
6. Breed or mutate it.
7. Compile it into another medium.
8. Preserve its ancestry and scars.

This matters because computational creativity normally separates generation from evaluation, while conceptual blending adds explicit input spaces, correspondence, selective projection,

and emergent blend structure ^[1] ^[2]. AI SLOP already contains richer mutation procedures than many computational-creativity systems; what it lacks is a consistent runtime for executing and comparing them.

The Operator Workbench

The most natural major project is an **Operator Ecology Workbench**: a browser application where every operator is executable, inspectable, breedable, and testable.

Runtime objects

Each experiment would manipulate these objects:

- **Task:** The human's actual requirement.
- **State:** Current ontology, constraints, symbols, media variables, and unresolved conflicts.
- **Operator:** A transition procedure.
- **Invariant:** What must survive the transition.
- **Candidate family:** Multiple results, not one lucky artifact.
- **Behavior descriptor:** Measurable properties used to classify the result.
- **Scar:** A state change that persists after its initiating condition disappears.
- **Archive cell:** A niche containing the strongest artifact found for one behavioral profile.
- **Lineage:** Parent operators, artifacts, mutations, and selection events.
- **Compiler:** A translator from conceptual state into text, image, GLSL, audio, or video instructions.

Execution cycle

1. Generate an ordinary baseline.
2. Select one operator and explicit strength.
3. Produce several candidates.
4. Run hard validators before aesthetic ranking.
5. Extract behavioral descriptors.
6. Place candidates into an archive.
7. Let the user preserve, reject, breed, or mutate them.
8. Compile selected states into a target medium.
9. Remove the operator and test whether claimed scars persist.
10. Save the complete lineage as a reproducible experiment.

This would make “operator burnout,” “conceptual allergy,” and “afterlife” computationally real. Novelty search rewards behavioral difference rather than resemblance to an ultimate target, specifically to avoid deceptive objective landscapes ^[3] ^[4]. MAP-Elites extends that idea by

maintaining many high-performing but qualitatively different solutions across user-selected behavioral dimensions instead of returning one winner [5].

Behavioral archives

The repository’s regulators map almost perfectly onto quality-diversity algorithms.

AI SLOP concept	Computational implementation
Conceptual allergy	Penalize proximity to occupied archive cells.
Cliché mortality	Maintain a blacklist of common structural signatures, not words.
Operator burnout	Remove successful operators from the active mutation pool.
Minority-axis sovereignty	Preserve rare but legitimate descriptor combinations.
Descendant fitness	Evaluate a parent by the diversity and quality of later offspring.
Meta-genomic speciation	Mutate operator definitions, not merely their outputs.
Observer scar	Compare state before and after evaluation or inspection.
Generative leverage	Divide structured downstream change by intervention complexity.
Alien distance	Swap the archive’s behavior descriptors during a run.
Afterlife module	Mark previously productive regions as exhausted.

Possible archive axes include:

- Recognizability versus ontological mutation.
- Causal propagation depth.
- Task survival.
- Number of unresolved jurisdictions.
- Degree of literal versus metaphorical transformation.
- Temporal persistence.
- Descendant diversity.
- Human surprise versus model-judge surprise.
- Surface novelty versus structural novelty.
- Compression level versus recoverability.
- Semantic recoil magnitude.
- Dependency-graph disruption.

A particularly appropriate interface would show a glowing two-dimensional field of archive cells. Each cell contains an artifact, but selecting it reveals the full operator lineage, prompt compilation, state diff, model/version, seed, validators, and later descendants.

Expanded operator schema

The existing JSON registry could support execution by adding fields like these:

```
{
  "id": "primitive_deletion",
  "class": "ontology",
  "scope": ["text", "image", "audio", "controller"],
  "inputs": ["task", "state", "primitive"],
  "preconditions": ["primitive_is_load_bearing"],
  "hard_constraints": ["no_synonym_substitution"],
  "state_effects": ["remove_node", "invalidate_dependents"],
  "invariants": ["task_requirement"],
  "escape_routes": ["renaming", "metaphorical_absence"],
  "behavior_descriptors": ["dependency_loss", "reconstruction_depth"],
  "falsification": "Could the output be unchanged if the primitive were restored?",
  "compatible_operators": [],
  "conflicting_operators": [],
  "media_adapters": {},
  "benchmark_cases": [],
  "parents": [],
  "version": "0.3"
}
```

The important addition is a distinction between **instruction text** and **operator semantics**. A Temporary Mind should have one abstract transition definition and multiple execution adapters: direct prompting, multi-agent execution, rule-based state mutation, shader compilation, music specification, and local-model intervention.

New Temporary Minds

The repository’s idea garden identifies missing territories such as social cognition, mathematics, humor, beauty, salience, completion, affordance, uncertainty, scale, self-modeling, and irreducible value conflict. These are candidate procedures designed to fill those gaps without duplicating the existing 26.

Candidate	Core operation	Decisive test
Norm Migration	Hold individual goals stable while progressively changing the rules by which agents recognize promises, ownership, consent, or authority.	Do outcomes change because coordination rules changed, rather than because an agent acquired a new personality?
Proof-Object Reciprocity	Treat every constructed object as a proof obligation and every valid proof as an object with manipulable geometry or behavior.	Does changing proof structure force a corresponding change in the artifact?
Premise Debt	Every local explanation incurs an assumption debt that must later be repaid through an unexpected but valid consequence.	Is the final reversal structurally required by earlier debt, or merely a joke added afterward?
Aesthetic Conservation	Beauty cannot be removed, only transferred between pattern, failure, material, timing, absence, or observer effort.	Can the location and form of aesthetic value be tracked across the transformation?

Candidate	Core operation	Decisive test
Irreversibility Salience	Rank details by the future options they permanently destroy rather than by immediate visual or narrative prominence.	Does the system notice quiet irreversible transitions before spectacular reversible ones?
Fertile Completion	A work is complete when it can generate several irreducible descendants, not when its immediate form feels resolved.	Does the completion rule choose different stopping points than ordinary polish criteria?
Negative Affordance	Define entities by actions they prevent, delay, redirect, or make socially unthinkable.	Does removing the prevention structure alter the entity's identity?
Belief Parliament	Maintain incompatible hypotheses with separate evidence ledgers and grant each limited action rights without averaging them into one probability.	Do minority hypotheses affect action before winning consensus?
Capability Shadow	Construct the system's self-model only from repeatable failure boundaries, blind spots, and dependence on external tools.	Does the self-model predict failures rather than merely narrate a persona?
Scale Jurisdiction	Give different spatial or temporal scales incompatible laws and define translation only at their boundaries.	Does changing observation scale alter the governing operation rather than merely reveal more detail?
Pareto Veto Mind	Treat important values as non-convertible jurisdictions; no single score may compensate for violating another.	Does the system preserve genuine conflict instead of manufacturing a weighted compromise?
Counterfactual Custody	Every discarded possibility must leave one constraint on the selected reality.	Would the final artifact differ if rejected branches had never existed?
Audience Refraction	Generate several observer models whose interpretations modify different artifact variables, without selecting a privileged audience.	Do multiple audiences causally deform the work rather than receive tailored descriptions?
Unknown-Unknown Transducer	Detect repeated residuals that none of the current categories explains, then protect the residual from immediate naming.	Does the residual influence later selection before being classified?

The last operator connects directly to Strategic Ambiguity: the anomaly must produce consequences before explanation neutralizes it.

Operator breeding

The repository's Structural Ancestry Requirement is strong enough to support an actual **operator genetics system**. Conceptual blending systems distinguish input spaces, shared structure, selective projection, completion, and elaboration; importantly, the blend may contain emergent relations that appear in neither input ^[1] [6].

A breeding engine could therefore require:

- Two parent operators with different intervention targets.
- A declared shared invariant.
- One rule inherited from each parent.
- One incompatibility that offspring must resolve.

- An emergent transition rule not reducible to sequential application.
- An ablation showing that removing either parent destroys the offspring's characteristic behavior.

Example offspring:

Fossilized Taboo

Parents: Procedural Taboo + Recall Mutation.

A forbidden direct operation cannot be performed, but each failed attempt mutates memory. Eventually the accumulated memory scars create an indirect path that did not exist at the beginning.

Migratory Sensorium

Parents: Synthetic Transducer + Property Unbundling.

A detector's sensory ownership moves among entities. Each new host experiences the same signal differently because its available actions differ.

Crystalline Recoil

Parents: Observational Crystallization + Semantic Recoil.

Attention fixes an interpretation temporarily, but later evidence can alter the rule that made the fixation possible. The old interpretation remains physically present as a fossil rather than being overwritten.

Jurisdictional Organism

Parents: Separated Jurisdictions + Kinetic Dissolution.

Different metabolic flows obey different legal systems. The "organism" exists only as negotiations among incompatible maintenance regimes.

Semantic visualizer

The GLSL visualizer should become a visualization of **external semantic state**, not a fake representation of neural activations.

A useful mapping would be:

- State dimensions → field channels.
- Operator strength → displacement amplitude.
- Contradiction → curl or interference.
- Uncertainty → diffusion.
- Salience → local brightness or growth rate.

- Retired concepts → sediment layers.
- Semantic recoil → backward-propagating waves.
- Observer actions → local crystallization.
- Archive density → viscosity.
- Unresolved jurisdictions → overlapping coordinate systems.
- Scars → persistent defects that survive parameter reset.

The artistic leap is to let shader state affect the semantic controller in return. Dense regions might reduce available language, vortices might change conceptual distance, and crystallized areas might become temporarily immutable.

Artificial-life direction

The operator system is highly compatible with artificial life because it already treats ideas as processes with mutation, inheritance, niches, scars, and speciation. Lenia is a continuous cellular-automaton family known for producing diverse, resilient, moving autonomous patterns; its original survey reported more than 400 identified forms across 18 families ^[7]. Flow-Lenia adds mass conservation, multispecies behavior, and emergent evolutionary dynamics ^[8].

A project called **Semantic Lenia** could assign each conceptual operator:

- A convolution kernel describing what state regions it can influence.
- A growth function describing conditions under which it strengthens or decays.
- A conserved resource such as attention, descriptive bandwidth, or interpretive certainty.
- A mutation rule.
- A compatibility matrix with other operators.
- A phenotype compiled into text, sound, or imagery.

This should not pretend that semantic concepts literally behave like cells. It would be an external dynamical artwork whose state is periodically interpreted by an LLM.

Useful implementation references include [HexLife](#), which combines browser-based cellular automata, multiple worlds, WebGL rendering, and evolutionary exploration; [Petri](#), which breeds multiple Lenia universes using MAP-Elites; and [Primordia](#), which combines Lenia, Physarum, particle life, and Gray-Scott systems in a GPU artificial-life environment. These are implementation references rather than evidence that their conceptual frameworks match AI SLOP.

Artwork families

Museum of Retired Metrics

Every room contains artifacts produced under a different alien distance metric. Once a metric yields a successful family, it is permanently retired, and the room becomes inaccessible except through its descendants.

The Absent Orchestra

Omission-Field Cognition reconstructs an unnamed musical object through rhythm, spectral holes, performer behavior, room response, and the reactions of other instruments. No channel may directly state the missing object.

Error Zoology

A classifier's coherent near-errors become literal species definitions. Each species receives a metabolism, morphology, habitat, failure behavior, and evolutionary history consistent with the misclassification.

Research on AI art has found that artists actively use model inconsistencies, imprecise representations, overinterpretation, and technological failures to create ambiguity rather than treating them only as defects ^[9]. This supports the repository's Structural-Error Sieve: preserve coherent errors, not indiscriminate corruption.

Codec Organs

Compression blocks, chroma subsampling, ringing, quantization, motion vectors, and packet loss become anatomical functions. A body does not receive a JPEG texture; it grows, remembers, heals, and decomposes according to codec behavior.

Algorithmic-information approaches have explicitly connected novelty generation to lossy compression and described generative systems as analogous to decompressors ^[10]. Glitch theory likewise treats file formats and codec errors as material agencies rather than arbitrary visual overlays ^{[11] [12]}.

Semantic Cemetery

Deleted concepts remain as gravestones whose positions alter navigation. Removing an idea therefore reduces descriptive clutter but changes the shape of every later route.

Noncommutative Film Lab

Apply transformation A and then B, followed by B and then A. The film consists of the growing difference between those histories, especially after occlusion, reappearance, identity reassignment, or temporal compression.

The Audience Is Weather

Several simulated audiences observe the same evolving piece. Their disagreement becomes wind, erosion, phase cancellation, or mutation pressure, but no audience directly controls the piece.

SLOP Radio

A semantic controller broadcasts its state as generative sound. Conflict becomes polyrhythmic jurisdiction, certainty narrows tuning, forgotten concepts survive as weak carrier frequencies, and operator burnout permanently removes an instrument-generation rule.

Quasicrystal Cognition

Temporary Minds are arranged according to an aperiodic tiling rather than a sequence. Neighboring tiles determine which operators may combine, while local patterns recur without global repetition.

This offers a strong connection to your interest in quasicrystals: the installation could preserve local recognizability while making global operator recurrence impossible to reduce to a loop.

Browser of Counterfactual Fossils

Every abandoned branch leaves a small executable residue. Users can click the residue to reveal how the current artifact would collapse if that rejected possibility had never constrained it.

The Work That Hates Being Named

An anomaly begins as an unlabelled behavior. Each successful label causes the anomaly to migrate into another medium or state dimension, operationalizing Strategic Ambiguity as an adversarial relationship with classification.

Real model interventions

The repository correctly separates prompt-level procedural mutation from actual intervention in open models. That boundary should remain absolute.

A local-model laboratory could support:

Intervention	Artistic use	Required control
Activation addition	Apply a behavior or concept direction during generation.	Layer and coefficient sweep; off-target task checks.
Sparse-feature steering	Amplify or suppress features discovered by sparse autoencoders.	Feature interpretation, reconstruction loss, and collateral-feature monitoring.
Activation patching	Replace one internal activation with an activation from another run.	Position, layer, and source-run ablations.

Intervention	Artistic use	Required control
Logit modification	Restrict or bias lexical possibilities.	Check whether semantic behavior survives through paraphrase.
Targeted model editing	Change a factual association or narrow behavior.	Locality, specificity, generalization, and unrelated-fact tests.
Hook-based lesions	Zero, rotate, swap, quantize, or noise selected intermediate states.	Compare against random vectors and equal-magnitude controls.

Activation Addition constructs steering vectors from differences between prompt-induced activations and injects them during inference; published results show that this can alter high-level output properties while often preserving unrelated performance ^[13]. ROME instead edits feed-forward weights associated with factual recall and evaluates whether the change remains specific while generalizing to relevant contexts ^[14].

Sparse-space steering is especially relevant to Synthetic Transducers and Selective Cognitive Lesions because raw activation directions are often entangled. Current research explores steering in sparse feature spaces precisely to obtain more interpretable and targeted interventions ^[15] ^[16].

Practical code references include:

- [Mihaiiii/llm_steer](#) for lightweight activation engineering.
- [DLR-SC/style-vectors-for-steering-llms](#) for style-vector experiments.
- [dmbeaglehole/neural_controllers](#) for concept-vector monitoring and steering.
- [dthinkr/mlx-lens](#) for steering and sparse-autoencoder analysis on Apple Silicon.
- [corl-team/steering-reasoning](#) for newer work connecting small steering interventions with reasoning behavior.

These should become an explicitly separate **white-box** track inside the workbench. Prompt-only runs, API-mediated generations, and actual hidden-state interventions should never share the same mechanism labels.

A metacreative controller

The support architecture can become a real **metacreative controller** that examines its own search process rather than only scoring outputs. Computational metacreativity research defines this in terms of self-reflection and self-control: the system observes and adjusts its creative standards or process ^[17].

The controller could periodically ask:

- Which operators dominate the archive?
- Which descriptors have become easy to optimize?
- Which apparent novelties are surface variations of the same skeleton?
- Which rejected artifacts have unusually fertile descendants?
- Which validators reject nearly everything?

- Which operators only work with one model?
- Which human selections contradict automatic scores?
- Has an alien metric become an ordinary metric through repeated use?
- Is the archive expanding, or merely polishing occupied niches?
- Should the system mutate an operator, a validator, a distance metric, or the task representation?

This would implement Meta-Genomic Speciation at the appropriate level: not “the AI became creative,” but “the external search controller modified which creative procedures are available and how their success is measured.”

Evaluation expansion

The existing 0–3 benchmark dimensions should remain, but each run could also record:

- **Causal delta:** Difference from baseline attributable to the operator.
- **Task survival:** Percentage of hard task requirements retained.
- **Operator identifiability:** Whether blinded evaluators can infer the installed operator.
- **Literal consequence depth:** Number of downstream decisions changed by the mutation.
- **Archive novelty:** Distance from previous behavioral signatures.
- **Ancestry irreducibility:** Performance loss when either parent is ablated.
- **Scar persistence:** Whether a change survives removal of its initiating condition.
- **Order sensitivity:** Difference between $A \rightarrow B$ and $B \rightarrow A$.
- **Cross-model portability:** Whether the recognizable mutation recurs across models.
- **Compiler fidelity:** Whether an abstract state produces corresponding behavior in different media.
- **Evaluator disagreement:** Human/model/rule-based ranking divergence.
- **Descendant fertility:** Number of structurally distinct valid offspring within a fixed generation budget.

A first serious benchmark could test all 26 Temporary Minds across the repository’s five standard tasks, three models, and five runs per condition. That creates 1,950 operator runs plus 75 ordinary baselines, while remaining small enough to analyze as families rather than isolated screenshots.

Research questions

The richest questions are not “does this prompt make the model weirder?” They are:

- Which operators create recognizable cross-task signatures?
- Which pairs produce irreducible offspring rather than stacked effects?
- Which minds merely rename ordinary brainstorming procedures?
- Which models collapse separated jurisdictions into stylistic blending?

- Can operator identity be detected from outputs without seeing prompts?
- Does operator success transfer from text to image or audio compilation?
- Do human evaluators and embedding-based novelty measures identify the same clichés?
- Can archive pressure produce genuinely new structural families, or does it optimize superficial difference?
- Which scars persist after their originating instruction disappears?
- Does model-generated self-evaluation domesticate the anomalies it is asked to inspect?
- Can a human detect more valuable stepping stones than a fixed novelty metric?

Interactive evolutionary work has already combined human selection with novelty search to discover stepping stones that objective optimization misses ^[18]. That makes a mixed-initiative model especially appropriate here: the software manages archives and reproducibility, while the artist protects strange candidates that automatic evaluators would eliminate.

First build sequence

Phase A: Registry compiler

Convert operators, validators, Temporary Minds, protocols, and media adapters into one validated schema. Add automatic checks for missing falsification tests, hidden metaphor escapes, undocumented state effects, and duplicate functional signatures.

Phase B: Text laboratory

Build a local browser interface with baseline generation, operator installation, candidate families, side-by-side state diffs, scoring, lab notes, and exportable experiment bundles. Begin with five sharply distinct operators rather than all 37: Primitive Deletion, Alien Distance, Synthetic Transducer, Recall Mutation, and Procedural Taboo.

Phase C: Archive field

Add embeddings and hand-designed behavioral descriptors, then place valid artifacts in a MAP-Elites-style grid. Let the user replace automatic descriptors with alien metrics during a run.

Phase D: GLSL organism

Feed controller state into a WebGL or WebGPU simulation. Preserve deleted concepts as sediment, encode contradictions as interference, and allow visual events to perturb semantic state.

Phase E: Breeding chamber

Implement parent selection, ancestry requirements, ablation, offspring naming, and descendant-fitness testing. Reject hybrids that are equivalent to running parent A and then parent B.

Phase F: White-box lab

Add open-model hooks, activation capture, steering vectors, sparse-feature analysis, and controlled lesions. Keep all white-box results in a distinct evidence namespace from prompt-level procedures.

The most immediately distinctive project would be **the browser-based Operator Ecology Workbench plus Semantic Cemetery visualizer**. It uses the strongest existing material, fits JavaScript/GLSL practice, produces both research data and strange visual artifacts, and gives every future idea somewhere concrete to live.

✱✱

1. https://computationalcreativity.net/iccc21/wp-content/uploads/2021/09/ICCC_2021_paper_122.pdf
2. https://en.wikipedia.org/wiki/Computational_creativity
3. https://link.springer.com/chapter/10.1007/978-1-4614-1770-5_3
4. https://www.cs.swarthmore.edu/~meeden/DevelopmentalRobotics/lehman_ecj11.pdf
5. <https://arxiv.org/abs/1504.04909>
6. <https://www.iiia.csic.es/~enric/papers/UniformModelBlending.pdf>
7. <https://arxiv.org/abs/1812.05433>
8. <https://arxiv.org/abs/2506.08569>
9. <https://arxiv.org/html/2403.09374v1>
10. <https://computationalcreativity.net/ijwcc07/papers/mcgregor.pdf>
11. <https://www.tandfonline.com/doi/full/10.1080/17514517.2017.1295711>
12. <https://attachments.are.na/39014791/efab99a7098f41c12cca601ef118896a.pdf?1756116572>
13. <https://arxiv.org/abs/2308.10248>
14. <https://arxiv.org/abs/2202.05262>
15. <https://arxiv.org/abs/2503.00177>
16. <https://openreview.net/pdf/70f3af98996906bfa9bd84ba645f0cf45176844a.pdf>
17. https://computationalcreativity.net/iccc2017/ICCC_17_accepted_submissions/ICCC-17_paper_25.pdf
18. <https://arxiv.org/pdf/1207.6682.pdf>
19. <https://arxiv.org/html/2604.06339v1>
20. https://openaccess.thecvf.com/content/CVPR2025/html/Sun_T2V-CompBench_A_Comprehensive_Benchmark_for_Compositional_Text-to-video_Generation_CVPR_2025_paper.html
21. <https://openreview.net/pdf?id=uUofPYVMZH>
22. <https://worldmodelbenchmark.com/>

23. <https://arxiv.org/abs/2602.10113>
24. <https://openreview.net/pdf/891c726266f6938a57cd7132081734caab9a0703.pdf>
25. <https://jirt.org/index.php/JIRT/article/view/47>
26. <https://digitalmindnews.com/ai/ai-video-generation-models-capabilities-limits/>
27. <https://escholarship.org/uc/item/33g562jt>
28. https://bmva-archive.org.uk/bmvc/2025/assets/papers/Paper_159/paper.pdf
29. <https://www.emergentmind.com/topics/wowbench>
30. <http://arxiv.org/pdf/2310.19842.pdf>
31. <http://arxiv.org/pdf/2408.15176.pdf>
32. <https://arxiv.org/pdf/2211.09124.pdf>
33. <https://www.frontiersin.org/articles/10.3389/frai.2020.00014/pdf>
34. <https://arxiv.org/html/2407.10462v1>
35. <https://downloads.hindawi.com/journals/cin/2022/5083146.pdf>
36. <https://pmc.ncbi.nlm.nih.gov/articles/PMC7861321/>
37. <https://arxiv.org/pdf/2110.12855.pdf>
38. <https://arxiv.org/html/2606.09780>
39. <https://www.frontiersin.org/journals/computer-science/articles/10.3389/fcomp.2025.1543074/full>
40. <https://arxiv.org/html/2406.05873v1>
41. https://human-competitive.org/sites/default/files/paper_0.pdf
42. <https://periodicos.unespar.edu.br/vortex/article/download/4690/3218/18978>
43. <https://www.ijcai.org/proceedings/2024/0858.pdf>
44. <https://www.openbookpublishers.com/books/10.11647/obp.0301/chapters/10.11647/obp.0301.06>
45. https://cris.unibo.it/bitstream/11585/952310/3/camReady_preEditoriale.pdf
46. <https://arxiv.org/html/2506.05104v2>
47. <https://www.frontiersin.org/journals/robotics-and-ai/articles/10.3389/frobt.2021.680586/full>
48. <https://pdfs.semanticscholar.org/ef0e/dbd25fe006a4355a4de52467d3d928c0a3f1.pdf>
49. https://www.academia.edu/77508702/Quality_Diversity_for_Aesthetic_Evolution
50. <http://www.cmap.polytechnique.fr/~nikolaus.hansen/proceedings/2011/GECCO/proceedings/p419.pdf>
51. https://cris.unibo.it/retrieve/handle/11585/952310/1fdff8c7-381c-4bcc-a1e8-9bd05552c773/camReady_preEditoriale.pdf
52. https://edoc.ub.uni-muenchen.de/18040/1/Meier_Max.pdf
53. <https://arxiv.org/html/2503.17671v1>
54. <http://arxiv.org/pdf/2310.04378.pdf>
55. <https://arxiv.org/html/2411.14193>
56. <http://arxiv.org/pdf/2409.01392v1.pdf>
57. <https://arxiv.org/html/2410.07171>
58. <https://arxiv.org/html/2412.00177v2>

59. <http://arxiv.org/pdf/2303.11916.pdf>
60. <https://docs.comfy.org/built-in-nodes/LatentBlend>
61. <https://github.com/google/prompt-to-prompt/>
62. <https://github.com/icaros-usc/pyribs>
63. <https://github.com/ttulttul/ComfyUI-Iterative-Mixer>
64. <https://github.com/janeyeon/prompt-to-prompt-lcm>
65. <https://quality-diversity.github.io/libraries.html>
66. <https://prompt-to-prompt.github.io/>
67. <https://comfy.icu/node/FeedbackSampler>
68. <https://blenderneko.github.io/ComfyUI-docs/Core Nodes/Sampling/KSampler/>
69. https://yuanzhi-zhu.github.io/documents/slides/prompt_to_prompt.pdf
70. <https://github.com/bloc97/CrossAttentionControl>
71. <https://github.com/miguelCalado/prompt-to-prompt-tensorflow>
72. <https://github.com/Comfy-Org/ComfyUI>
73. <https://lunaleee.github.io/posts/prompt2prompt/>
74. <https://linn-2021.github.io/prompt-to-prompt-stable-diffusion/>
75. <https://www.cambridge.org/core/services/aop-cambridge-core/content/view/253B777A512E28573AACEAC6977EDCB4/S2326376823000256a.pdf/div-class-title-artificial-intelligence-and-archaeological-illustration-div.pdf>
76. <https://arxiv.org/pdf/2405.01901.pdf>
77. <https://arxiv.org/html/2503.13531v1>
78. <https://arxiv.org/pdf/2312.08056.pdf>
79. <https://arxiv.org/html/2210.06145>
80. <https://arxiv.org/pdf/2406.08739.pdf>
81. <https://arxiv.org/html/2412.01450>
82. <https://www.mdpi.com/2571-9408/7/11/295>
83. <https://grayarea.org/press/trillionth-webpage-net-art-commissions-internet-archive/>
84. <https://arxiv.org/pdf/2602.05710.pdf>
85. <https://www.frontiersin.org/journals/communication/articles/10.3389/fcomm.2026.1826494/pdf>
86. https://pure.uva.nl/ws/files/329503123/From_Slop_to_Slop_M_C_Journal.pdf
87. <https://www.musikundmedien.hu-berlin.de/de/medienwissenschaft/medientheorien/ernst-in-english/SCRIPTS/PDF/medarch-ready.pdf/@@download/file/MEDARCH-READY.pdf>
88. <https://www.tandfonline.com/doi/full/10.1080/17540763.2025.2607745>
89. https://rodin.uca.es/bitstream/handle/10498/36475/OA_2025_0236.pdf?sequence=1&isAllowed=y
90. https://www.academia.edu/144668877/Theory_as_Poor_Image_Hito_Steyerls_Text_in_Algorithmic_Circulation
91. <https://ghuneim.us/blog/the-perfect-crime-when-ai-completed-photographys-murder-of-reality/>
92. <https://journals.zycentre.com/cvca/article/download/98/70/1161>
93. <https://www.scribd.com/document/912404230/In-Defense-of-the-Poor-Image>

94. <https://aestheticinvestigations.eu/article/download/12085/15194/29333>
95. <https://ucrarts.ucr.edu/wp-content/uploads/2024/09/digital-capture.pdf>
96. <https://artreview.com/summer-2014-feature-hito-steyerl/>
97. https://spectrum.library.concordia.ca/id/eprint/993614/1/Cavdar_MA_S2024.pdf
98. <https://arxiv.org/html/2410.12877v1>
99. <https://arxiv.org/pdf/2410.12299.pdf>
100. <http://arxiv.org/pdf/2409.14026.pdf>
101. <http://arxiv.org/pdf/2403.05767.pdf>
102. <http://arxiv.org/pdf/2410.17245.pdf>
103. <https://arxiv.org/pdf/2502.19649.pdf>
104. <https://arxiv.org/html/2502.15507v1>
105. <https://www.alphaxiv.org/abs/2606.06735>
106. <https://arxiv.org/html/2410.12877v2>
107. https://aiwiki.ai/wiki/scaling_monosemanticity
108. <https://arxiv.org/html/2502.17601v1>
109. <http://arxiv.org/pdf/2405.02148.pdf>
110. <https://huggingface.co/papers/2308.10248>
111. <https://medium.com/@jac08h/scaling-monosemanticity-48715e3ef041>
112. <https://arxiv.org/html/2509.22739v3>
113. <https://www.paperfeed.app/2026-W22/arxiv-2605-29358.html>
114. <https://galileo.ai/blog/anthropic-ai-interpretability-breakthrough>
115. <https://research.mental-momentum.ai/r/anthropic-mechanistic-interpretability-2ic9sw>
116. <https://arxiv.org/html/2602.06801v3>
117. <https://www.alphaxiv.org/abs/2312.06681>
118. <https://arxiv.org/pdf/1610.02478.pdf>
119. <https://arxiv.org/vc/arxiv/papers/1812/1812.05556v1.pdf>
120. <https://arxiv.org/pdf/2012.02282.pdf>
121. <https://arxiv.org/pdf/2405.04985.pdf>
122. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10175717/>
123. <https://faculty.cc.gatech.edu/~riedl/pubs/iccc12.pdf>
124. https://computationalcreativity.net/iccc2014/wp-content/uploads/2014/06/14.1_Schorlemmer.pdf
125. <https://staff.fnwi.uva.nl/b.bredeweg/pdf/BSc/20122013/Steinbruck.pdf>
126. https://www.lirmm.fr/ecai2012/images/stories/ecai_doc/pdf/workshop/W40_c3gi_pre-proceedings_20120803.pdf
127. https://www.iiia.csic.es/media/filer_public/54/cc/54cc46e0-1276-4015-801e-34c5eb1d1d15/preprint.pdf
128. <https://eden.dei.uc.pt/~camara/files/FCTSubmission/GameOn03.pdf>
129. <https://computationalcreativity.net/iccc2019/papers/iccc19-paper-6.pdf>
130. <https://www.iiia.csic.es/~enric/papers/Falomir-Plaza-CreativeUnderstanding.pdf>

131. <https://arxiv.org/pdf/2107.00949.pdf>
132. <https://www.iiia.csic.es/~enric/papers/Concept+Invention.pdf>
133. <https://axon.cs.byu.edu/ICCC2015proceedings/7.1Besold.pdf>
134. https://www.scienceopen.com/document_file/4b2425b8-9cc4-424c-b9a6-70db160d59bf/ScienceOpen/354_Upton.pdf
135. https://www.scienceopen.com/document_file/2f4ebd71-01d6-4690-a1f5-f7729de8fbcc/ScienceOpen/322_Notaro.pdf
136. <https://arxiv.org/pdf/2401.16268.pdf>
137. <http://www.m-hikari.com/ces/ces2014/ces21-24-2014/dongjokimCES21-24-2014.pdf>
138. <https://www.mdpi.com/2076-0752/6/2/5/pdf?version=1491913106>
139. [https://zenodo.org/record/4663399/files/Dejan Grba and Vladimir Todorovic - 2021 - Rendering Life - Abstract and References.pdf](https://zenodo.org/record/4663399/files/Dejan_Grba_and_Vladimir_Todorovic_-_2021_-_Rendering_Life_-_Abstract_and_References.pdf)
140. https://www.mitpressjournals.org/doi/pdf/10.1162/isal_a_00215
141. <https://www.frontiersin.org/articles/10.3389/frobt.2014.00008/pdf>
142. <https://arxiv.org/pdf/1311.1213.pdf>
143. <https://en.wikipedia.org/wiki/Lenia>
144. <https://hammyasf.github.io/lenia.html>
145. <https://www.mysimulator.uk/content/tutorials/gray-scott-gpu.html>
146. <https://statox.github.io/particles-life/reaction-diffusion>
147. <https://lumitree.art/blog/cellular-automata>
148. <http://www.mrob.com/pub/comp/xmorphism/ogl/index.html>
149. <https://alife.org/encyclopedia/software-platforms/lenia/>
150. <https://github.com/piellardj/reaction-diffusion-webgl>
151. <https://github.com/amandaghassaei/ReactionDiffusionShader>
152. <https://hegl.mathi.uni-heidelberg.de/continuous-cellular-automata/>
153. <https://arxiv.org/pdf/1410.0281.pdf>
154. <https://arxiv.org/vc/arxiv/papers/1812/1812.05433v1.pdf>
155. <https://gist.github.com/robinhouston/ed597847175cf692ecce>
156. <https://direct.mit.edu/artl/article/30/1/106/119728/A-Survey-of-Recent-Practice-of-Artificial-Life-in>
157. <http://arxiv.org/pdf/1706.07068.pdf>
158. <https://arxiv.org/html/2504.04298v1>
159. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10914360/>
160. <https://dl.acm.org/doi/pdf/10.1145/3613904.3641971>
161. <https://computerresearch.org/index.php/computer/article/view/102283>
162. <https://www.mdpi.com/2076-0752/7/3/25/pdf?version=1531210008>
163. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10946898/>
164. <https://www.mdpi.com/2075-5309/11/1/29/pdf>
165. <https://arxiv.org/html/2405.04144v2>
166. <https://dl.acm.org/doi/pdf/10.1145/3613904.3642461>

167. <https://www.alphaxiv.org/abs/2501.09815>

168. <https://escholarship.org/content/qt1df6204v/qt1df6204v.pdf>

169. <https://www.emergentmind.com/papers/2406.12027>

170. https://link.springer.com/article/10.1007/s00146-025-02799-5?error=cookies_not_supported&code=e42552f3-6259-41e4-8562-2b0fa6fc1869

171. <https://arxiv.org/html/2106.10800v5>

172. <https://blog.ai.princeton.edu/2026/05/15/i-am-myself-a-lossy-compression/>

173. <https://attachments.are.na/35094923/abd5a9e1c63eeca05c663414769cba01.pdf?1741545276>

174. <https://bohrium.dp.tech/paper/arxiv/2401.15497>

175. <https://arxiv.org/pdf/2104.02726.pdf>

176. <https://pdfs.semanticscholar.org/83e3/6cc3360d18afcbeb960660d9c31cee703750.pdf>

177. <http://arxiv.org/pdf/2207.04857.pdf>

178. <https://arxiv.org/pdf/2003.11624.pdf>

179. <https://arxiv.org/pdf/2102.13434.pdf>

180. <http://arxiv.org/pdf/2408.13719.pdf>

181. <https://arxiv.org/pdf/2412.03227.pdf>

182. <https://arxiv.org/pdf/2209.03618.pdf>

183. https://www.techrxiv.org/articles/preprint/Looking_For_Novelty_in_Search-based_Software_Product_Line_Testing/12789203/files/24193250.pdf

184. <https://joellehman.com/lehman-dissertation.pdf>

185. <https://arxiv.org/pdf/2605.30570.pdf>

186. <https://arxiv.org/pdf/1609.03357.pdf>

187. <https://web.stanford.edu/~mrifaki/quality-diversity-readings.html>

188. https://dl.acm.org/doi/10.1162/EVCO_a_00025

189. <https://www.alphaxiv.org/abs/2504.08057v3>

190. <https://direct.mit.edu/evco/article/32/3/249/116787/Preliminary-Analysis-of-Simple-Novelty-Search>

191. <https://ar5iv.labs.arxiv.org/html/2003.03377>

192. <https://inria.hal.science/hal-02555202/file/multitask.pdf>

193. <https://www.cs.swarthmore.edu/~meeden/cs81/f17/projects/JeffKyle.pdf>

194. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10282580/>

195. <https://napier-repository.worktribe.com/preview/2296985/Urquhart.pdf>

196. https://www.eecs.tufts.edu/~jsinapov/papers/Gizzi_ICCC_2020.pdf

197. https://computationalcreativity.net/iccc2018/sites/default/files/papers/IICCC_2018_paper_7.pdf

198. <https://arxiv.org/abs/2107.05599>

199. <https://computationalcreativity.net/ijwcc08/ijwcc08proceedings.pdf>

200. <https://computationalcreativity.net/iccc20/papers/164-iccc20.pdf>

201. <https://ar5iv.labs.arxiv.org/html/2202.01961>

202. <https://homepages.abdn.ac.uk/g.ritchie/pages/papers/newgencomp.pdf>

203. <https://qmro.qmul.ac.uk/xmlui/handle/123456789/105581>

204. <https://repository.ubn.ru.nl/bitstream/handle/2066/132144/132144.pdf>

205. <https://arxiv.org/pdf/2211.08861.pdf>

206. https://antoniosliapis.com/research/research_divsearch.php

207. <https://dspace.cuni.cz/bitstream/handle/20.500.11956/193425/120486128.pdf?sequence=1&isAllowed=y>

208. <https://theses.liacs.nl/pdf/2021-2022-RicchiutiMauro.pdf>

209. <https://arxiv.org/html/1907.04053v1>

210. https://computationalcreativity.net/iccc22/papers/ICCC-2022_paper_60.pdf

211. <http://arxiv.org/pdf/2308.08576.pdf>

212. <https://arxiv.org/pdf/2205.01605.pdf>

213. <http://arxiv.org/pdf/1805.11057.pdf>

214. <https://arxiv.org/pdf/2205.10271.pdf>

215. <http://arxiv.org/pdf/2410.14925.pdf>

216. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10584976/>

217. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11982568/>

218. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11639127/>

219. <https://aprra.net/article/download/115426/163737>

220. <https://www.anthropocenes.net/article/666/galley/4360/download/>

221. <https://aprra.net/article/view/115426>

222. https://openresearch.lsbu.ac.uk/download/82a7bd30bda69f2a71cf5e9606f78f254385f992af12d92b9c0be17028d0d9bc/49067766/2018_PhD_Lynch.pdf

223. <https://cinej.pitt.edu/ojs/cinej/article/download/905/803>

224. <https://www.anthropocenes.net/article/id/666/>

225. https://ejournals.eu/en/journal_article_files/full_text/0197d00e-464e-734a-8a5e-b214c6aa8716/download

226. <https://median.newmediacaucus.org/caa-edition/glitch-as-an-algorithmic-art/>

227. <https://www.yin.roma.it/2026/08/who-creates-whom-art-ai-and-the-transformation-of-the-maker/>

228. <https://www.iiisci.org/journal/pdv/sci/pdfs/OL109XF.pdf>

229. <https://billseaman.com/Papers/Ascott - Early Interactive Work and Some Cybernetic Relationalities 2018 final.pdf>

230. https://www.academia.edu/1411717/Datamoshing_and_the_emergence_of_digital_complexity_from_digital_chaos

231. <https://research-repository.griffith.edu.au/server/api/core/bitstreams/be9c4da7-726f-5fed-b07b-76876982e226/content>

232. <https://unitesi.unive.it/retrieve/fca8119b-8b6d-457f-92c2-06db46992013/867800-1272313.pdf>

233. https://openresearch.ocadu.ca/id/eprint/3353/1/Li_Jun_2021_MDES_DIGF.pdf

234. <https://arxiv.org/pdf/2310.09336.pdf>

235. <https://arxiv.org/html/2410.16719>

236. <https://arxiv.org/html/2502.04549v1>

237. <https://arxiv.org/html/2404.12141>

238. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10783868/>

239. <http://arxiv.org/pdf/2403.07902.pdf>

240. <https://arxiv.org/html/2408.13256v2>

241. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10443900/>

242. <https://arxiv.org/html/2406.07844v2>

243. <https://openreview.net/pdf/fe0e2f5e09a4f719c6cca20e6d6350ee6d6c0900.pdf>

244. https://downloads.webis.de/theses/papers/imambayeva_2024.pdf

245. https://openaccess.thecvf.com/content/ICCV2025/papers/Dat_VSC_Visual_Search_Compositional_Text-to-Image_Diffusion_Model_ICCV_2025_paper.pdf

246. <https://neurips.cc/virtual/2025/poster/115203>

247. <https://openreview.net/pdf?id=PUlqjT4rzq7>

248. <https://web3.arxiv.org/pdf/2508.20783>

249. <https://arxiv.org/html/2607.03831v1>

250. https://proceedings.neurips.cc/paper_files/paper/2023/file/f8ad010cdd9143dbb0e9308c093aff24-Paper-Datasets_and_Benchmarks.pdf

251. <https://arxiv.org/html/2604.04575v1>

252. <https://arxiv.org/html/2410.22775v2>

253. https://proceedings.neurips.cc/paper_files/paper/2024/file/cdccc6d47c1627350014a3076112ab824-Paper-Conference.pdf

254. <https://openreview.net/pdf?id=VQYTs3Ymtu>

255. <https://www.emergentmind.com/topics/binding-id-mechanism-in-vlms>

256. <https://openreview.net/forum?id=4mzGiMooXM&referrer=the%20profile%20of%20Pan%20Gao>

257. <http://arxiv.org/pdf/2401.01651.pdf>

258. <https://arxiv.org/html/2501.08545v6>

259. <https://arxiv.org/pdf/2206.03429.pdf>

260. <https://arxiv.org/html/2412.16211v1>

261. <https://arxiv.org/html/2504.04907v1>

262. <https://arxiv.org/html/2407.21408v2>

263. <https://arxiv.org/html/2502.20694v1>

264. <http://arxiv.org/pdf/2403.15377v4.pdf>

265. <https://arxiv.org/html/2508.00144v1>

266. <https://arxiv.org/html/2507.18107v1>

267. <https://www.alphaxiv.org/abs/2606.27537v4>

268. <https://world-bench.github.io/>